

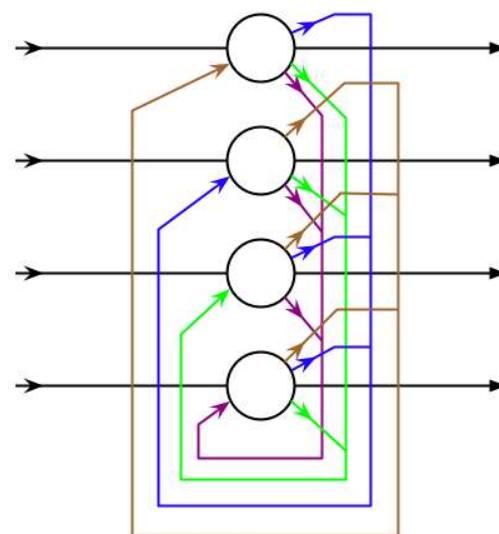
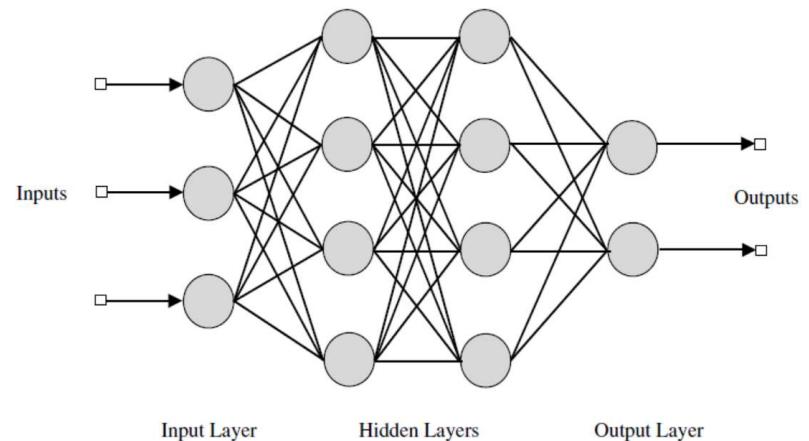
# **Neural Networks**

**Hopfield Nets and Auto Associators**

**Fall 2022**

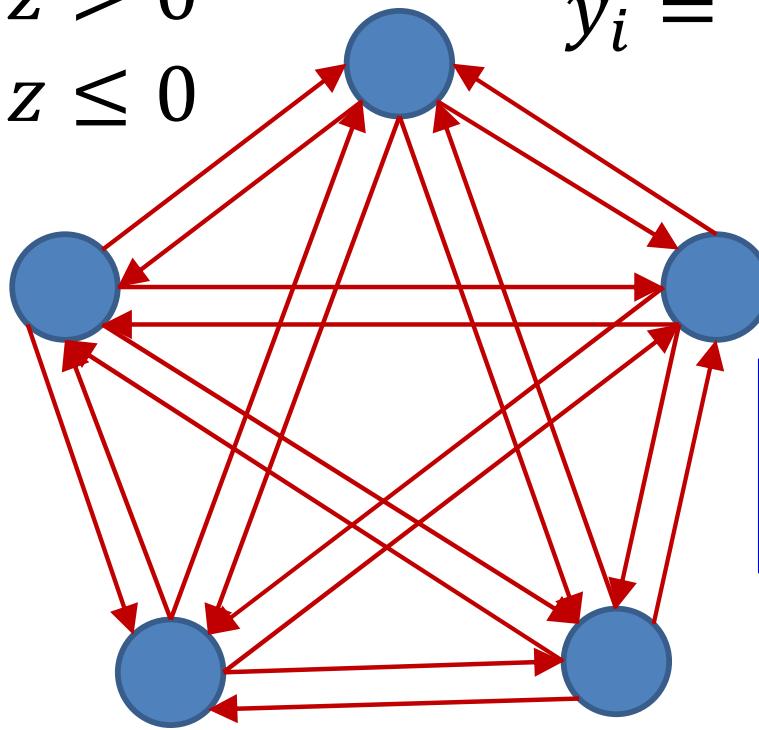
# Story so far

- Neural networks for computation
- All feedforward structures
- But what about..



# Consider this loopy network

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$

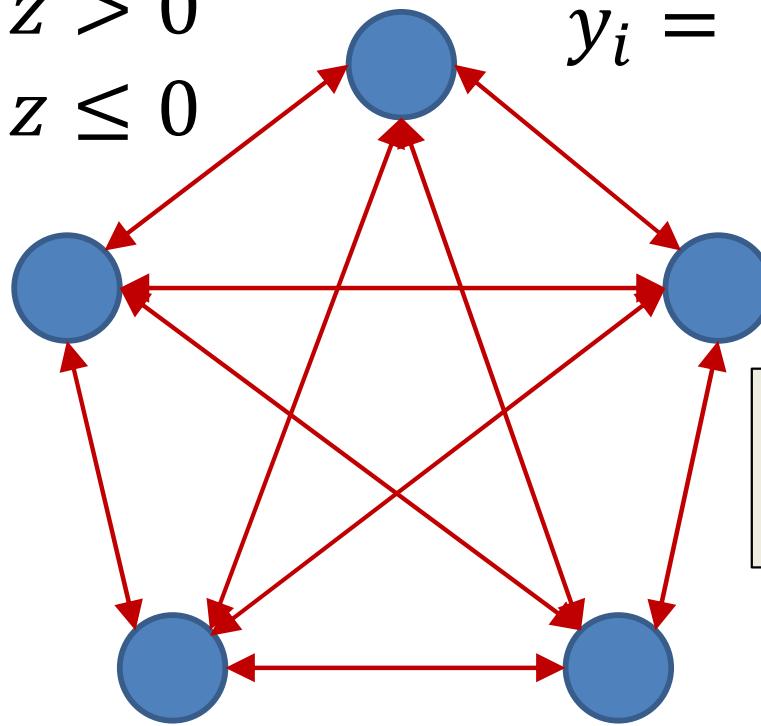


The output of a neuron affects the input to the neuron

- Each neuron is a perceptron with +1/-1 output
- Every neuron *receives* input from every other neuron
- Every neuron *outputs* signals to every other neuron

# Consider this loopy network

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$

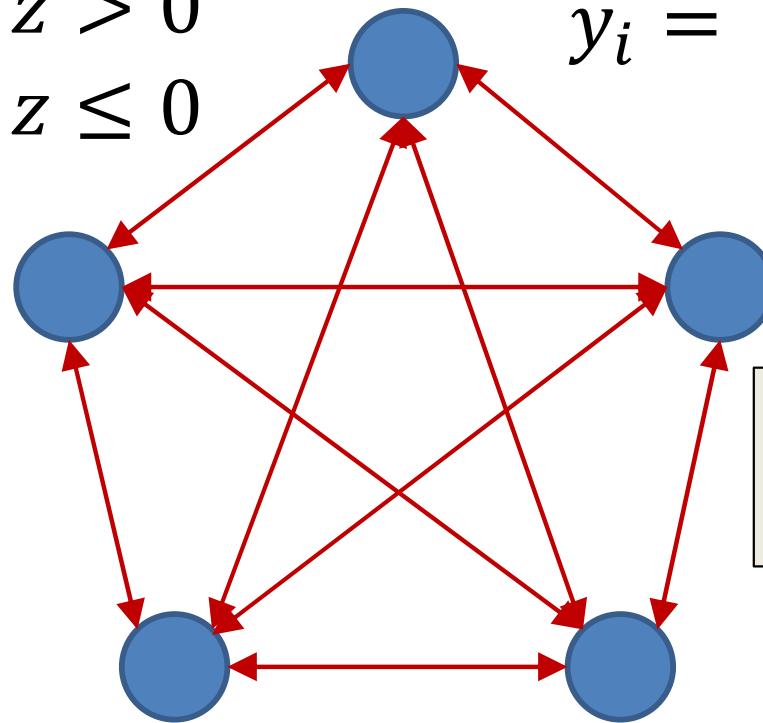


A symmetric network:  
 $w_{ij} = w_{ji}$

- Each neuron is a perceptron with +1/-1 output
- Every neuron *receives* input from every other neuron
- Every neuron *outputs* signals to every other neuron

# Hopfield Net

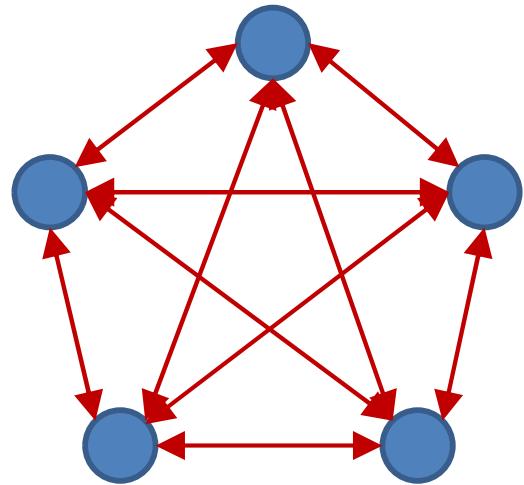
$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$



A symmetric network:  
 $w_{ij} = w_{ji}$

- Each neuron is a perceptron with +1/-1 output
- Every neuron *receives* input from every other neuron
- Every neuron *outputs* signals to every other neuron

# Loopy network

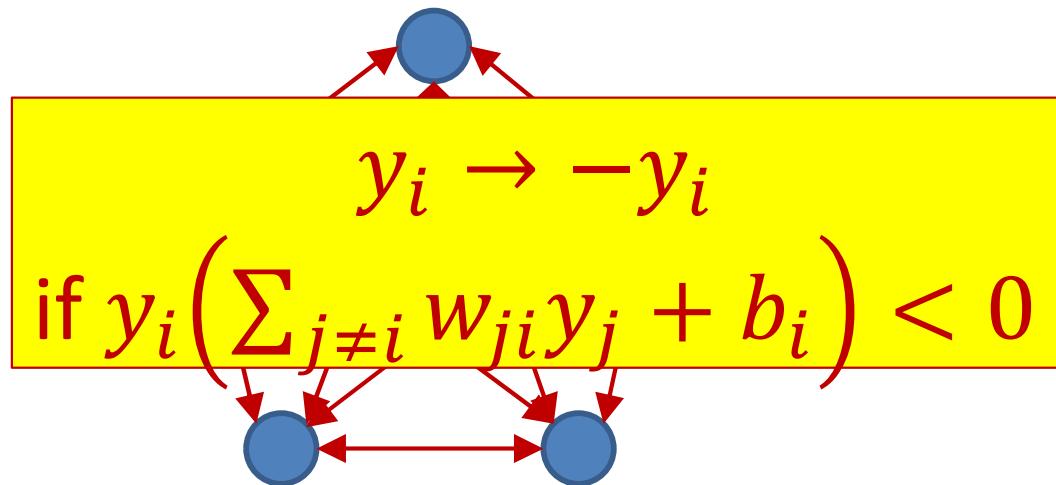


$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

- At each time each neuron receives a “field”  $\sum_{j \neq i} w_{ji} y_j + b_i$
- If the sign of the field matches its own sign, it does not respond
- If the sign of the field opposes its own sign, it “flips” to match the sign of the field

# Loopy network

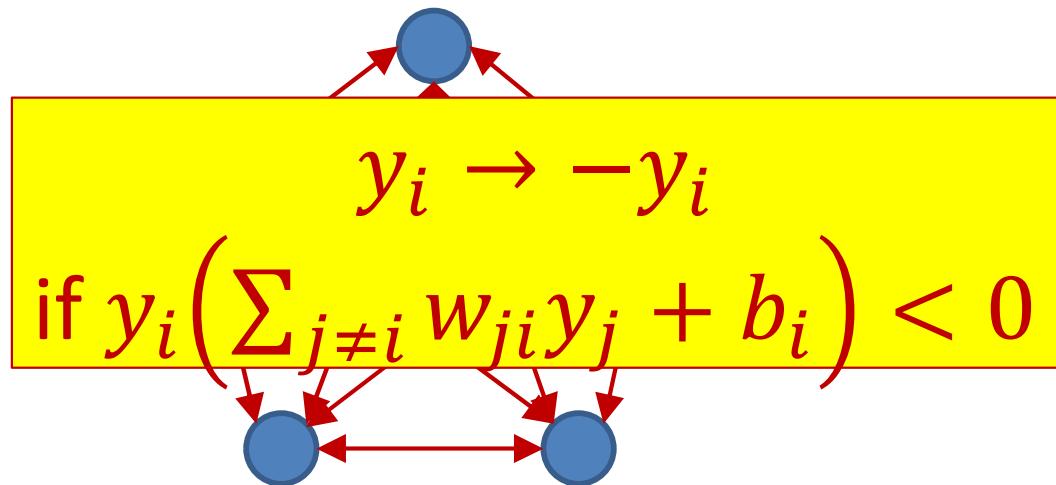


$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

- At each time each neuron receives a “field”  $\sum_{j \neq i} w_{ji} y_j + b_i$
- If the sign of the field matches its own sign, it does not respond
- If the sign of the field opposes its own sign, it “flips” to match the sign of the field

# Loopy network



$$y_i = \Theta \left( \sum_{j \neq i} w_{ji}y_j + b_i \right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

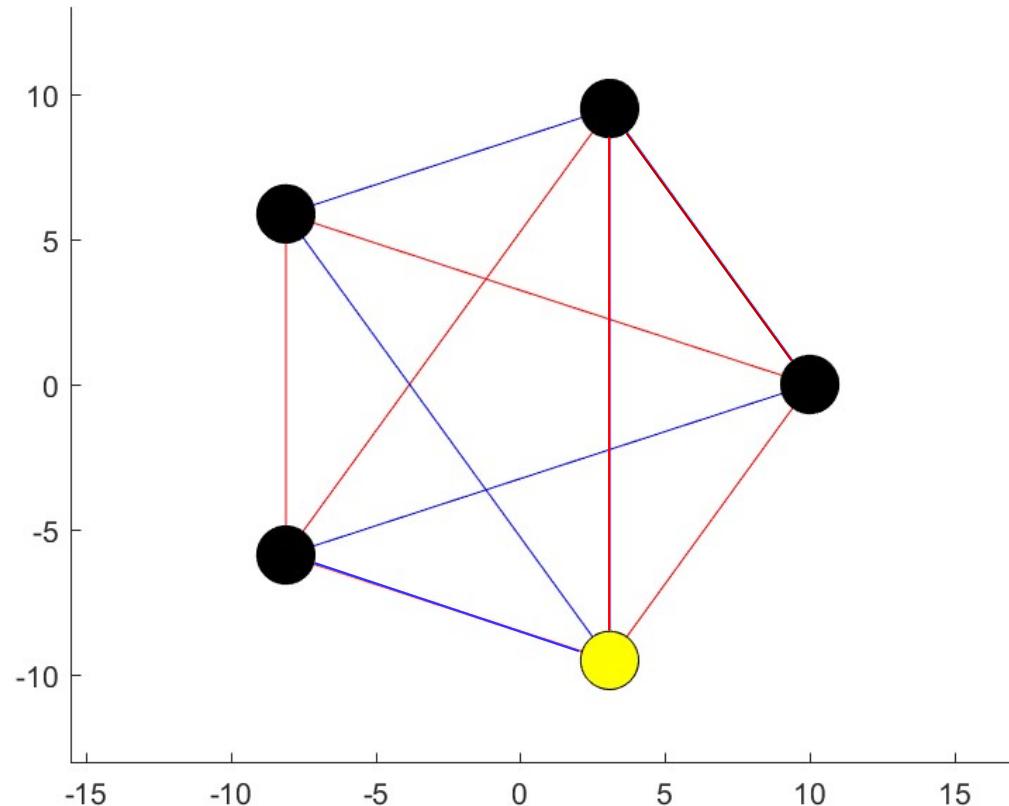
- A neuron “flips” if weighted sum of other neurons’ outputs is of the opposite sign to its own current (output) value

But this may cause other neurons to flip!

es a “field”  $\sum_{j \neq i} w_{ji}y_j + b_i$   
s own sign, it does not

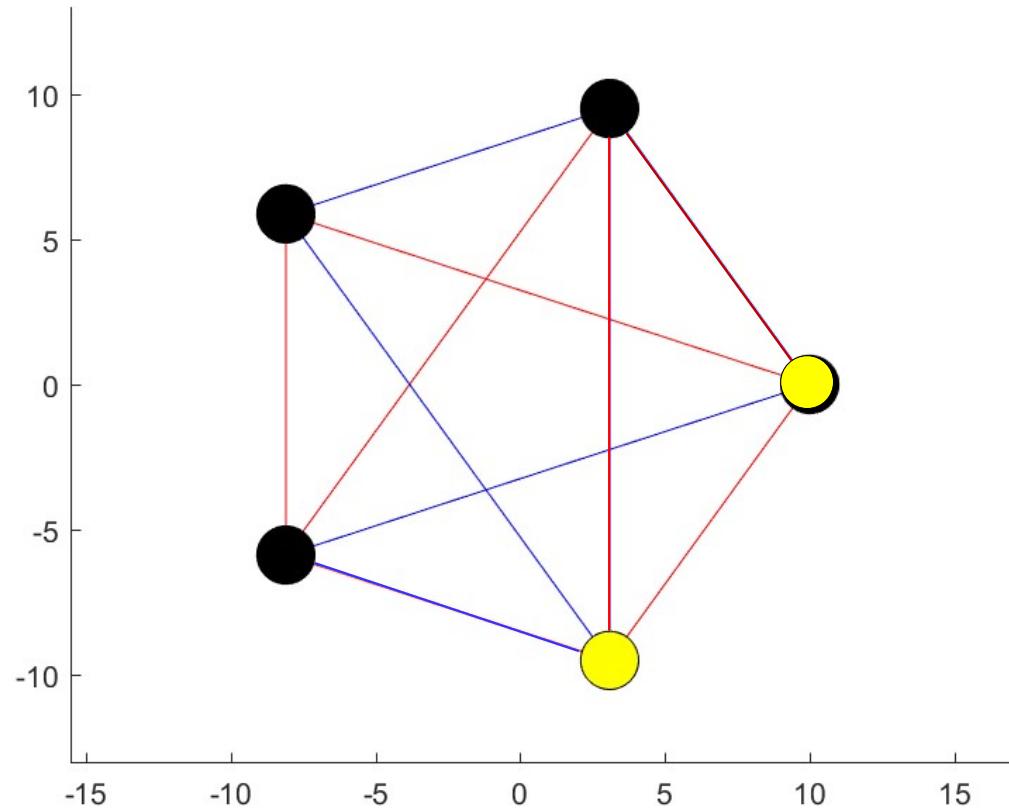
- If the sign of the field opposes its own sign, it “flips” to match the sign of the field

# Example



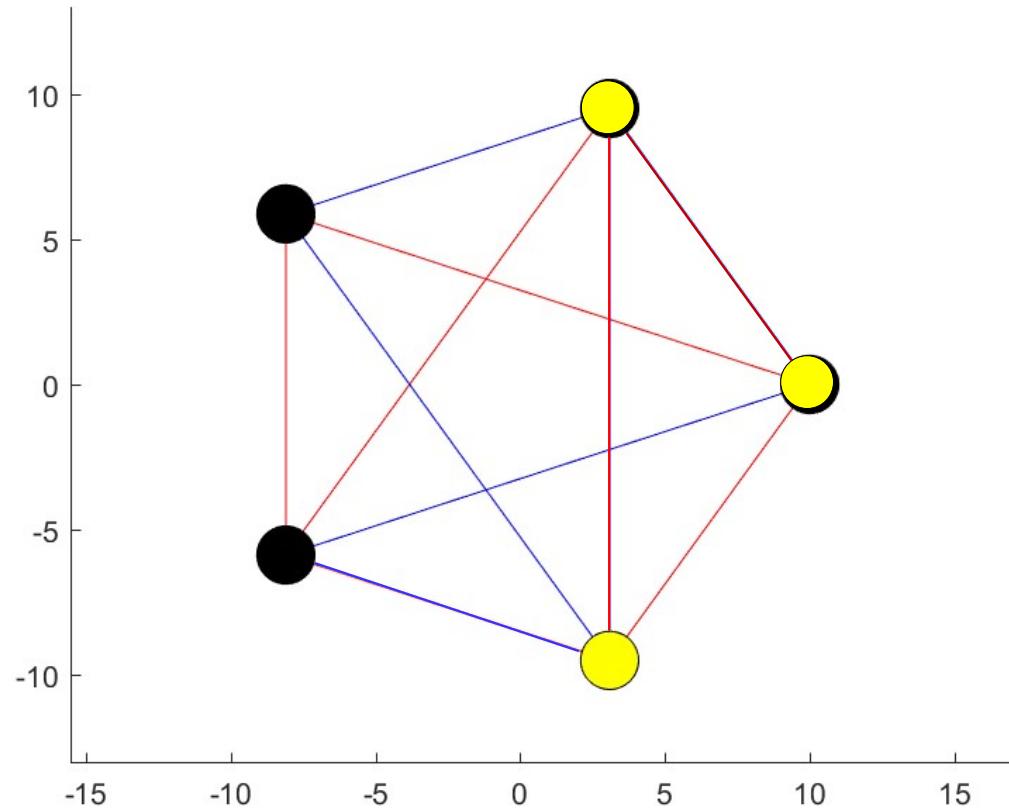
- Red edges are  $+1$ , blue edges are  $-1$
- Yellow nodes are  $-1$ , black nodes are  $+1$

# Example



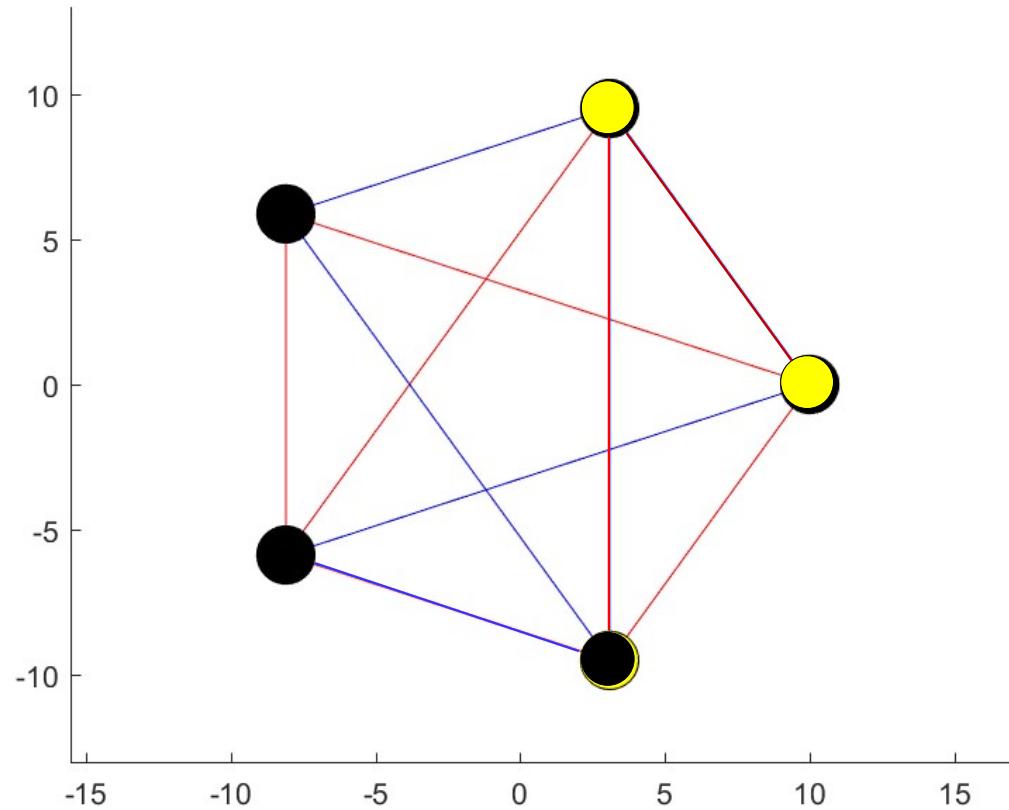
- Red edges are  $+1$ , blue edges are  $-1$
- Yellow nodes are  $-1$ , black nodes are  $+1$

# Example



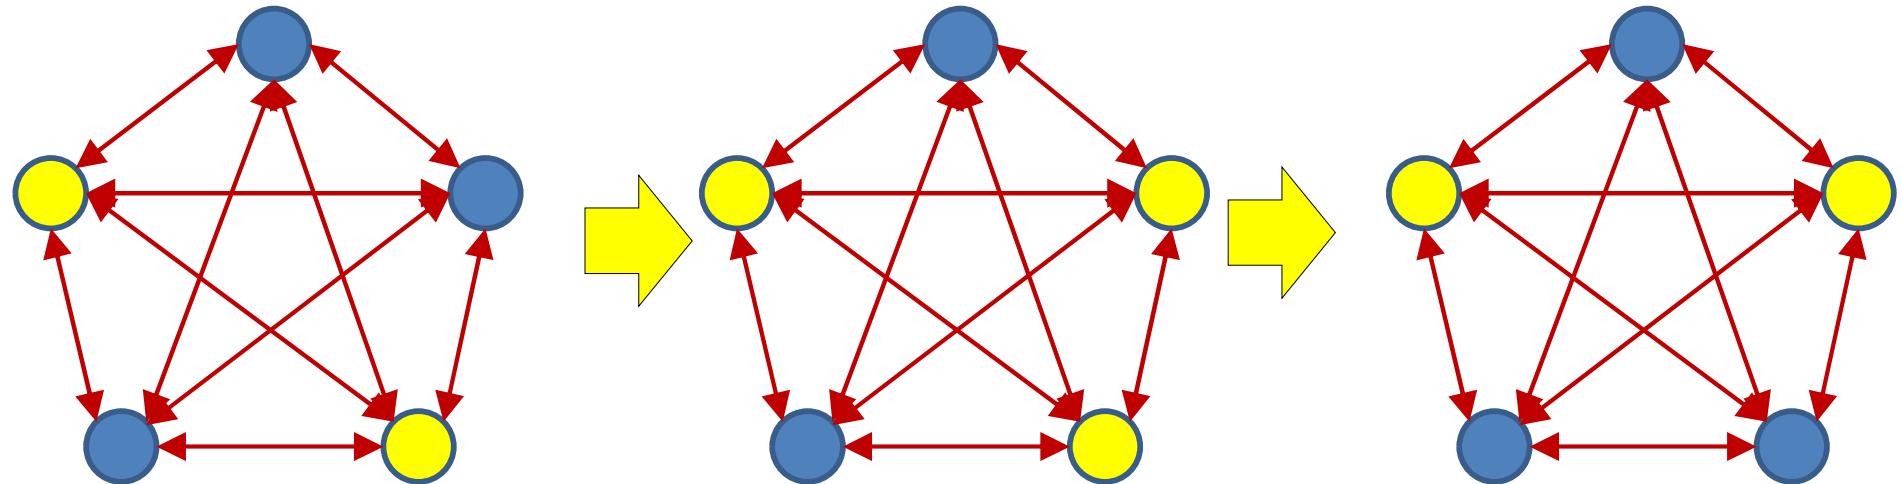
- Red edges are +1, blue edges are -1
- Yellow nodes are -1, black nodes are +1

# Example



- Red edges are +1, blue edges are -1
- Yellow nodes are -1, black nodes are +1

# Loopy network

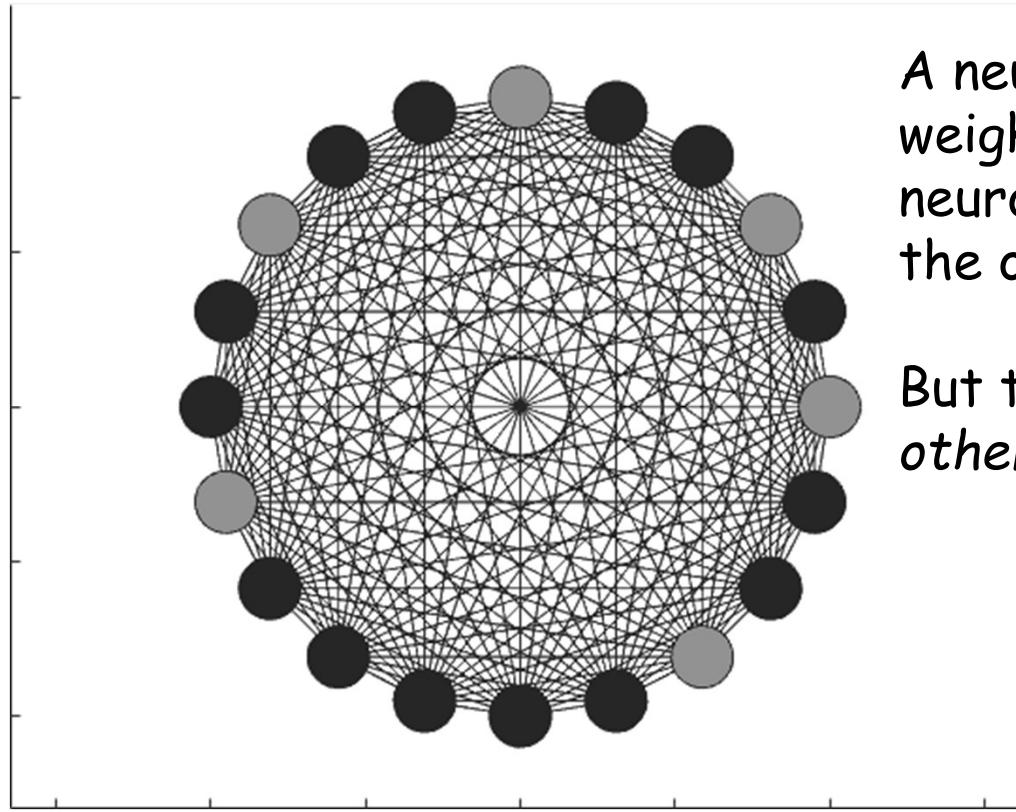


- If the sign of the field at any neuron opposes its own sign, it “flips” to match the field
  - Which will change the field at other nodes
    - Which may then flip
      - Which may cause other neurons including the first one to flip...
        - » And so on...

# 20 evolutions of a loopy net

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$

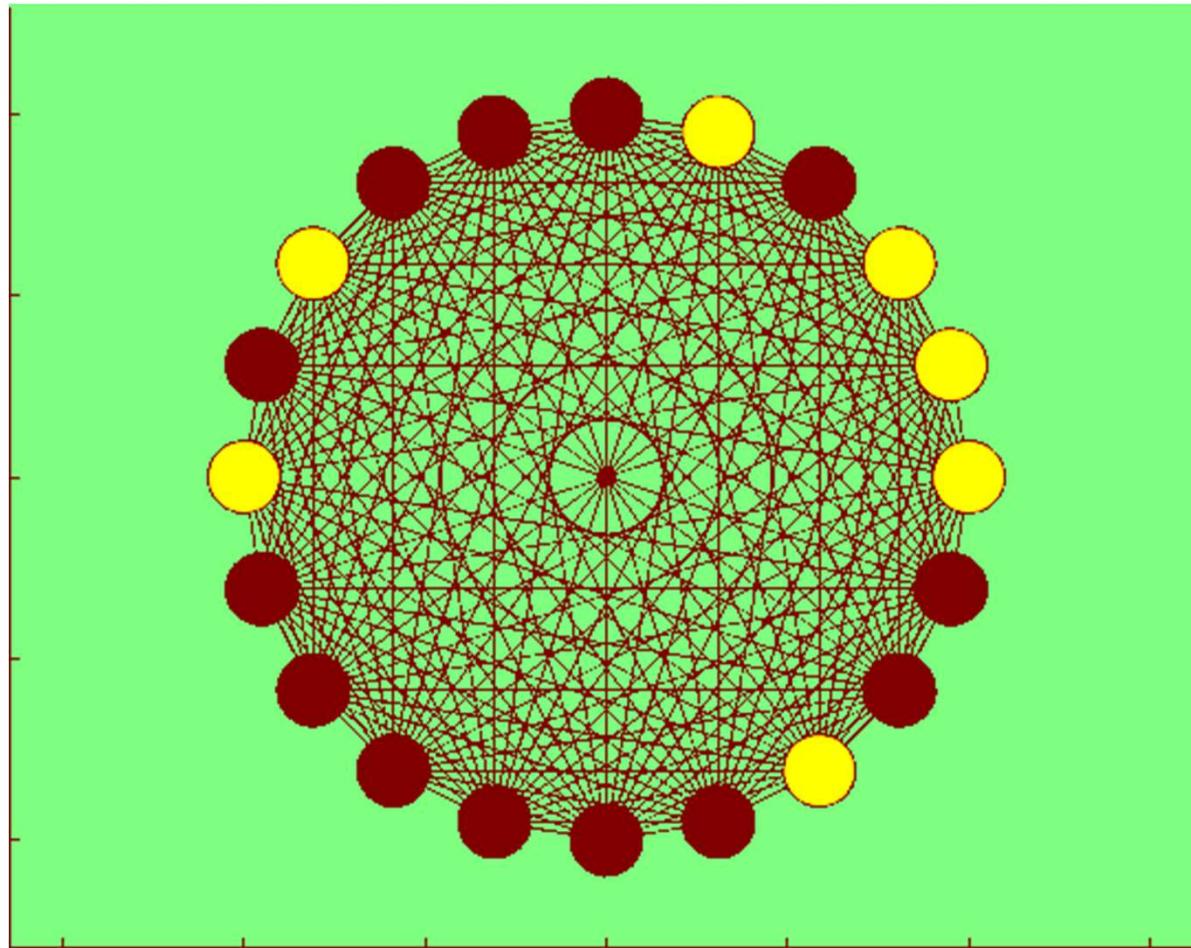


A neuron “flips” if weighted sum of other neuron’s outputs is of the opposite sign

But this may cause other neurons to flip!

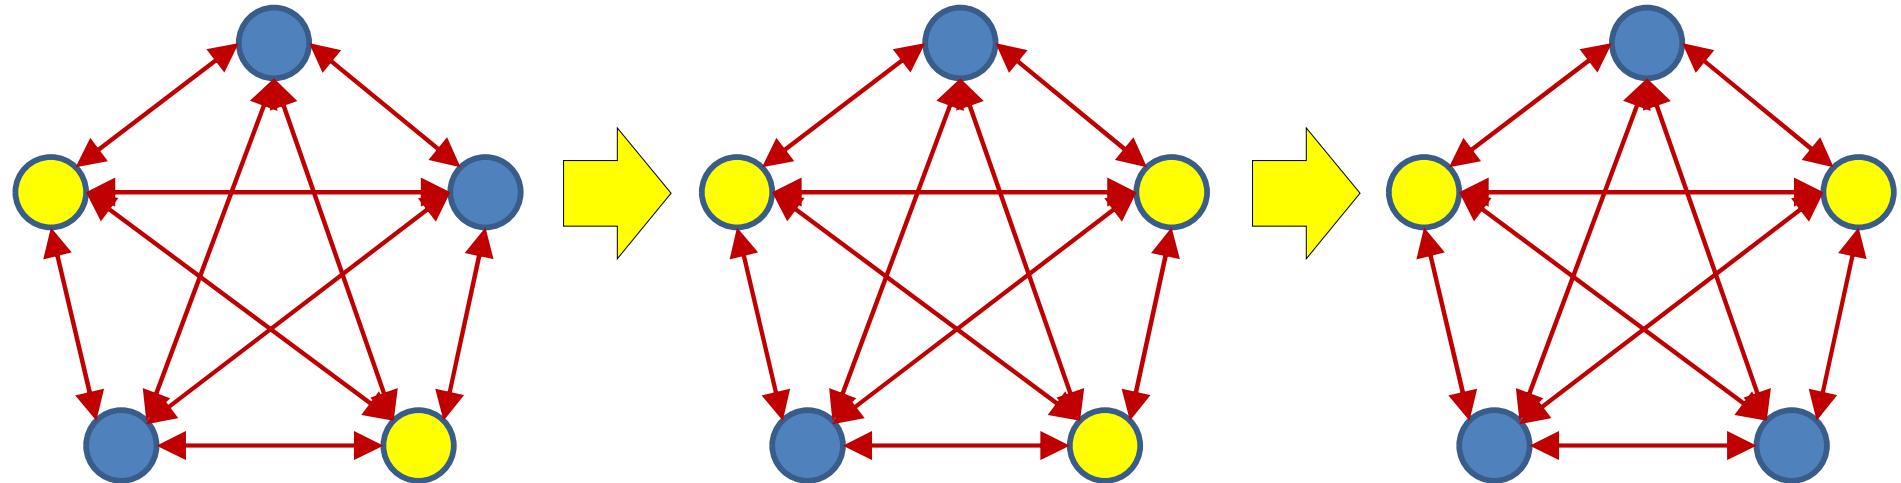
- All neurons which do not “align” with the local field “flip”

# 120 evolutions of a loopy net



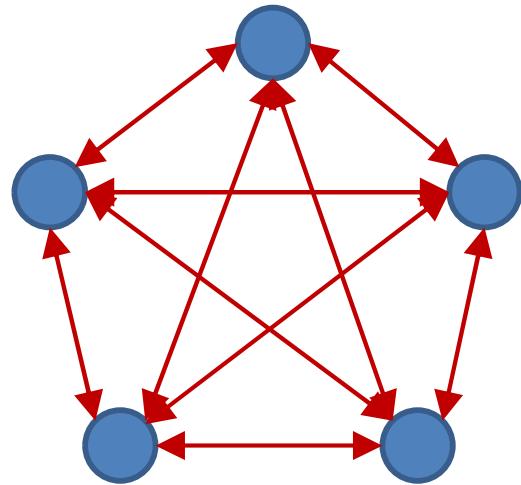
- All neurons which do not “align” with the local field “flip”

# Loopy network



- If the sign of the field at any neuron opposes its own sign, it “flips” to match the field
  - Which will change the field at other nodes
    - Which may then flip
      - Which may cause other neurons including the first one to flip...
- *Will this behavior continue for ever??*

# Loopy network



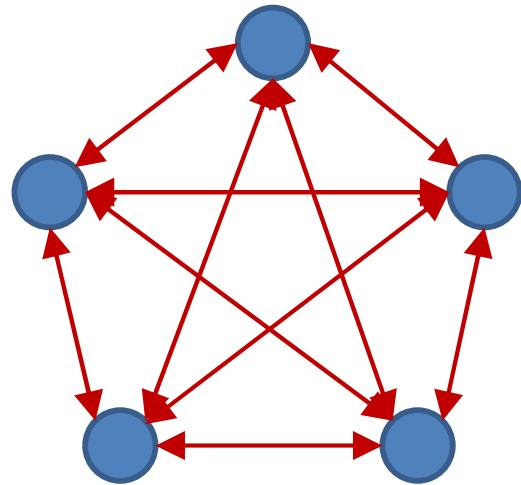
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

- Let  $y_i^-$  be the output of the  $i$ -th neuron just *before* it responds to the current field
- Let  $y_i^+$  be the output of the  $i$ -th neuron just *after* it responds to the current field
- If  $y_i^- = \text{sign}\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$ , then  $y_i^+ = y_i^-$ 
  - If the sign of the field matches its own sign, it does not flip

$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 0$$

# Loopy network



$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

- If  $y_i^- \neq \text{sign}\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$ , then  $y_i^+ = -y_i^-$

$$y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) - y_i^- \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) = 2y_i^+ \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

- This term is always positive!
- Every flip of a neuron is guaranteed to locally increase*

$$y_i \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

# Globally

- Consider the following sum across *all* nodes

$$\begin{aligned} D(y_1, y_2, \dots, y_N) &= \sum_i y_i \left( \sum_{j \neq i} w_{ji} y_j + b_i \right) \\ &= \sum_{i,j \neq i} w_{ij} y_i y_j + \sum_i b_i y_i \end{aligned}$$

- Assume  $w_{ii} = 0$
- For any unit  $k$  that “flips” because of the local field

$$\Delta D(y_k) = D(y_1, \dots, y_k^+, \dots, y_N) - D(y_1, \dots, y_k^-, \dots, y_N)$$

- This is strictly positive

$$\Delta D(y_k) = 2y_k^+ \left( \sum_{j \neq k} w_{jk} y_j + b_k \right)$$

# Upon flipping a single unit

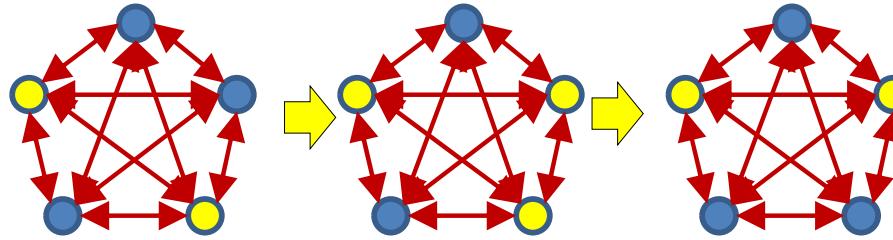
$$\Delta D(y_k) = D(y_1, \dots, y_k^+, \dots, y_N) - D(y_1, \dots, y_k^-, \dots, y_N)$$

- Expanding

$$\Delta D(y_k) = (y_k^+ - y_k^-) \left( \sum_{j \neq k} w_{jk} y_j + b_k \right)$$

- All other terms that do not include  $y_k$  cancel out
- This is always positive!
- *Every flip of a unit results in an increase in  $D$*

# Hopfield Net



- Flipping a unit will result in an increase (non-decrease) of

$$D = \sum_{i,j \neq i} w_{ij} y_i y_j + \sum_i b_i y_i$$

- $D$  is bounded

$$D_{max} = \sum_{i,j \neq i} |w_{ij}| + \sum_i |b_i|$$

- The minimum increment of  $D$  in a flip is

$$\Delta D_{min} = \min_{i, \{y_i, i=1..N\}} 2 \left| \sum_{j \neq i} w_{ji} y_j + b_i \right|$$

- Any sequence of flips must converge in a finite number of steps

# The Energy of a Hopfield Net

- Define the *Energy* of the network as

$$E = -\frac{1}{2} \left( \sum_{i,j \neq i} w_{ij} y_i y_j - \sum_i b_i y_i \right)$$

- Just 0.5 times the negative of  $D$ 
  - The 0.5 is only needed for convention
- The evolution of a Hopfield network constantly decreases its energy

# Story so far

- A Hopfield network is a loopy binary network with symmetric connections
- Every neuron in the network attempts to “align” itself with the sign of the weighted combination of outputs of other neurons
  - The local “field”
- Given an initial configuration, neurons in the net will begin to “flip” to align themselves in this manner
  - Causing the field at other neurons to change, potentially making them flip
- Each evolution of the network is guaranteed to decrease the “energy” of the network
  - The energy is lower bounded and the decrements are upper bounded, so the network is guaranteed to converge to a stable state in a finite number of steps

# Poll 1 @1794, @1795, @1796

Hopfield networks are loopy networks whose output activations “evolve” over time

- True
- False

Hopfield networks will evolve continuously, forever

- True
- False

Hopfield networks can also be viewed as infinitely deep shared parameter MLPs

- True
- False

# Poll 1

Hopfield networks are loopy networks whose output activations “evolve” over time

- True
- False

Hopfield networks will evolve continuously, forever

- True
- False

Hopfield networks can also be viewed as infinitely deep shared parameter MLPs

- True
- False

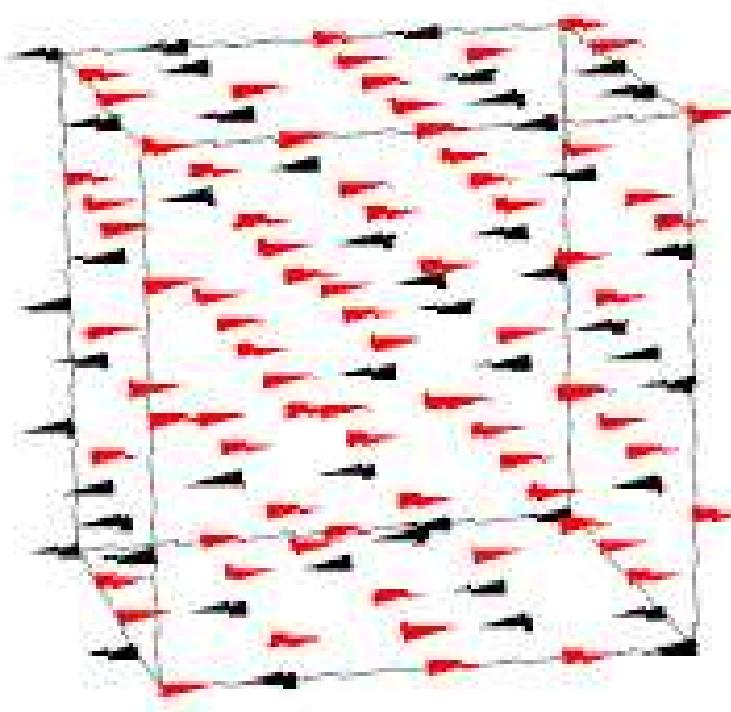
# The Energy of a Hopfield Net

- Define the *Energy* of the network as

$$E = -\frac{1}{2} \left( \sum_{i,j \neq i} w_{ij} y_i y_j - \sum_i b_i y_i \right)$$

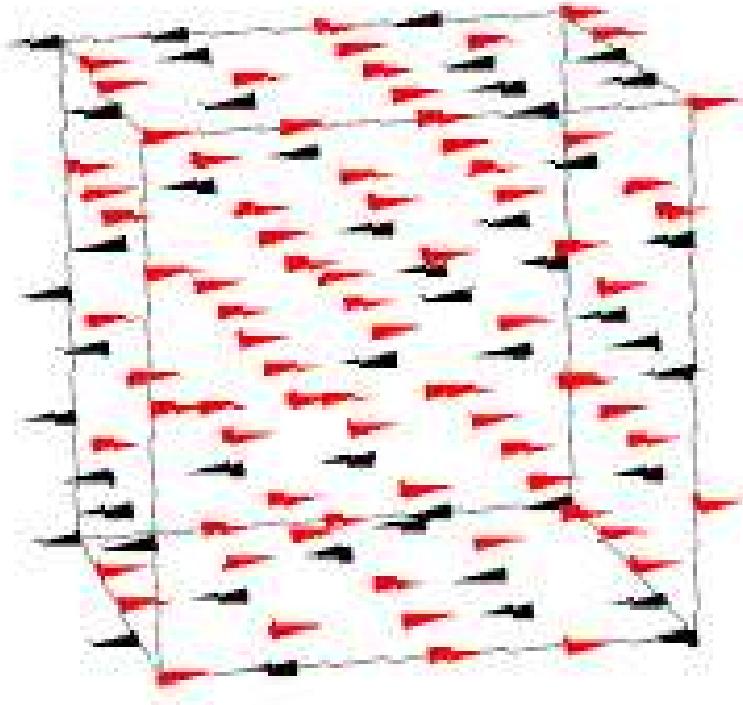
- Just 0.5 times the negative of  $D$
- The evolution of a Hopfield network constantly decreases its energy
- Where did this “energy” concept suddenly sprout from?

# Analogy: Spin Glass



- Magnetic dipoles in a disordered magnetic material
- Each dipole tries to *align* itself to the local field
  - In doing so it may flip
- This will change fields at *other* dipoles
  - Which may flip
- Which changes the field at the current dipole...

# Analogy: Spin Glasses



Total field at current dipole:

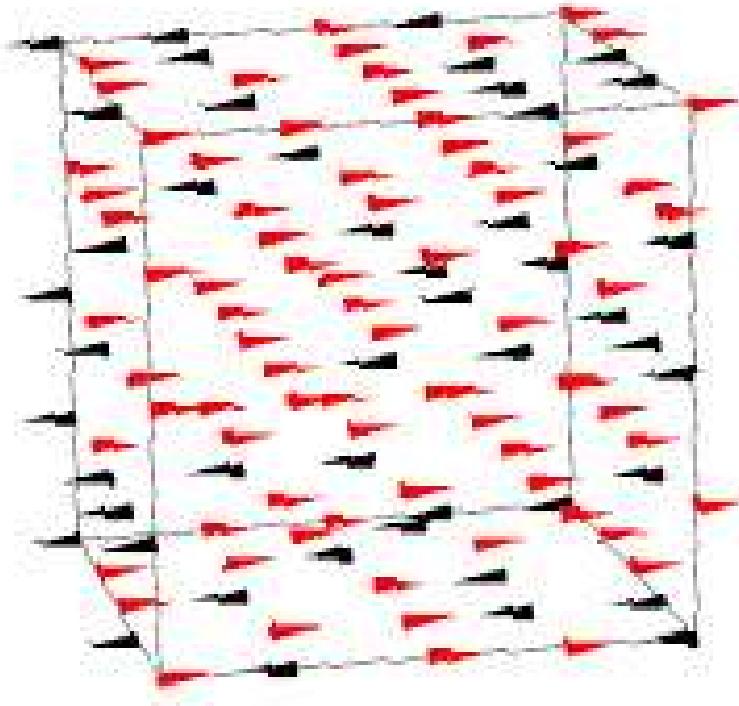
$$f(p_i) = \sum_{j \neq i} J_{ji} x_j + b_i$$

intrinsinc

external

- $p_i$  is vector position of  $i$ -th dipole
- The field at any dipole is the sum of the field contributions of all other dipoles
- The contribution of a dipole to the field at any point depends on interaction  $J$ 
  - Derived from the “Ising” model for magnetic materials (Ising and Lenz, 1924)

# Analogy: Spin Glasses



- A Dipole flips if it is misaligned with the field in its location

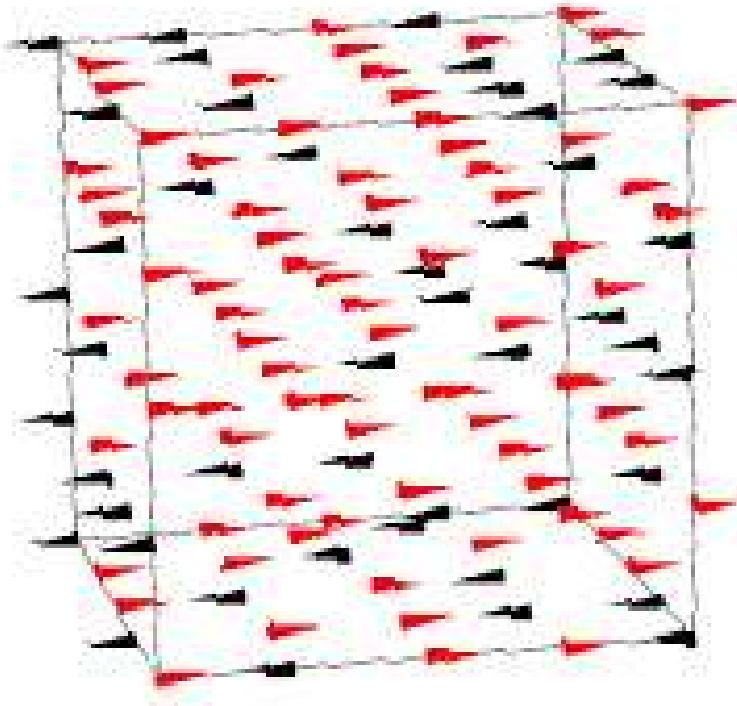
Total field at current dipole:

$$f(p_i) = \sum_{j \neq i} J_{ji}x_j + b_i$$

Response of current dipole

$$x_i = \begin{cases} x_i & \text{if } \text{sign}(x_i f(p_i)) = 1 \\ -x_i & \text{otherwise} \end{cases}$$

# Analogy: Spin Glasses



- Dipoles will keep flipping
  - A flipped dipole changes the field at other dipoles
    - Some of which will flip
  - Which will change the field at the current dipole
    - Which may flip
  - Etc..

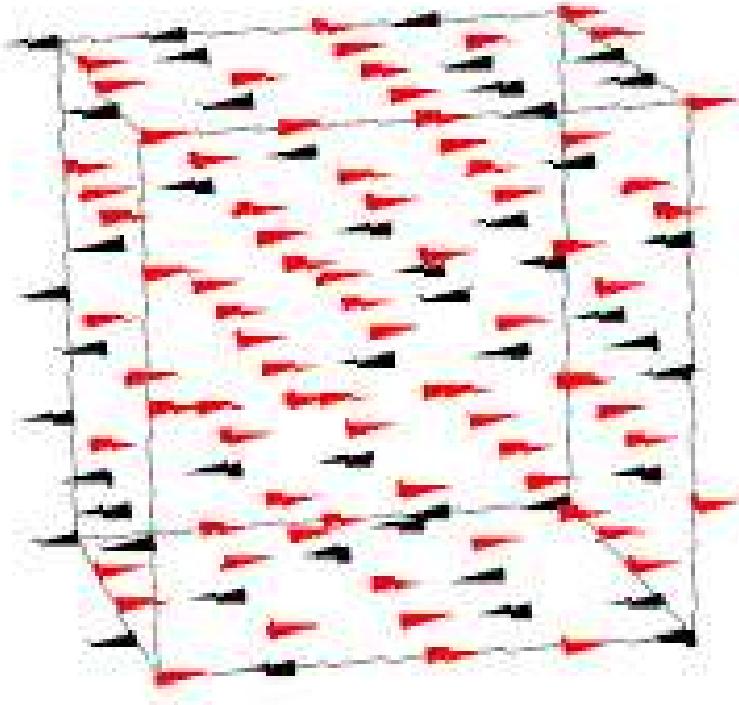
Total field at current dipole:

$$f(p_i) = \sum_{j \neq i} J_{ji}x_j + b_i$$

Response of current dipole

$$x_i = \begin{cases} x_i & \text{if } \text{sign}(x_i f(p_i)) = 1 \\ -x_i & \text{otherwise} \end{cases}$$

# Analogy: Spin Glasses



Total field at current dipole:

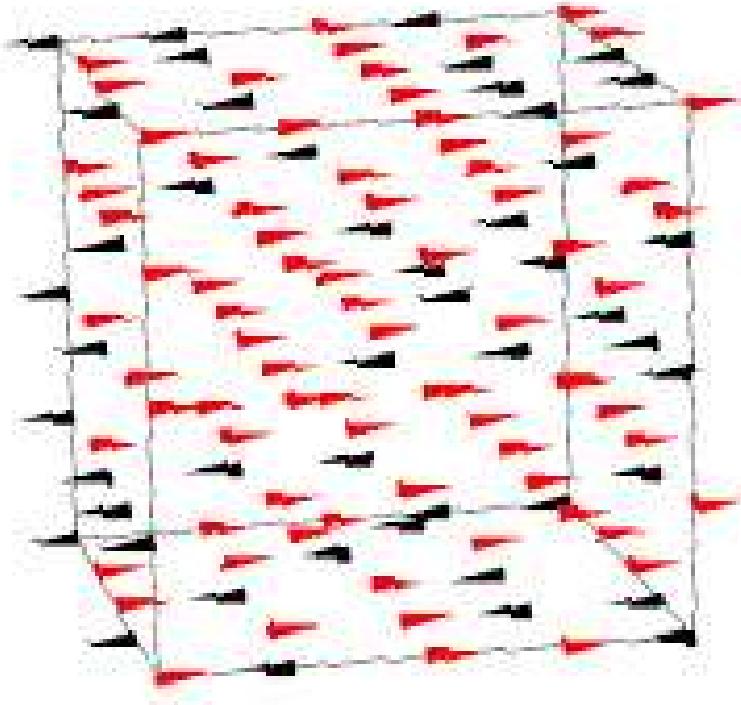
$$f(p_i) = \sum_{j \neq i} J_{ji} x_j + b_i$$

Response of current dipole

$$x_i = \begin{cases} x_i & \text{if } \text{sign}(x_i f(p_i)) = 1 \\ -x_i & \text{otherwise} \end{cases}$$

- When will it stop???

# Analogy: Spin Glasses



Total field at current dipole:

$$f(p_i) = \sum_{j \neq i} J_{ji}x_j + b_i$$

Response of current dipole

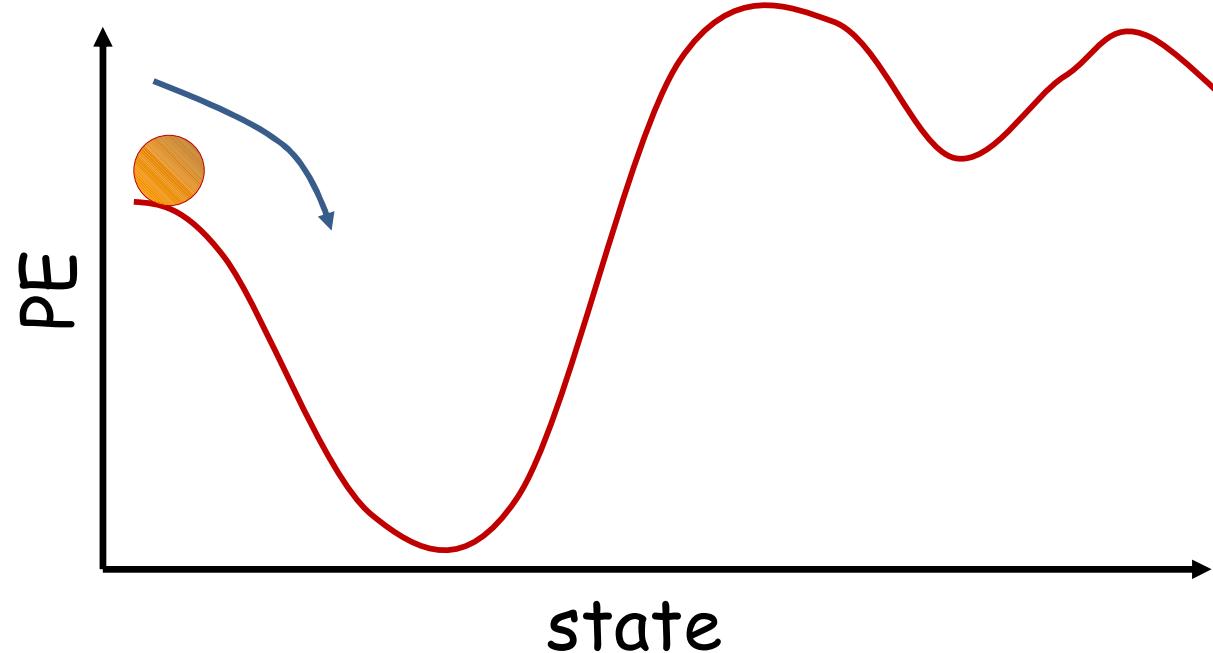
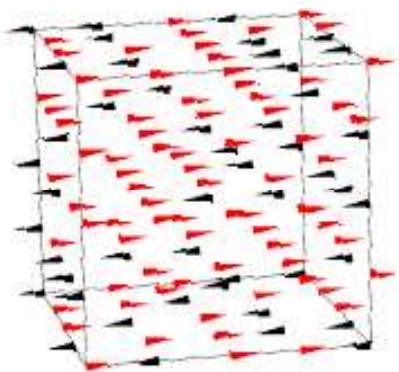
$$x_i = \begin{cases} x_i & \text{if } \text{sign}(x_i f(p_i)) = 1 \\ -x_i & \text{otherwise} \end{cases}$$

- The “Hamiltonian” (total energy) of the system

$$E = -\frac{1}{2} \sum_i x_i f(p_i) = -\sum_i \sum_{j>i} J_{ji}x_i x_j - \sum_i b_i x_i$$

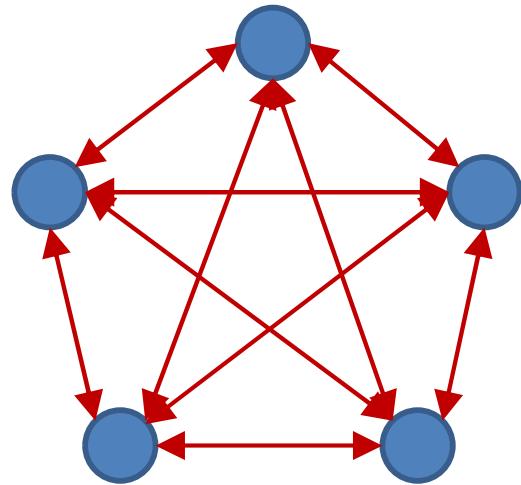
- The system *evolves* to minimize the energy
  - Dipoles stop flipping if any flips result in increase of energy

# Spin Glasses



- The system stops at one of its *stable* configurations
  - Where energy is a local minimum
- Any small jitter from this stable configuration *returns it* to the stable configuration
  - I.e. the system *remembers* its stable state and returns to it

# Hopfield Network



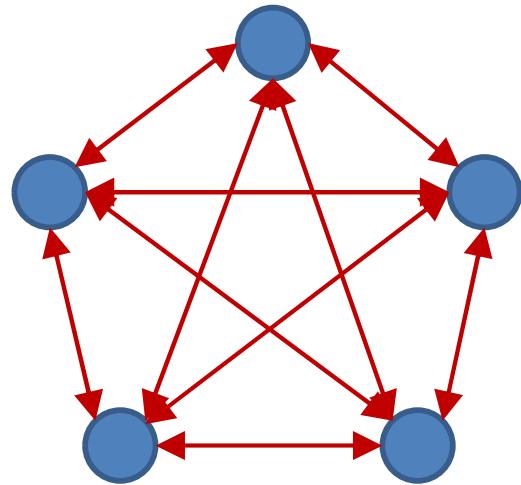
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

$$E = -\frac{1}{2} \left( \sum_{i,j \neq i} w_{ij}y_iy_j + \sum_i b_i y_i \right)$$

- This is analogous to the potential energy of a spin glass
  - The system will evolve until the energy hits a local minimum

# Hopfield Network



$$y_i = \Theta \left( \sum_{j \neq i} w_{ji} y_j + b_i \right)$$

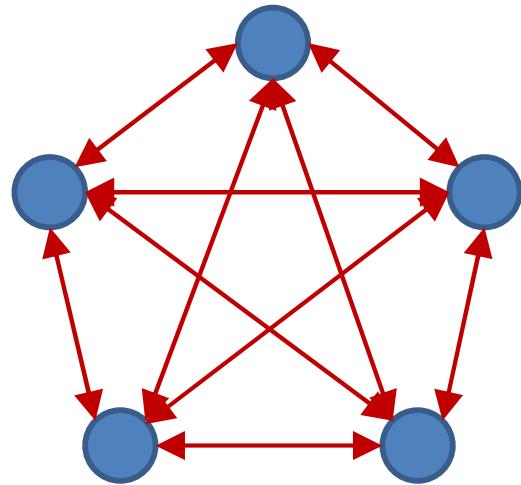
$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

The bias is equivalent to having a single extra unit pegged at 1

We will not always explicitly show the bias

Often, in fact, a bias is not used, although in our case we are just being lazy in not showing it explicitly

# Hopfield Network



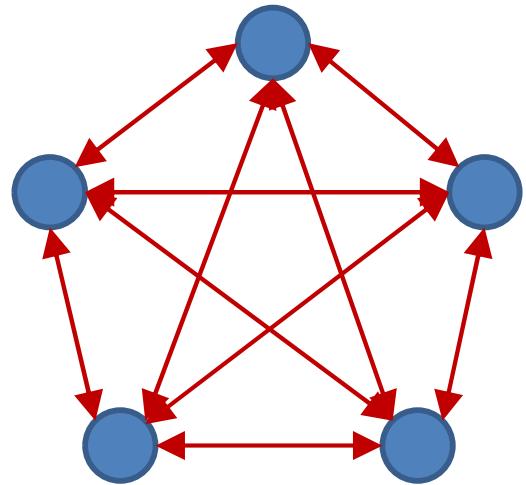
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j\right)$$

$$\Theta(z) = \begin{cases} +1 & \text{if } z > 0 \\ -1 & \text{if } z \leq 0 \end{cases}$$

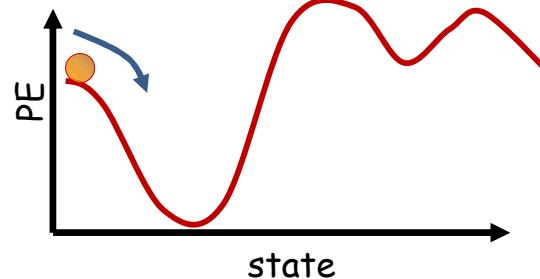
$$E = -\frac{1}{2} \sum_{i,j < i} w_{ij}y_i y_j$$

- This is analogous to the potential energy of a spin glass
  - The system will evolve until the energy hits a local minimum
    - Above equation is a factor of 0.5 off from earlier definition for conformity with thermodynamic system

# Evolution

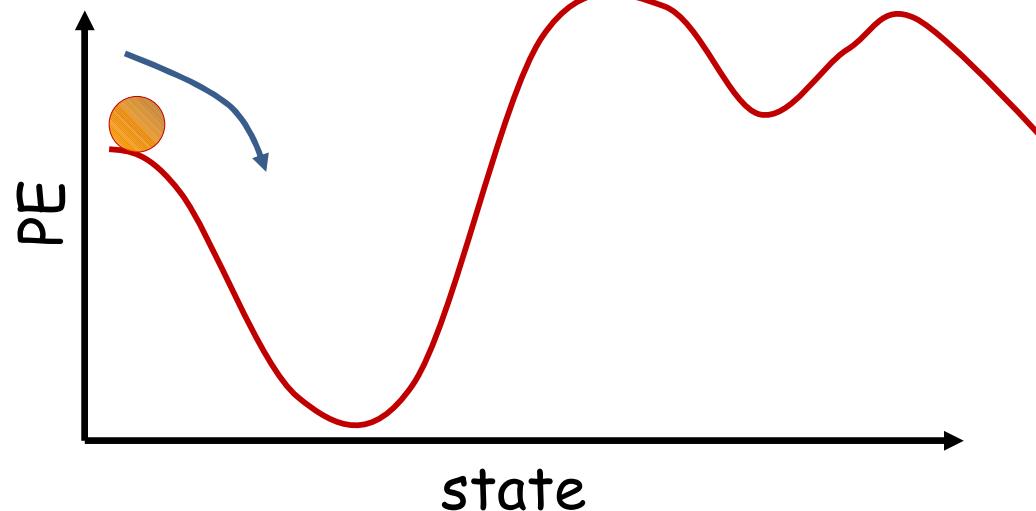
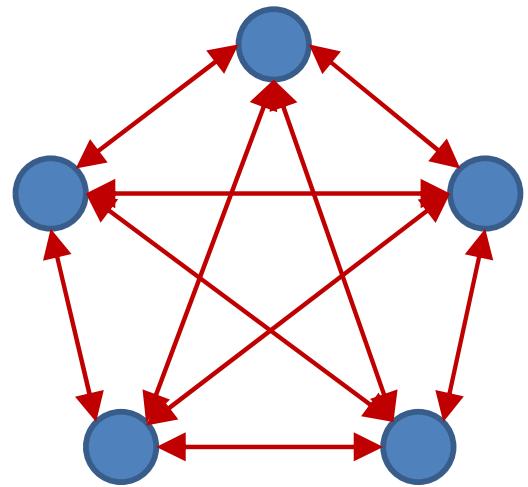


$$E = -\frac{1}{2} \sum_{i,j < i} w_{ij} y_i y_j$$



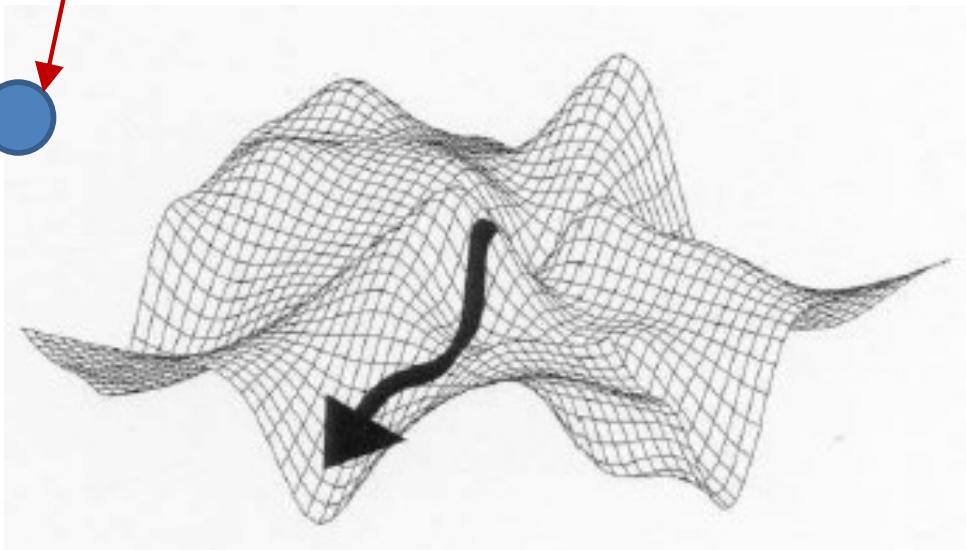
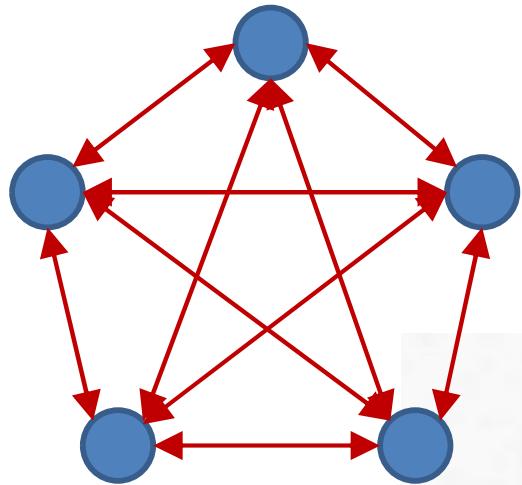
- The network will evolve until it arrives at a local minimum in the energy contour

# **Content-addressable memory**



- Each of the minima is a “stored” pattern
  - If the network is initialized close to a stored pattern, it will inevitably evolve to the pattern
- **This is a *content addressable memory***
  - Recall memory content from partial or corrupt values
- Also called ***associative memory***

# Evolution

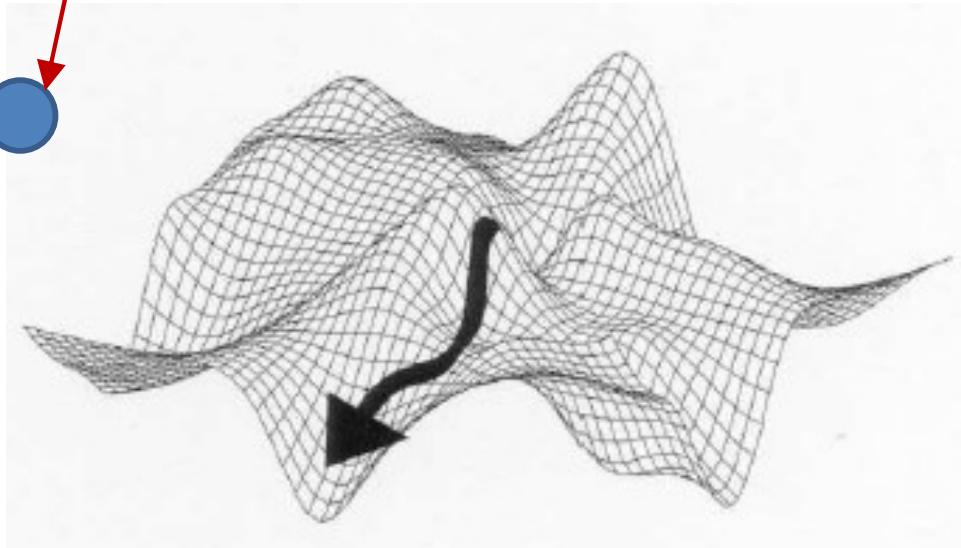
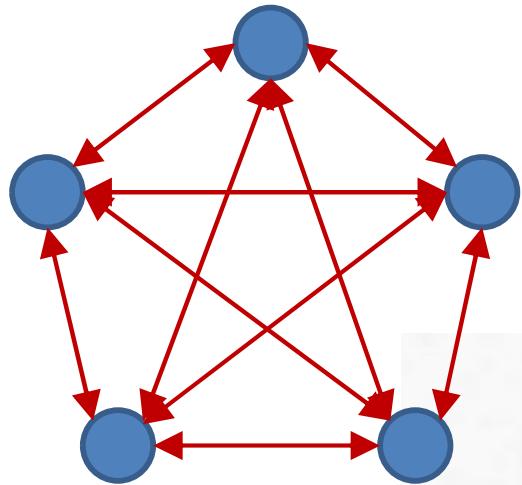


$$E = -\frac{1}{2} \sum_{i,j < i} w_{ij} y_i y_j$$

Image pilfered from  
unknown source

- The network will evolve until it arrives at a local minimum in the energy contour

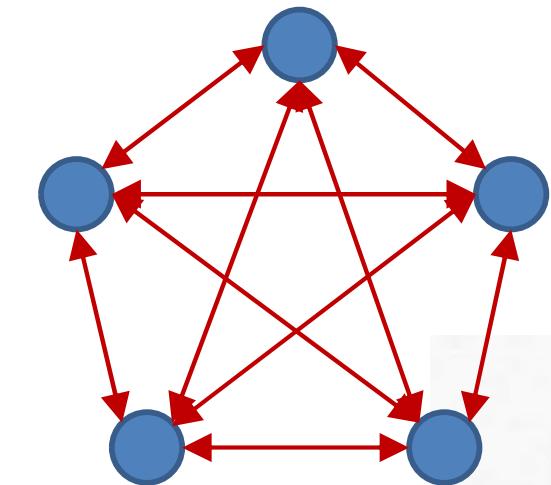
# Evolution



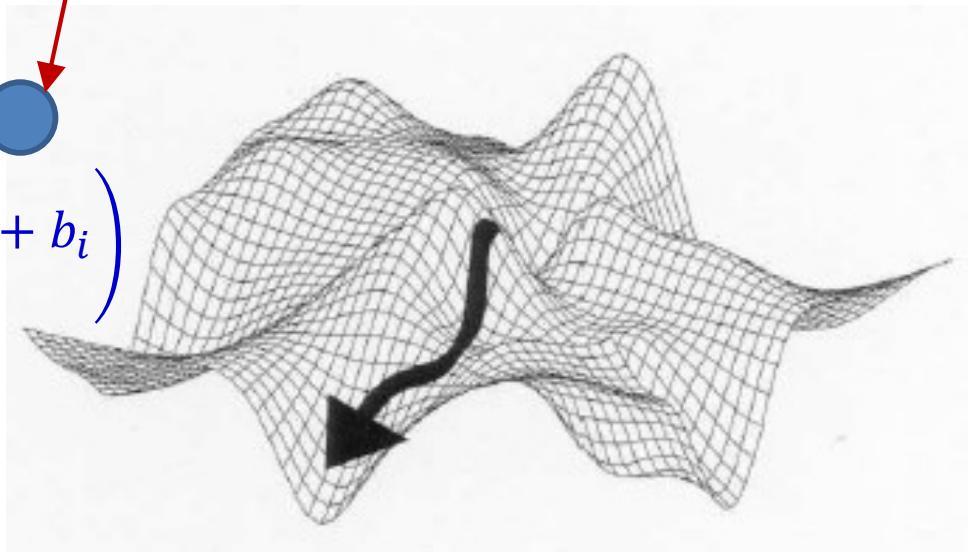
$$E = -\frac{1}{2} \sum_{i,j < i} w_{ij} y_i y_j$$

- The network will evolve until it arrives at a local minimum in the energy contour
- We proved that *every* change in the network will result in *decrease* in energy
  - So path to energy minimum is monotonic

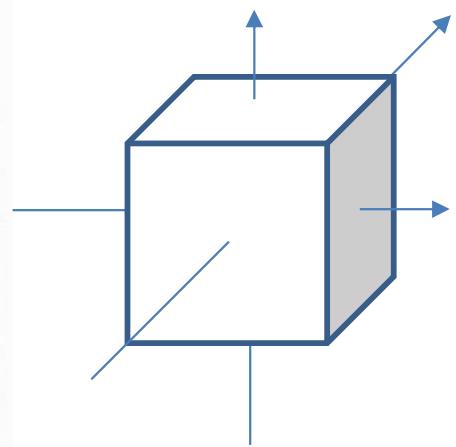
# Evolution



$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$

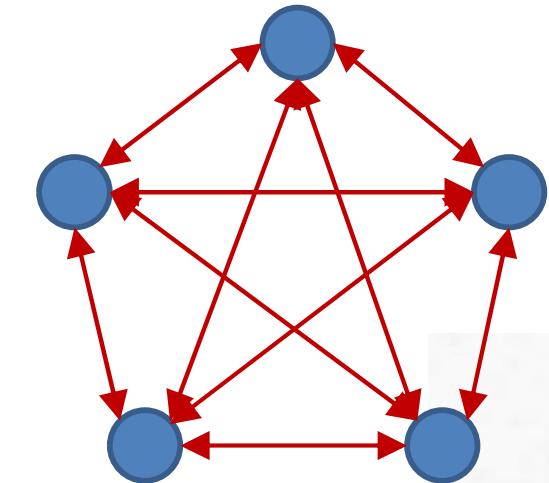


$$E = -\frac{1}{2} \sum_{i,j < i} w_{ij} y_i y_j$$

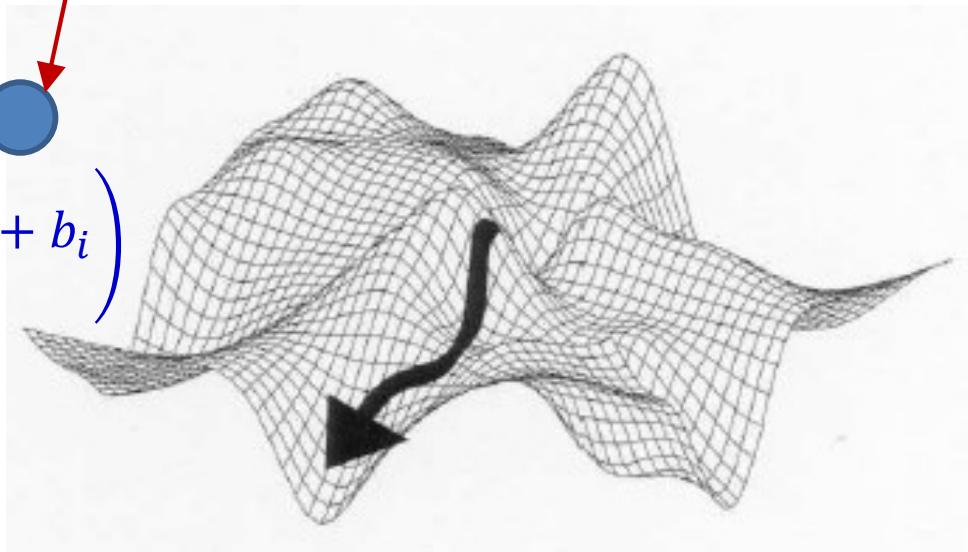


- For threshold activations the energy contour is only defined on a lattice
  - Corners of a unit cube on  $[-1,1]^N$

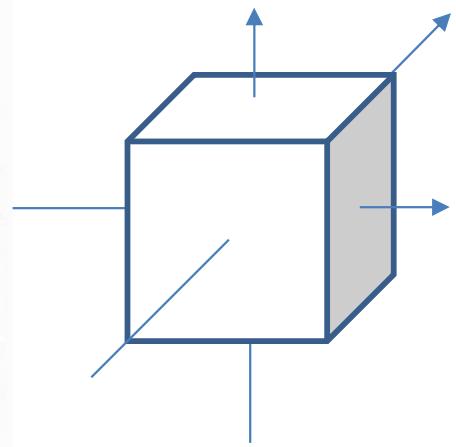
# Evolution



$$y_i = \Theta\left(\sum_{j \neq i} w_{ji} y_j + b_i\right)$$

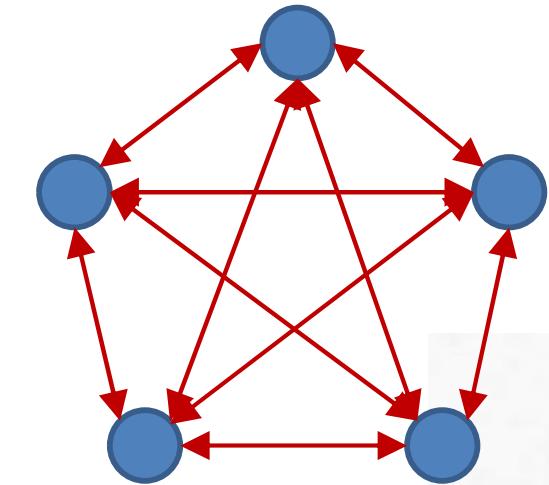


$$E = -\frac{1}{2} \sum_{i,j < i} w_{ij} y_i y_j$$

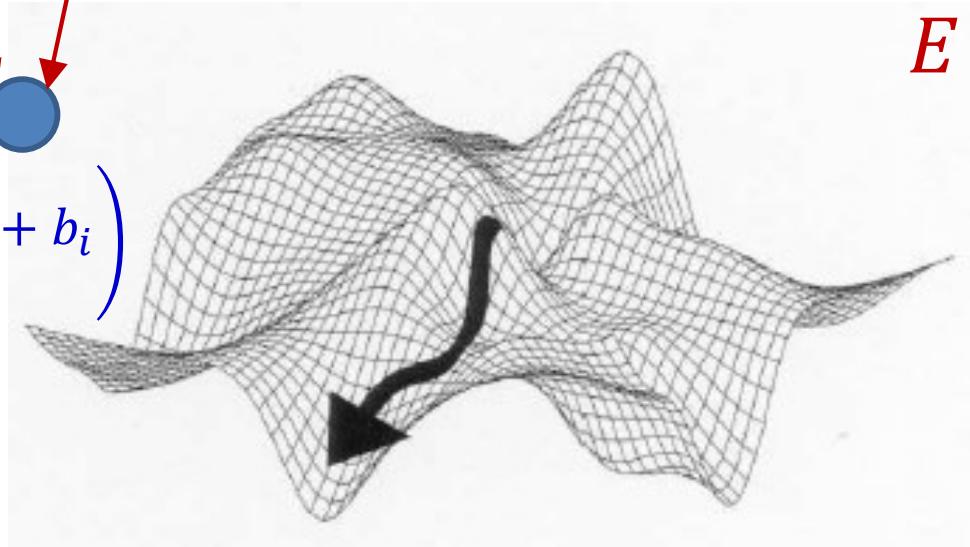


- For threshold activations the energy contour is only defined on a lattice
  - Corners of a unit cube on  $[-1,1]^N$
- For tanh activations it will be a continuous function

# Evolution



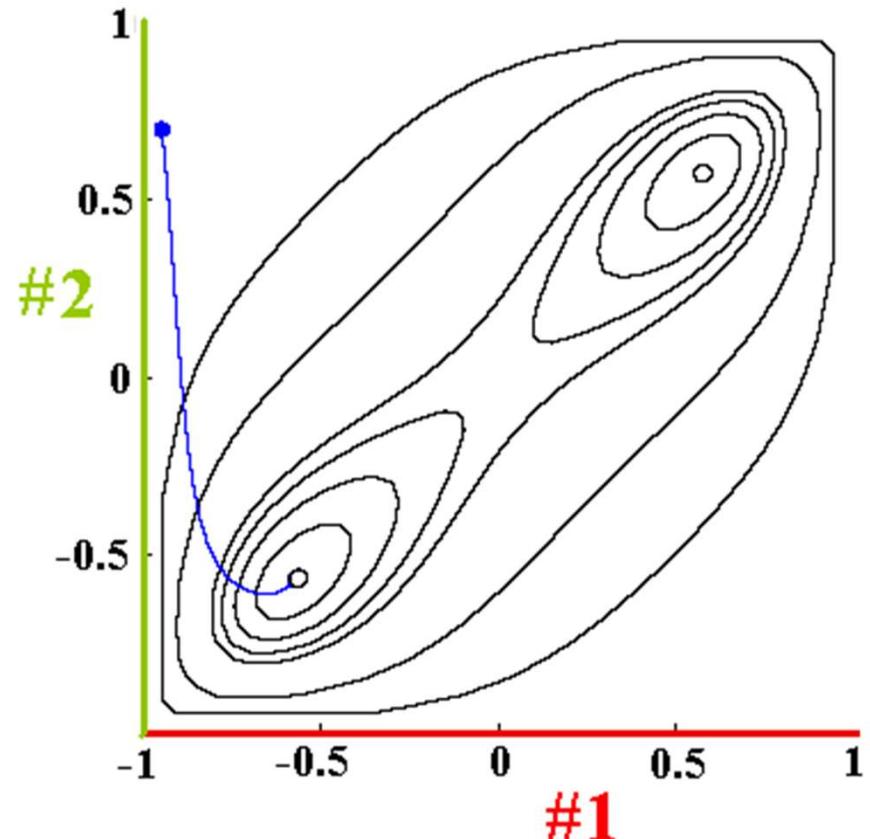
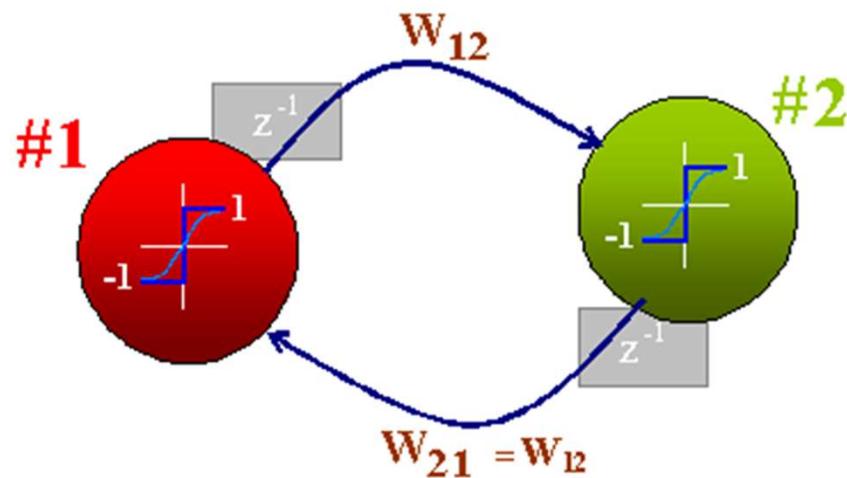
$$y_i = \Theta\left(\sum_{j \neq i} w_{ji}y_j + b_i\right)$$



$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$

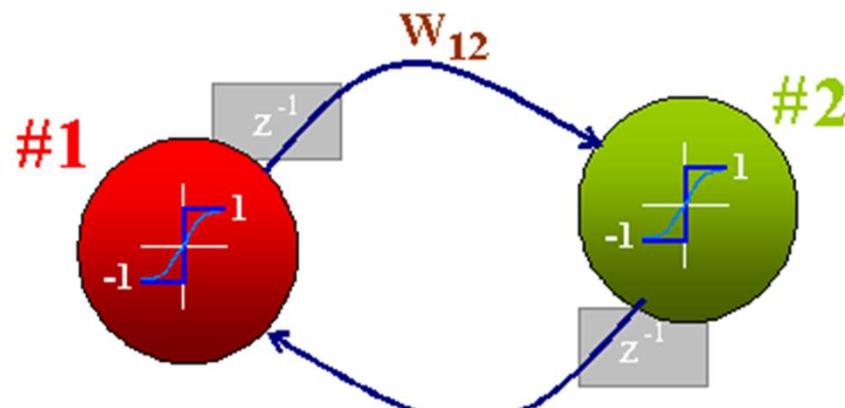
- For threshold activations the energy contour is only defined on a lattice
  - Corners of a unit cube
- For tanh activations it will be a continuous function
  - With output in [-1 1]

# “Energy”contour for a 2-neuron net



- Two stable states (tanh activation)
  - Symmetric, not at corners
  - Blue arc shows a typical trajectory for tanh activation

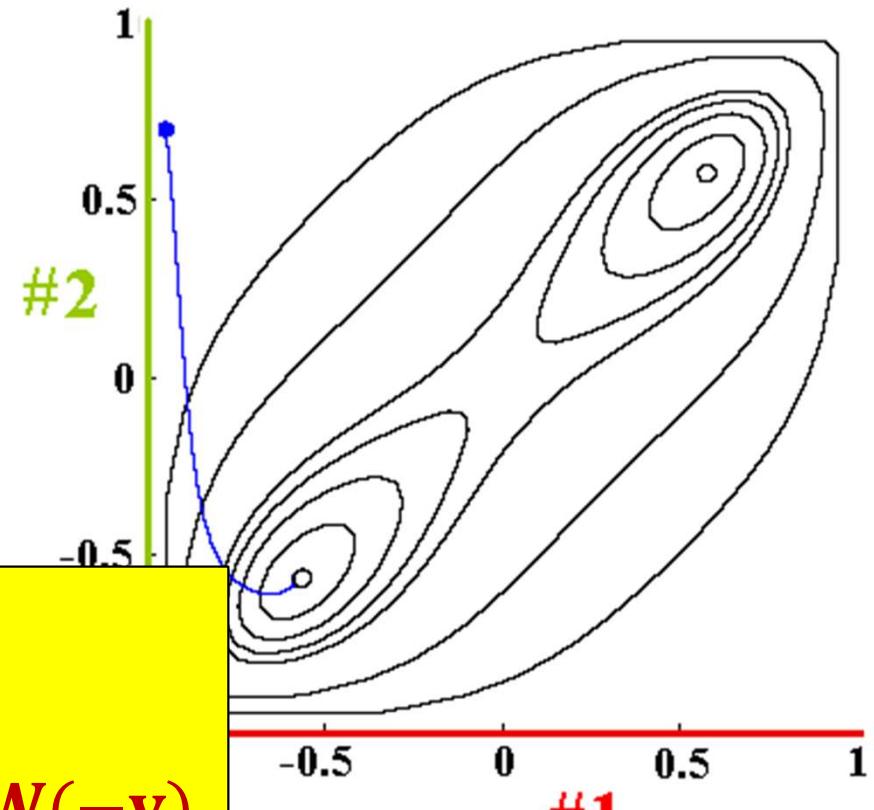
# “Energy”contour for a 2-neuron net



Why symmetric?

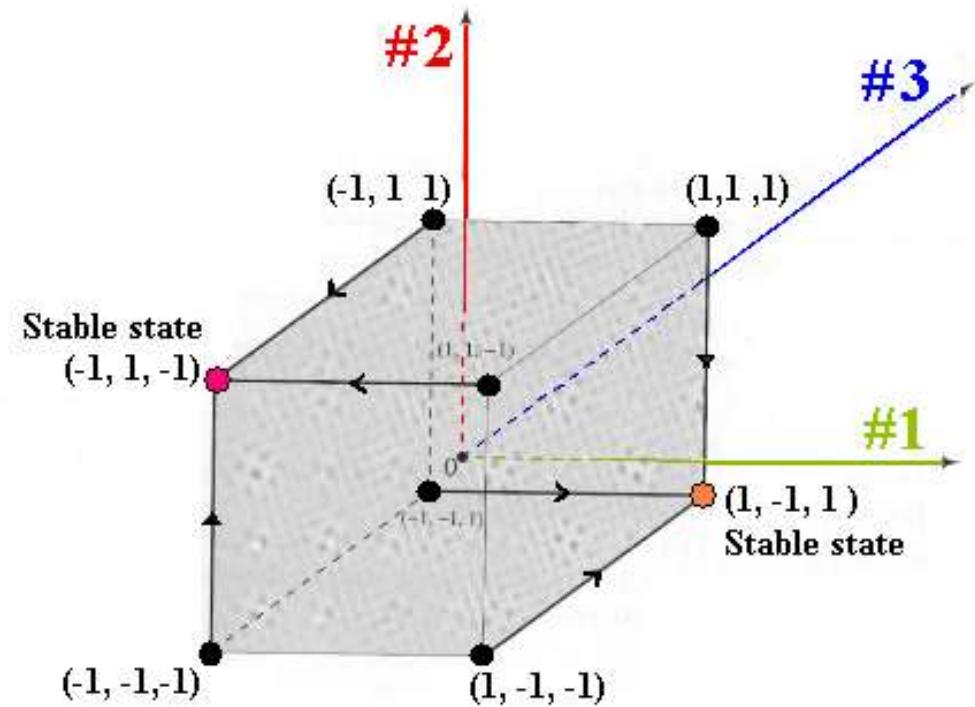
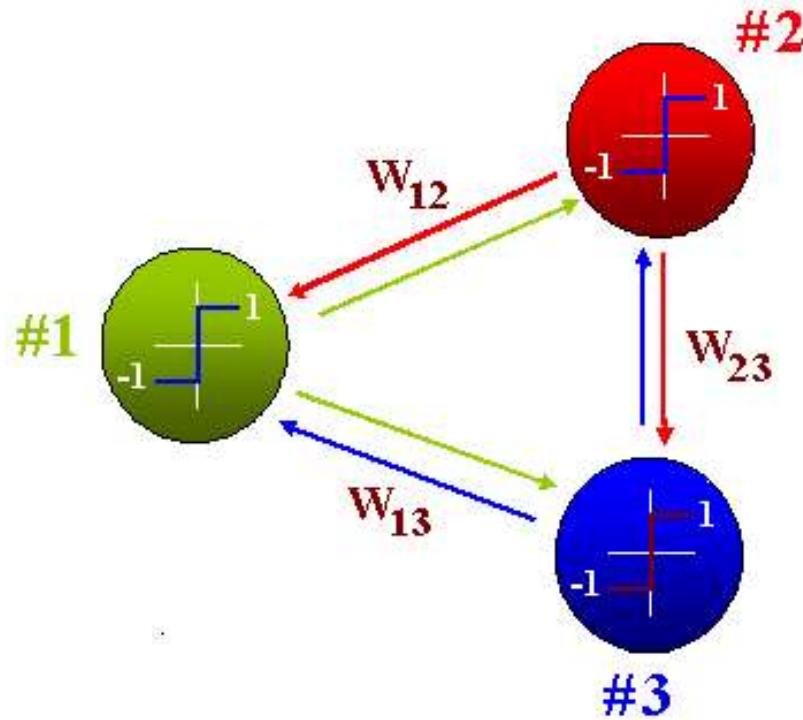
Because  $-\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y} = -\frac{1}{2} (-\mathbf{y})^T \mathbf{W} (-\mathbf{y})$

If  $\hat{\mathbf{y}}$  is a local minimum, so is  $-\hat{\mathbf{y}}$



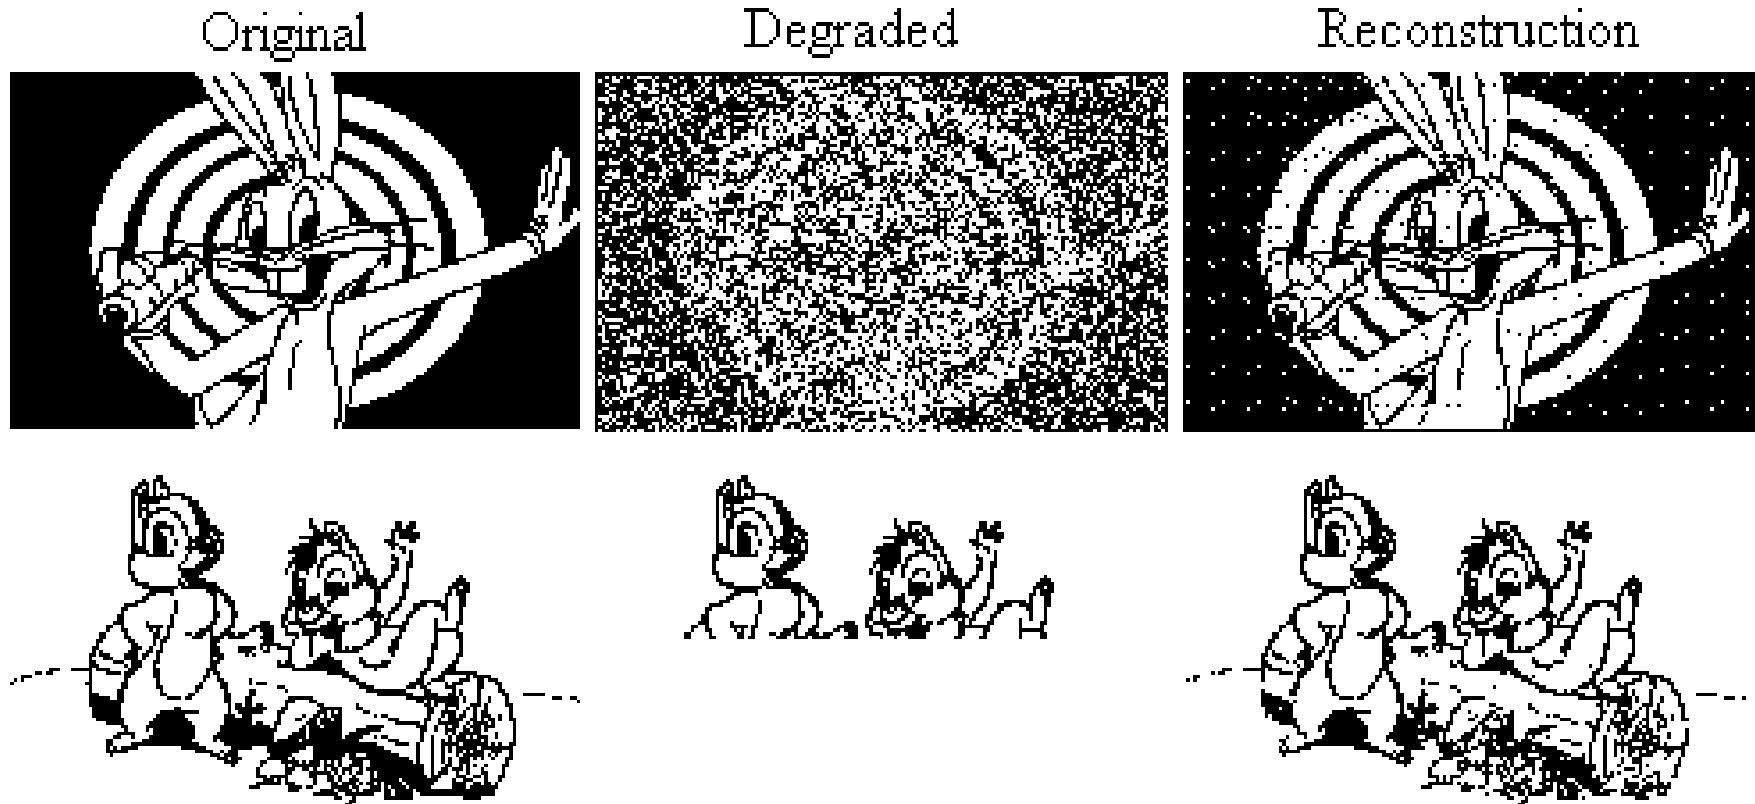
- Blue arc shows a typical trajectory for sigmoid activation

# 3-neuron net



- 8 possible states
- 2 stable states (hard thresholded network)

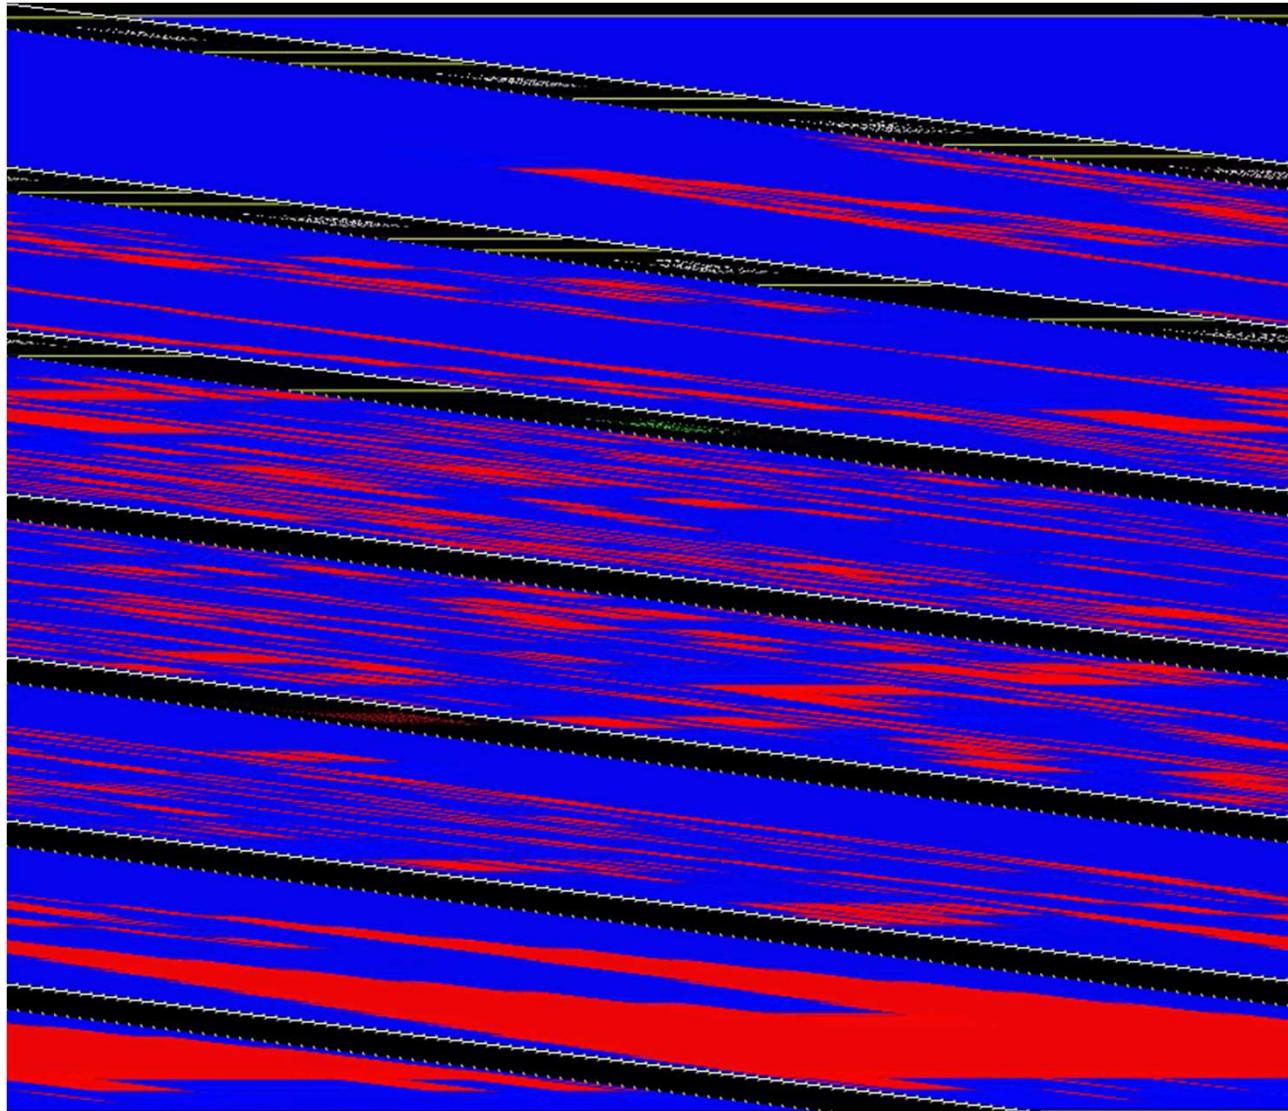
# Examples: Content addressable memory



Hopfield network reconstructing degraded images  
from noisy (top) or partial (bottom) cues.

- <http://staff.itee.uq.edu.au/janetw/cmc/chapters/Hopfield/> 47

# Hopfield net examples



# Computational algorithm

1. Initialize network with initial pattern

$$y_i(0) = x_i, \quad 0 \leq i \leq N - 1$$

2. Iterate until convergence

$$y_i(t + 1) = \Theta\left(\sum_{j \neq i} w_{ji} y_j\right), \quad 0 \leq i \leq N - 1$$

- Very simple
- Updates can be done sequentially, or all at once
- Convergence

$$E = - \sum_i \sum_{j > i} w_{ji} y_j y_i$$

does not change significantly any more

# Computational algorithm

1. Initialize network with initial pattern

$$\mathbf{y} = \mathbf{x}, \quad 0 \leq i \leq N - 1$$

2. Iterate until convergence

$$\mathbf{y} = \Theta(\mathbf{W}\mathbf{y})$$

Writing  $\mathbf{y} = [y_1, y_2, y_3, \dots, y_N]^\top$   
and arranging the weights as a matrix  $\mathbf{W}$

- Very simple
- Updates can be done sequentially, or all at once
- Convergence

$$E = -0.5\mathbf{y}^\top \mathbf{W}\mathbf{y}$$

does not change significantly any more

# Story so far

- A Hopfield network is a loopy binary network with symmetric connections
  - Neurons try to align themselves to the local field caused by other neurons
- Given an initial configuration, the patterns of neurons in the net will evolve until the “energy” of the network achieves a local minimum
  - The evolution will be monotonic in total energy
  - The dynamics of a Hopfield network mimic those of a spin glass
  - The network is symmetric: if a pattern  $Y$  is a local minimum, so is  $-Y$
- The network acts as a *content-addressable* memory
  - If you initialize the network with a somewhat damaged version of a local-minimum pattern, it will evolve into that pattern
  - Effectively “recalling” the correct pattern, from a damaged/incomplete version

# Poll 2, @1797

Mark all that are correct about Hopfield nets

- The network activations evolve until the energy of the net arrives at a local minimum
- Hopfield networks are a form of content addressable memory
- It is possible to analytically determine the stored memories by inspecting the weights matrix

# Poll 2

**Mark all that are correct about Hopfield nets**

- **The network activations evolve until the energy of the net arrives at a local minimum**
- **Hopfield networks are a form of content addressable memory**
- It is possible to analytically determine the stored memories by inspecting the weights matrix

# Issues

- How do we make the network store *a specific* pattern or set of patterns?
- How many patterns can we store?
- How to “retrieve” patterns better..

# Issues

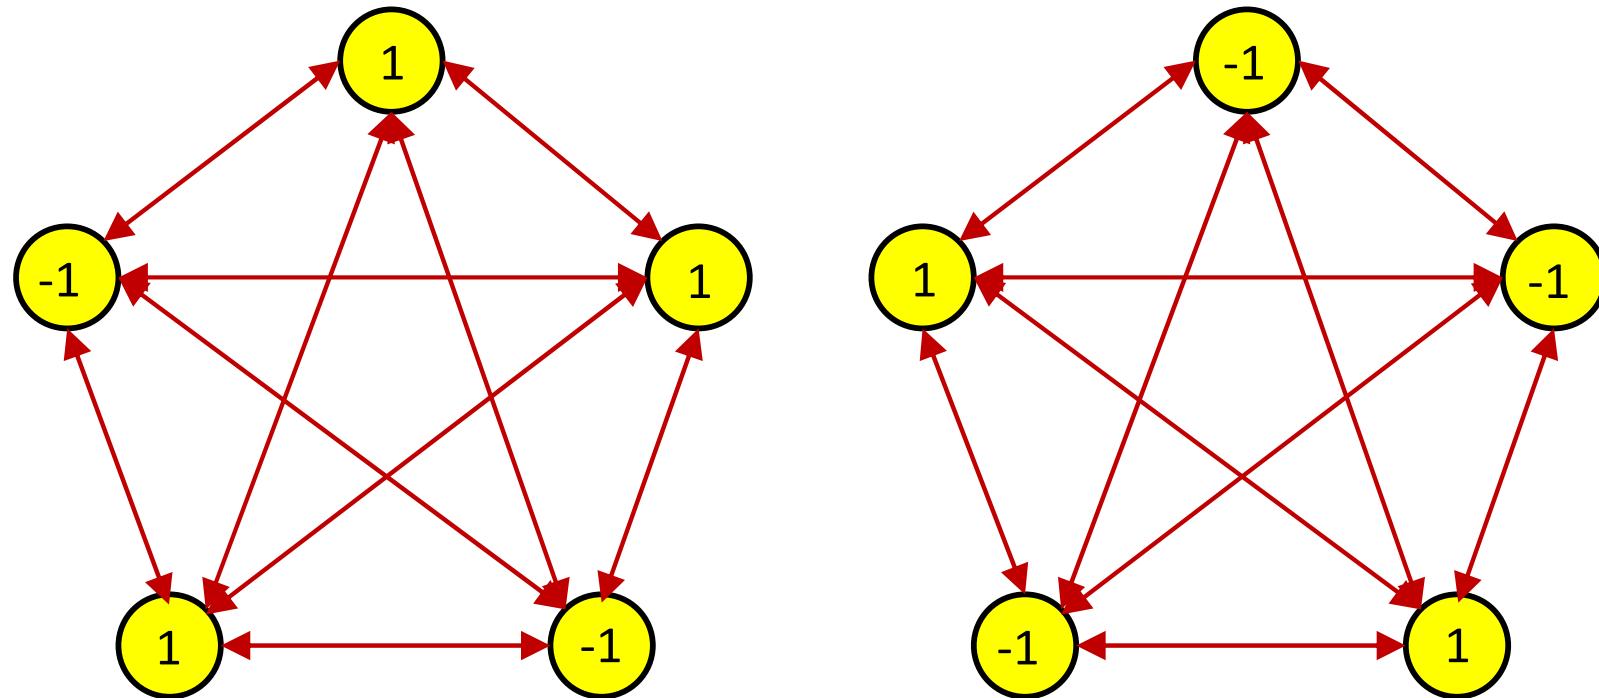
- How do we make the network store *a specific* pattern or set of patterns?
- How many patterns can we store?
- How to “retrieve” patterns better..

# How do we remember a *specific* pattern?

- How do we teach a network to “remember” this image
- For an image with  $N$  pixels we need a network with  $N$  neurons
- Every neuron connects to every other neuron
- Weights are symmetric (not mandatory)
- $\frac{N(N-1)}{2}$  weights in all



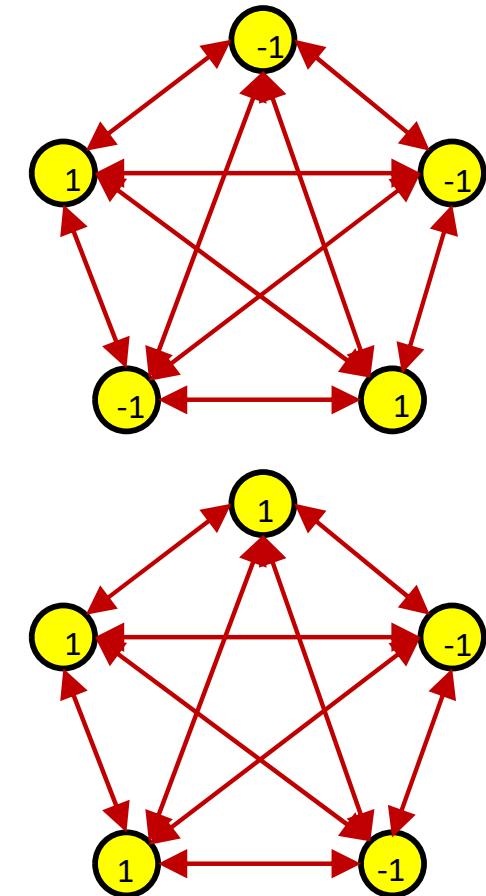
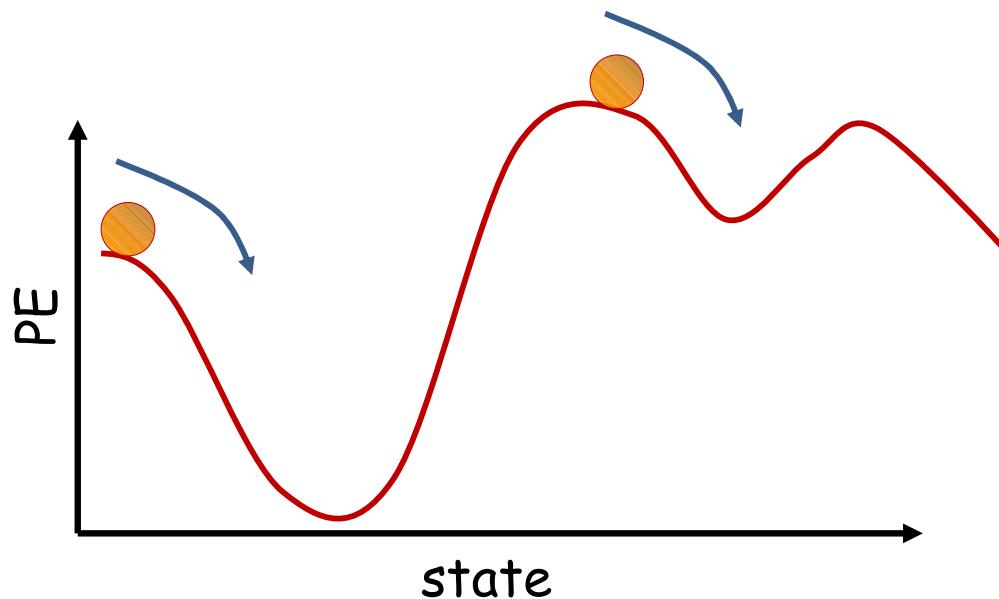
# Storing patterns: Training a network



- A network that stores pattern  $P$  also naturally stores  $-P$ 
  - Symmetry  $E(P) = E(-P)$  since  $E$  is a function of  $y_i y_j$

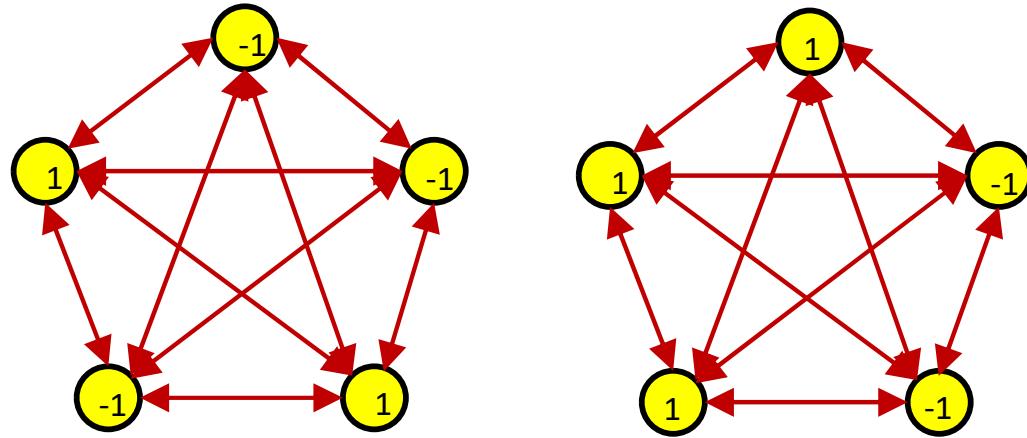
$$E = - \sum_i \sum_{j < i} w_{ji} y_j y_i$$

# A network can store *multiple* patterns



- Every stable point is a stored pattern
- So we could design the net to store multiple patterns
  - Remember that every stored pattern  $P$  is actually *two* stored patterns,  $P$  and  $-P$

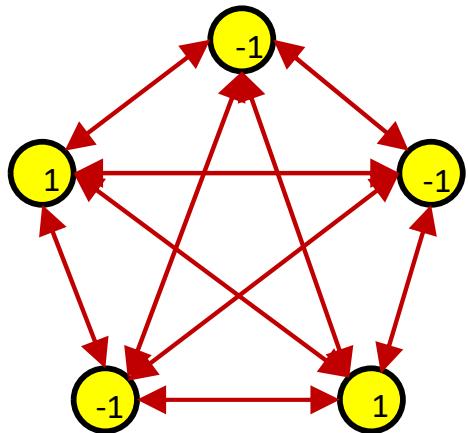
# Storing a pattern



$$E = - \sum_i \sum_{j < i} w_{ji} y_j y_i$$

- Design  $\{w_{ij}\}$  such that the energy is a local minimum at the desired  $P = \{y_i\}$

# Storing specific patterns

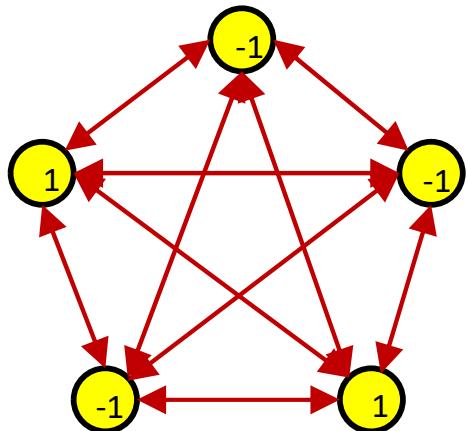


- Storing 1 pattern: We want

$$\text{sign} \left( \sum_{j \neq i} w_{ji} y_j \right) = y_i \quad \forall i$$

- This is a stationary pattern

# Storing specific patterns



HEBBIAN LEARNING:

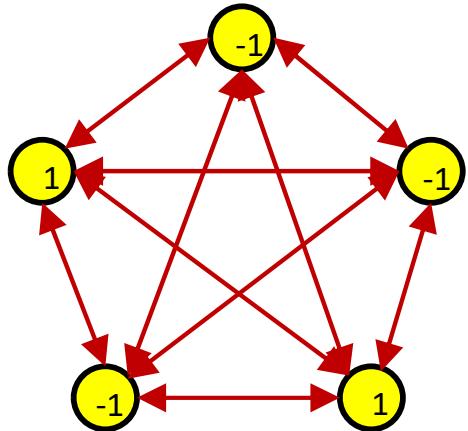
$$w_{ji} = y_j y_i$$

- Storing 1 pattern: We want

$$\text{sign} \left( \sum_{j \neq i} w_{ji} y_j \right) = y_i \quad \forall i$$

- This is a stationary pattern

# Storing specific patterns

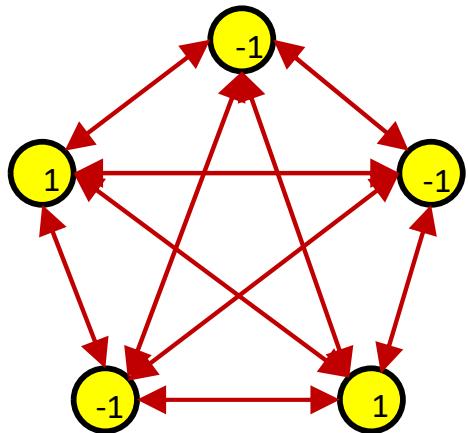


HEBBIAN LEARNING:

$$w_{ji} = y_j y_i$$

- $\text{sign}\left(\sum_{j \neq i} w_{ji} y_j\right) = \text{sign}\left(\sum_{j \neq i} y_j y_i y_j\right)$   
 $= \text{sign}\left(\sum_{j \neq i} y_j^2 y_i\right) = \text{sign}(y_i) = y_i$

# Storing specific patterns



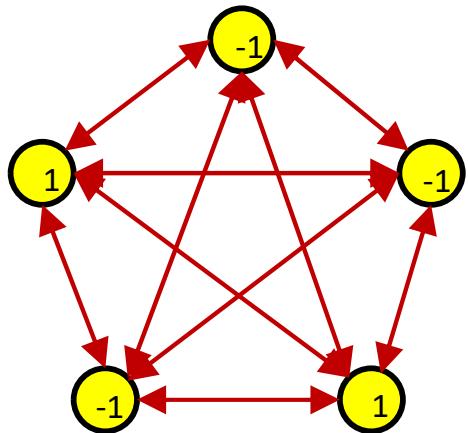
HEBBIAN LEARNING:

$$w_{ji} = y_j y_i$$

The pattern is stationary

- $\text{sign}\left(\sum_{j \neq i} w_{ji} y_j\right) = \text{sign}\left(\sum_{j \neq i} y_j y_i y_j\right)$   
 $= \text{sign}\left(\sum_{j \neq i} y_j^2 y_i\right) = \text{sign}(y_i) = y_i$

# Storing specific patterns



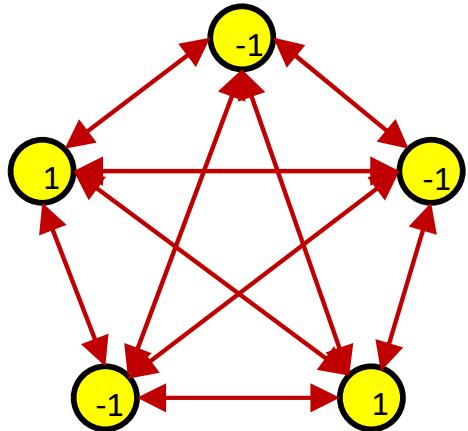
HEBBIAN LEARNING:

$$w_{ji} = y_j y_i$$

$$\begin{aligned} E &= - \sum_i \sum_{j < i} w_{ji} y_j y_i = - \sum_i \sum_{j < i} y_i^2 y_j^2 \\ &= - \sum_i \sum_{j < i} 1 = -0.5N(N-1) \end{aligned}$$

- This is the lowest possible energy value for the network for binary weights

# Storing specific patterns



HEBBIAN LEARNING:

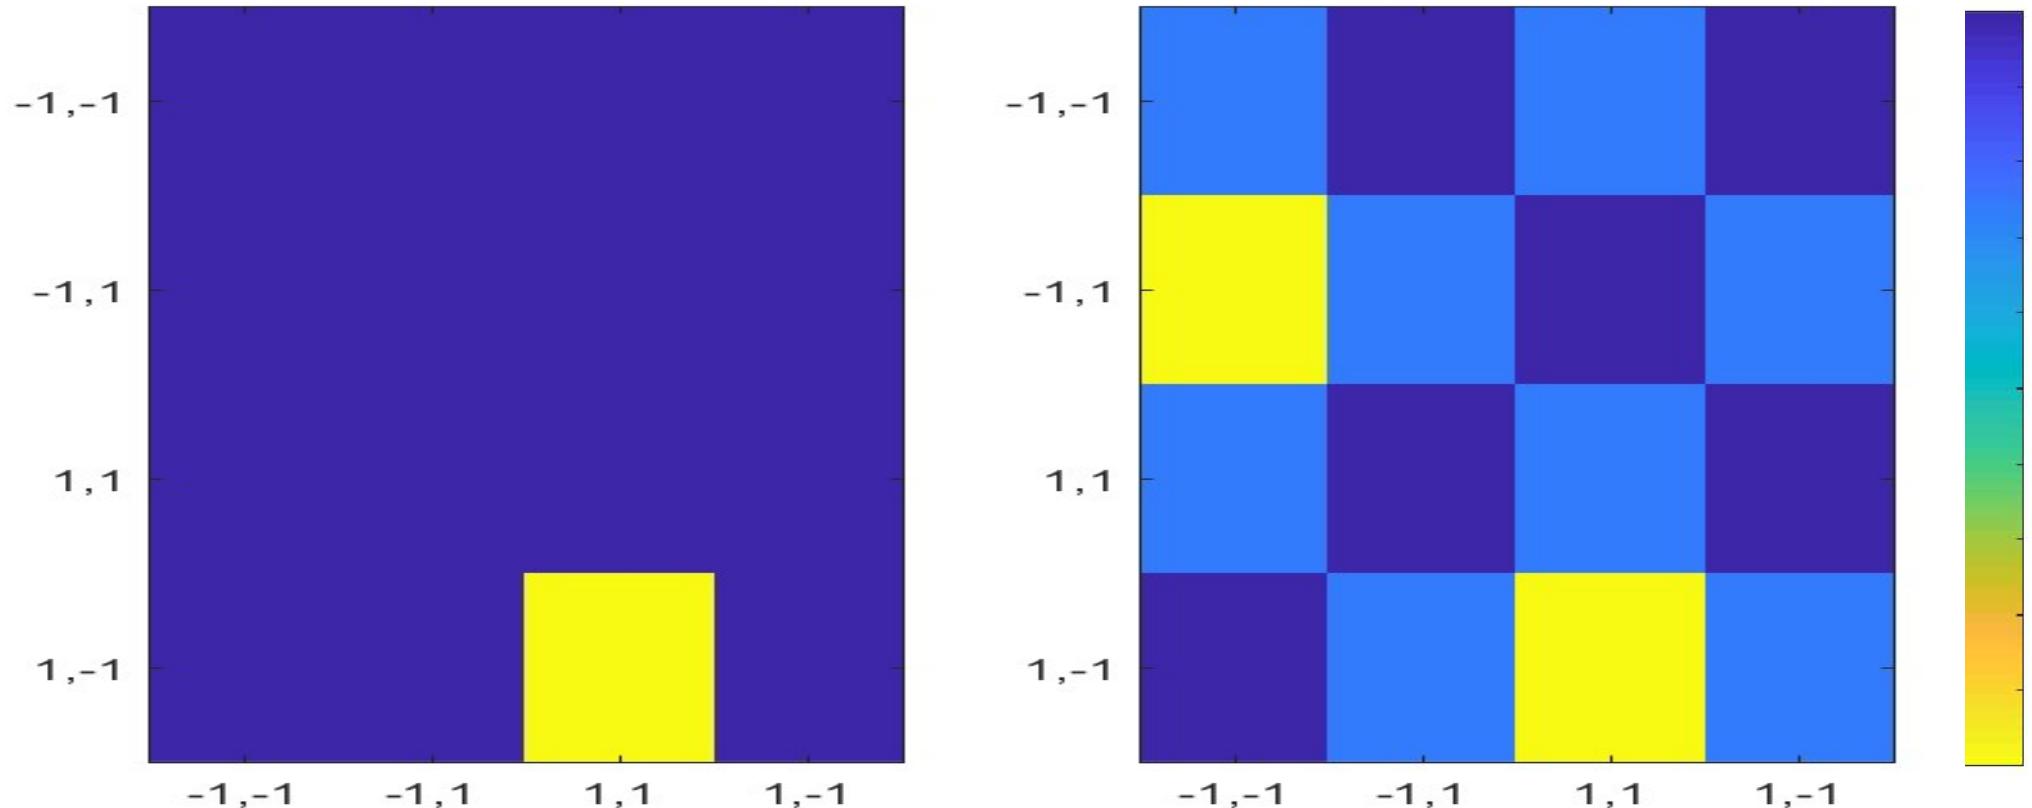
$$w_{ji} = y_j y_i$$

The pattern is *STABLE*

$$\begin{aligned} E &= - \sum_i \sum_{j < i} w_{ji} y_j y_i = - \sum_i \sum_{j < i} y_i^- y_j^- \\ &= - \sum_i \sum_{j < i} 1 = -0.5N(N - 1) \end{aligned}$$

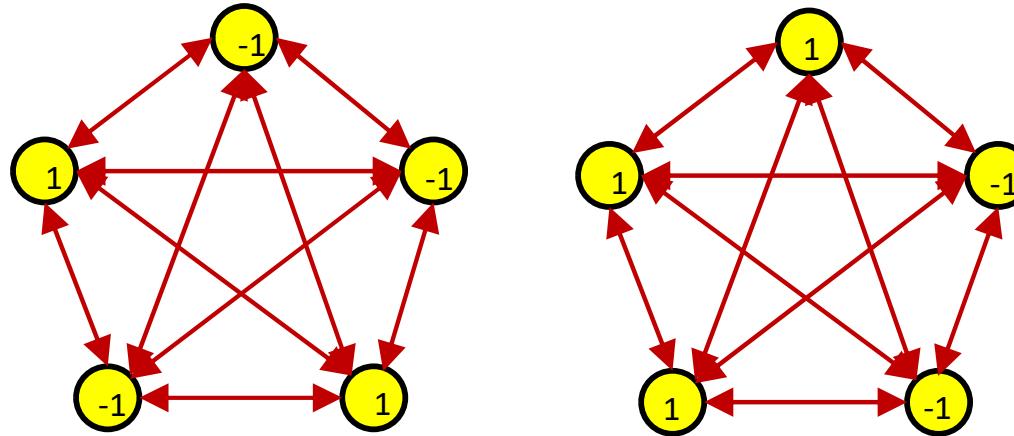
- This is the lowest possible energy value for the network for binary weights

# Hebbian learning: Storing a 4-bit pattern



- Left: Pattern stored. Right: Energy map
- Stored pattern has lowest energy
- Gradation of energy ensures stored pattern (or its ghost) is recalled from everywhere

# Storing multiple patterns

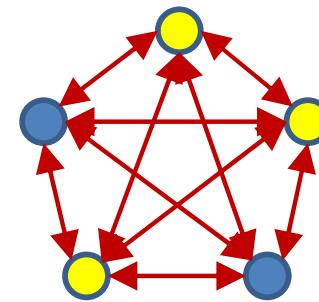
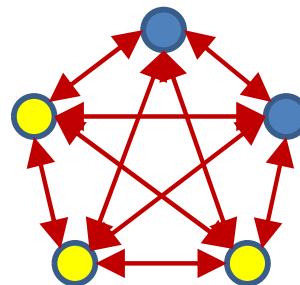
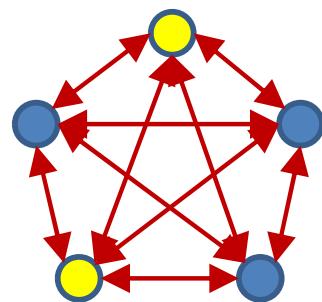


- To store *more* than one pattern

$$w_{ji} = \frac{1}{N} \sum_{\mathbf{y}_p \in \{\mathbf{y}_p\}} y_i^p y_j^p$$

- $\{\mathbf{y}_p\}$  is the set of patterns to store
- Super/subscript  $p$  represents the specific pattern
- $N$  is the number of patterns

# How many patterns can we store?



- **Hopfield**: For a network of  $N$  neurons can store up to  $\sim 0.15N$  random patterns through Hebbian learning
  - Provided they are “far” enough
- Where did this number come from?

# The limits of Hebbian Learning

- Consider the following: We must store  $K$   $N$ -bit patterns of the form

$$\mathbf{y}_k = [y_1^k, y_2^k, \dots, y_N^k], k = 1 \dots K$$

- Hebbian learning (scaling by  $\frac{1}{N}$  for normalization, this does not affect actual pattern storage):

$$w_{ij} = \frac{1}{N} \sum_k y_i^k y_j^k$$

- For any pattern  $\mathbf{y}_p$  to be stable:**

$$y_i^p \sum_j w_{ij} y_j^p > 0 \quad \forall i$$

$$y_i^p \frac{1}{N} \sum_j \sum_k y_i^k y_j^k y_j^p > 0 \quad \forall i$$

# The limits of Hebbian Learning

- For any pattern  $\mathbf{y}_p$  to be stable:

$$y_i^p \frac{1}{N} \sum_j \sum_k y_i^k y_j^k y_j^p > 0 \quad \forall i$$

$$y_i^p \frac{1}{N} \sum_j y_i^p y_j^p y_j^p + y_i^p \frac{1}{N} \sum_j \sum_{k \neq p} y_i^k y_j^k y_j^p > 0 \quad \forall i$$

- Note that the first term equals 1 (because  $y_j^p y_j^p = y_i^p y_i^p = 1$ )
  - i.e. for  $\mathbf{y}_p$  to be stable the requirement is that the second *crosstalk term*:

$$y_i^p \frac{1}{N} \sum_j \sum_{k \neq p} y_i^k y_j^k y_j^p > -1 \quad \forall i$$

- The pattern will *fail* to be stored if the *crosstalk*

$$y_i^p \frac{1}{N} \sum_j \sum_{k \neq p} y_i^k y_j^k y_j^p < -1 \quad \text{for any } i$$

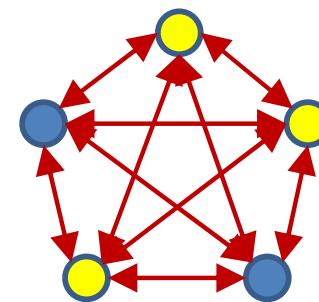
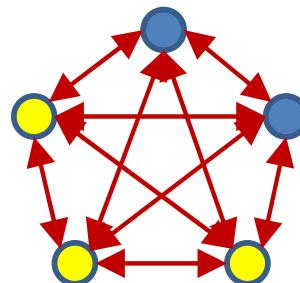
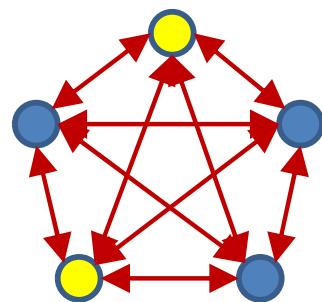
# The limits of Hebbian Learning

- For any random set of  $K$  patterns to be stored, the probability of the following must be low

$$\left( C_i^p = \frac{1}{N} \sum_j \sum_{k \neq p} y_i^p y_i^k y_j^k y_j^p \right) < -1$$

- For large  $N$  and  $K$  the probability distribution of  $C_i^p$  approaches a Gaussian with 0 mean, and variance  $K/N$ 
  - Considering that individual bits  $y_i^l \in \{-1, +1\}$  and have variance 1
- For a Gaussian,  $C \sim N(0, K/N)$ 
  - $P(C < -1 | \mu = 0, \sigma^2 = K/N) < 0.004$  for  $K/N < 0.14$
- I.e. To have less than 0.4% probability that stored patterns will *not* be stable,  $K < 0.14N$

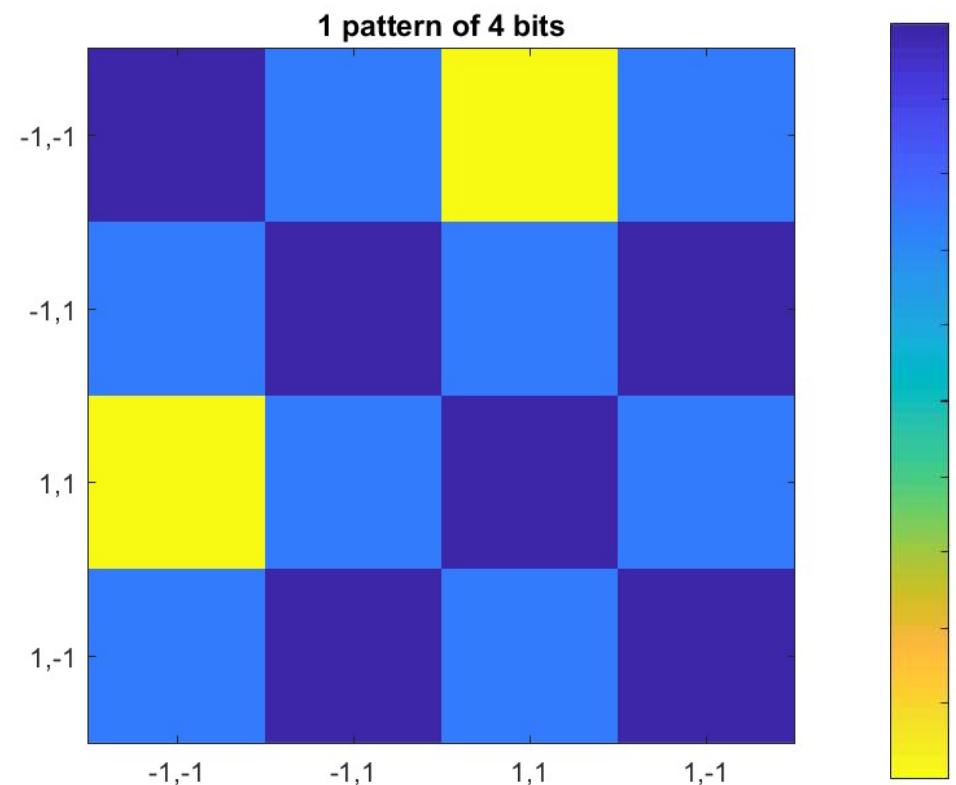
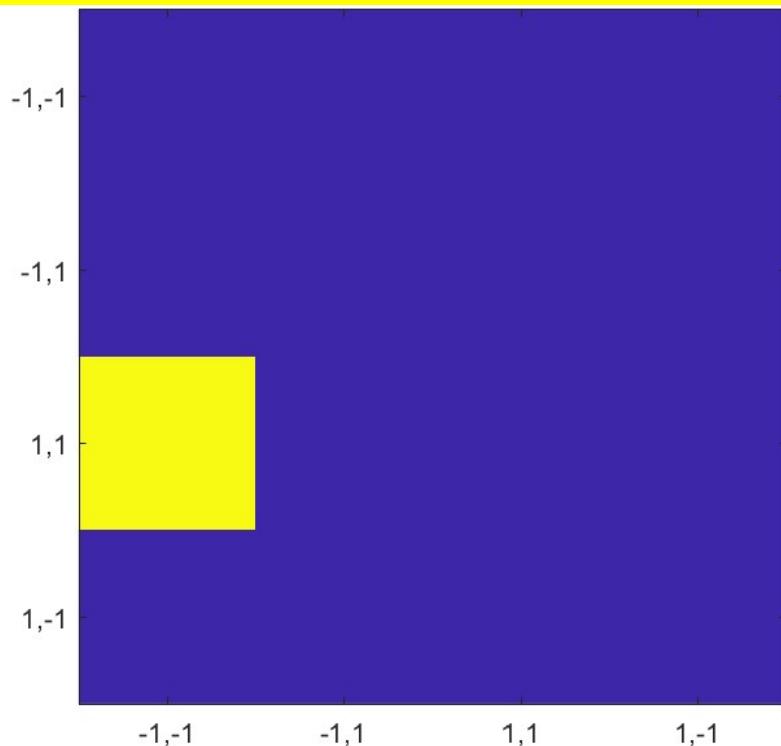
# How many patterns can we store?



- A network of  $N$  neurons trained by Hebbian learning can store up to  $\sim 0.14N$  random patterns with low probability of error
  - Computed assuming  $\text{prob}(\text{bit} = 1) = 0.5$ 
    - On average no. of matched bits in any pair = no. of mismatched bits
      - Patterns are “orthogonal” – maximally distant – from one another
    - Expected behavior for *non-orthogonal* patterns?
- To get some insight into what is stored, lets see some examples

# Hebbian learning: One 4-bit pattern

Topological representation on a Karnaugh map

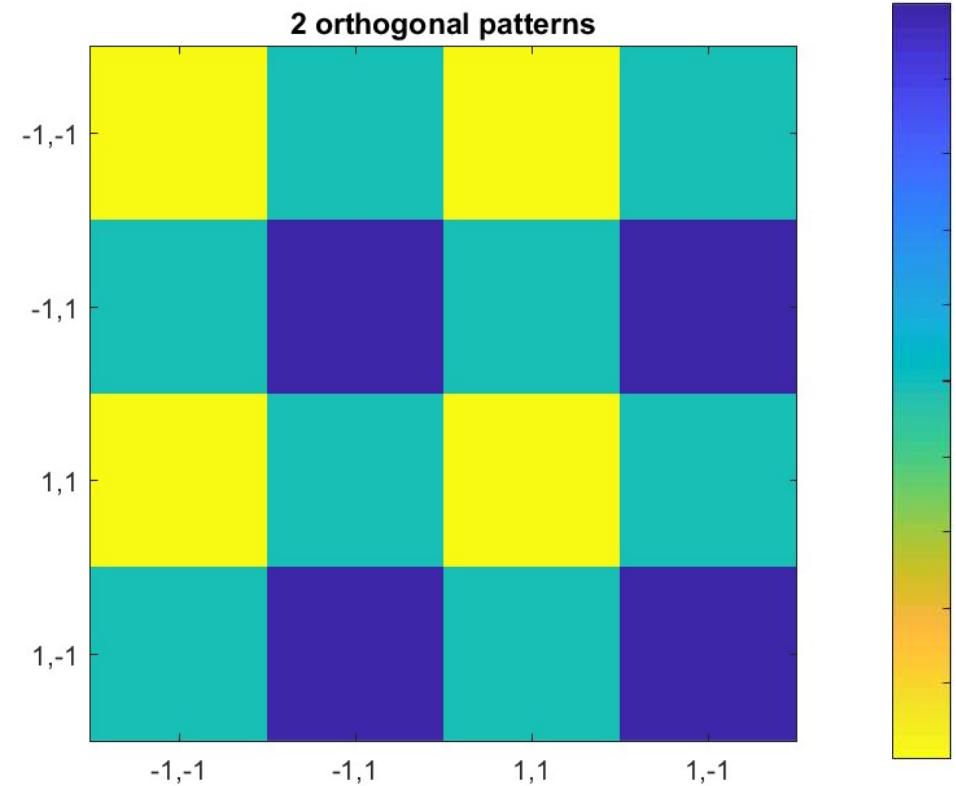
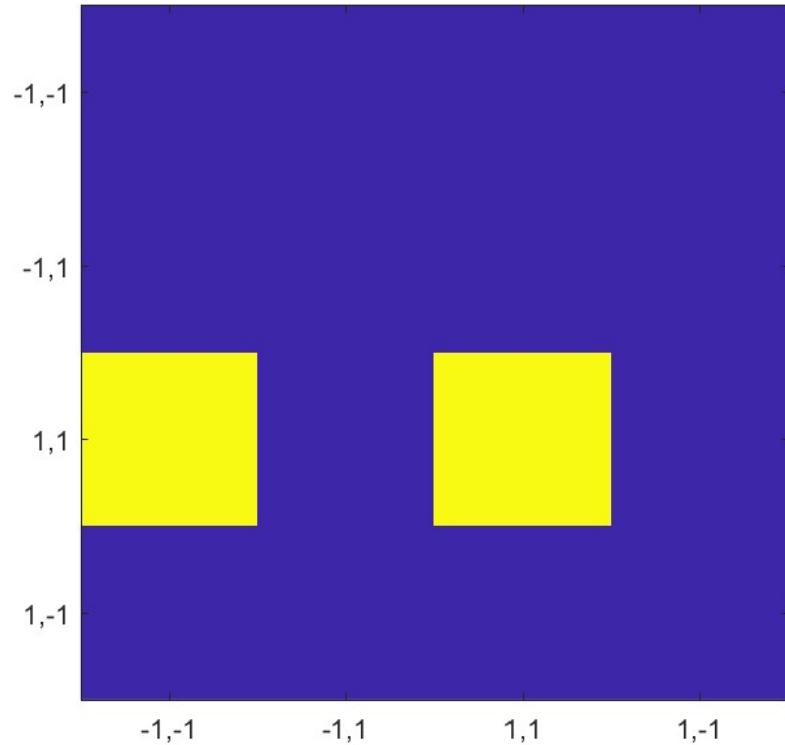


- Left: Pattern stored. Right: Energy map
- Note: Pattern is an energy well, but there are other local minima
  - Where?
  - Also note “shadow” pattern

# Storing multiple patterns: Orthogonality

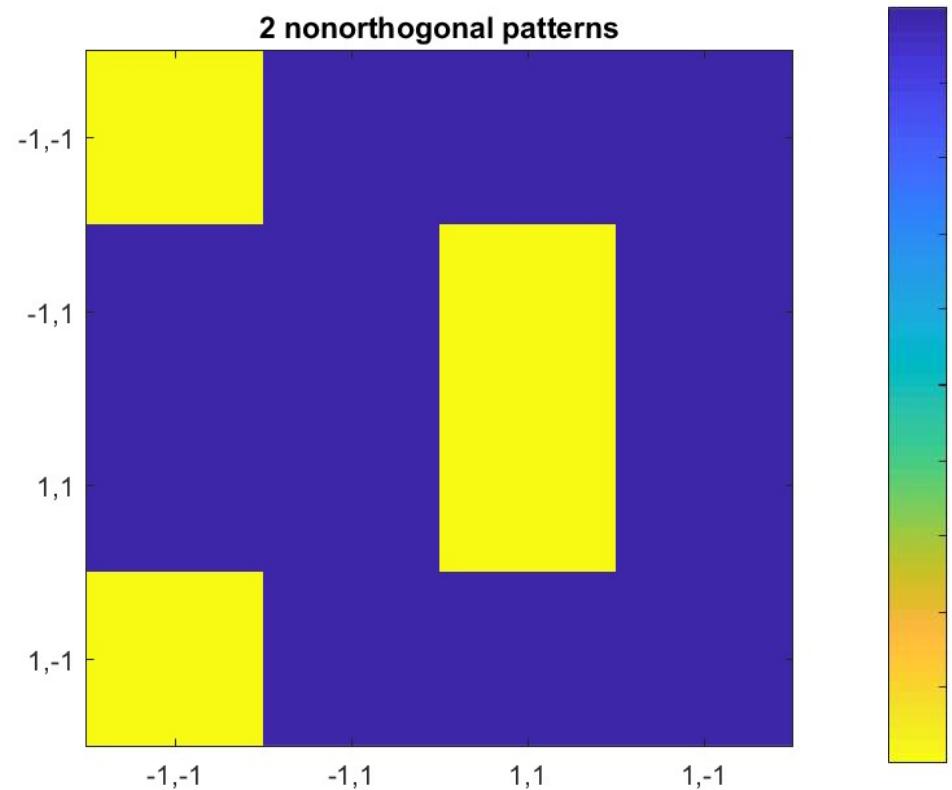
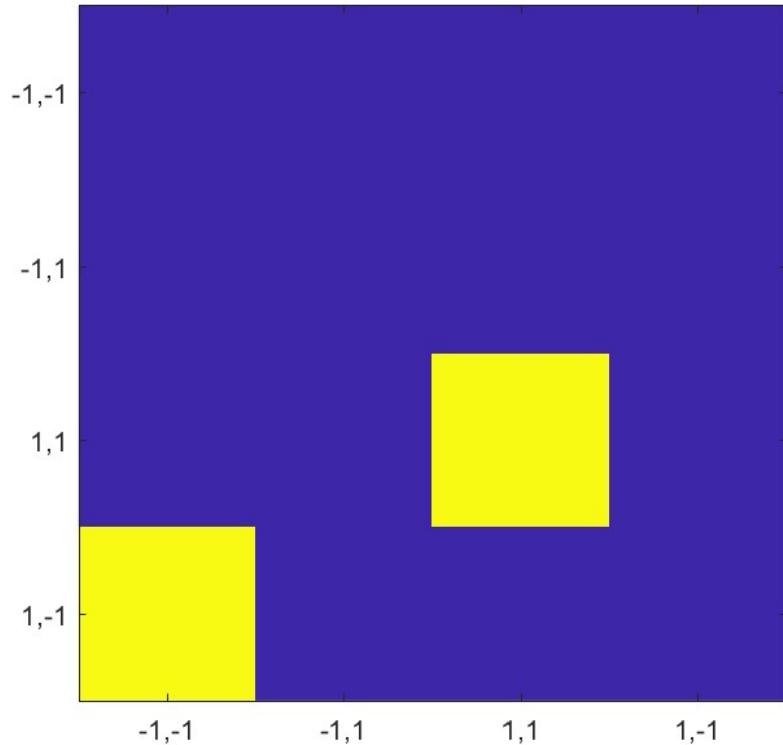
- The maximum Hamming distance between two  $N$ -bit patterns is  $N/2$ 
  - Because any pattern  $Y = -Y$  for our purpose
- Two patterns  $y_1$  and  $y_2$  that differ in  $N/2$  bits are *orthogonal*
  - Because  $y_1^T y_2 = 0$
- For  $N = 2^M L$ , where  $L$  is an odd number, there are at most  $2^M$  orthogonal binary patterns
  - Others may be *almost* orthogonal

# Two orthogonal 4-bit patterns



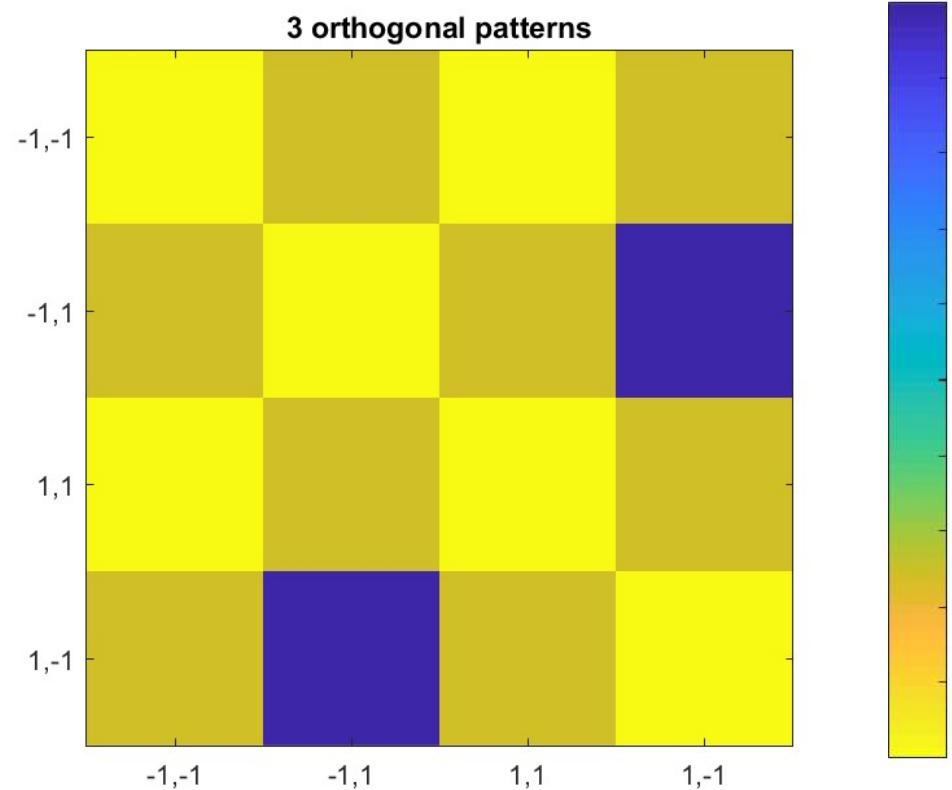
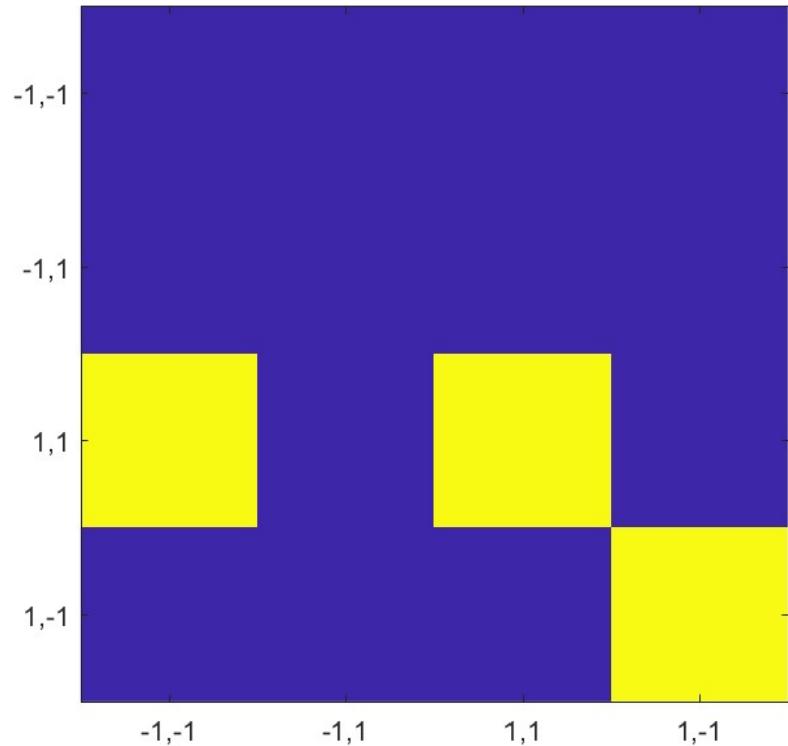
- Patterns are local minima (stationary and stable)
  - No other local minima exist
  - But patterns perfectly confusable for recall

# Two *non-orthogonal* 4-bit patterns



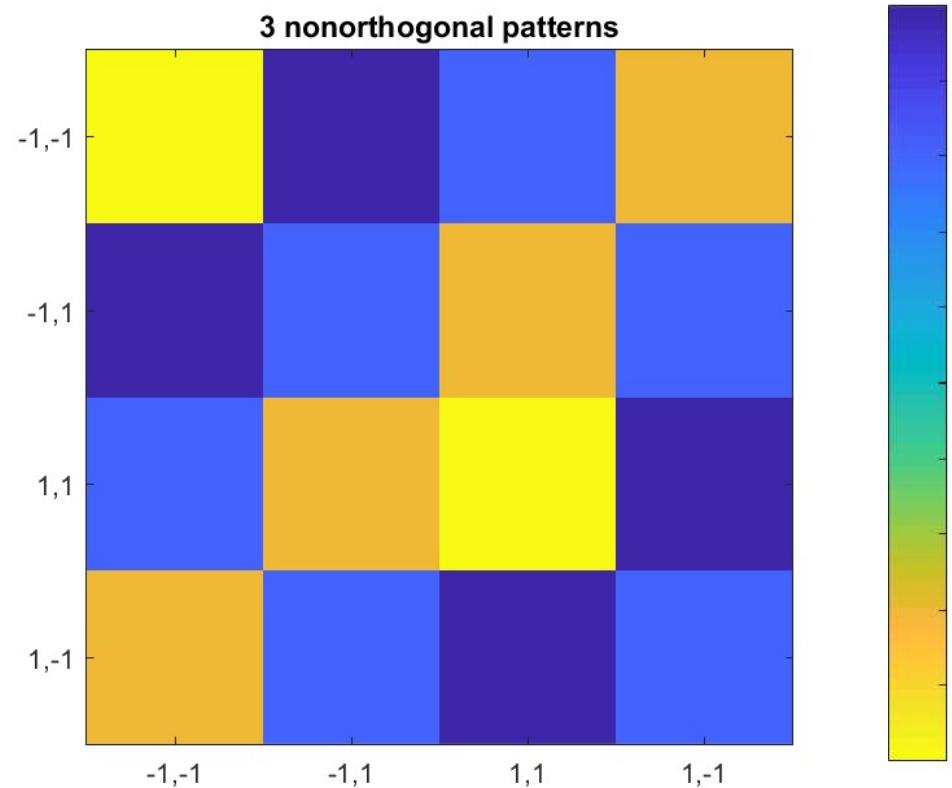
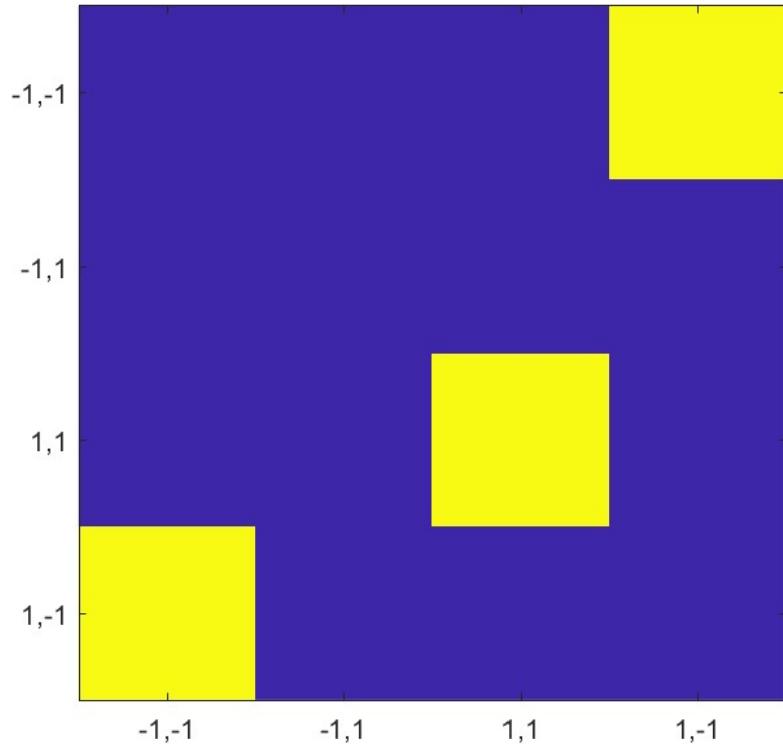
- Patterns are local minima (stationary and stable)
  - No other local minima exist
  - Actual *wells* for patterns
    - Patterns may be perfectly recalled!
  - Note  $K > 0.14 N$

# *Three orthogonal 4-bit patterns*



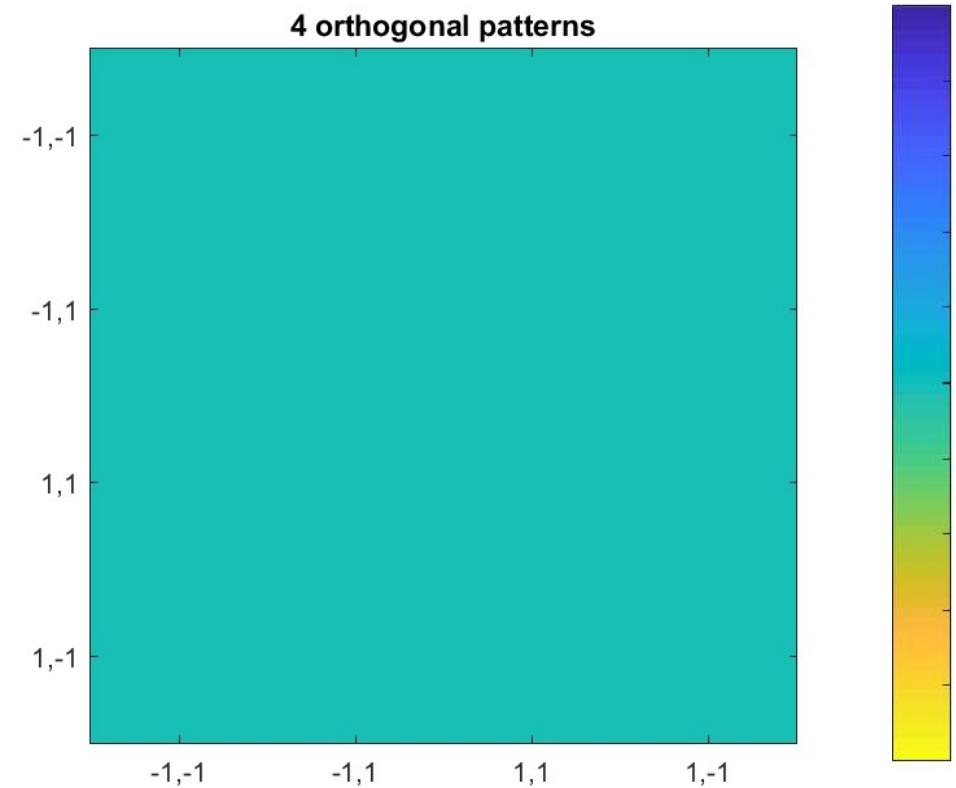
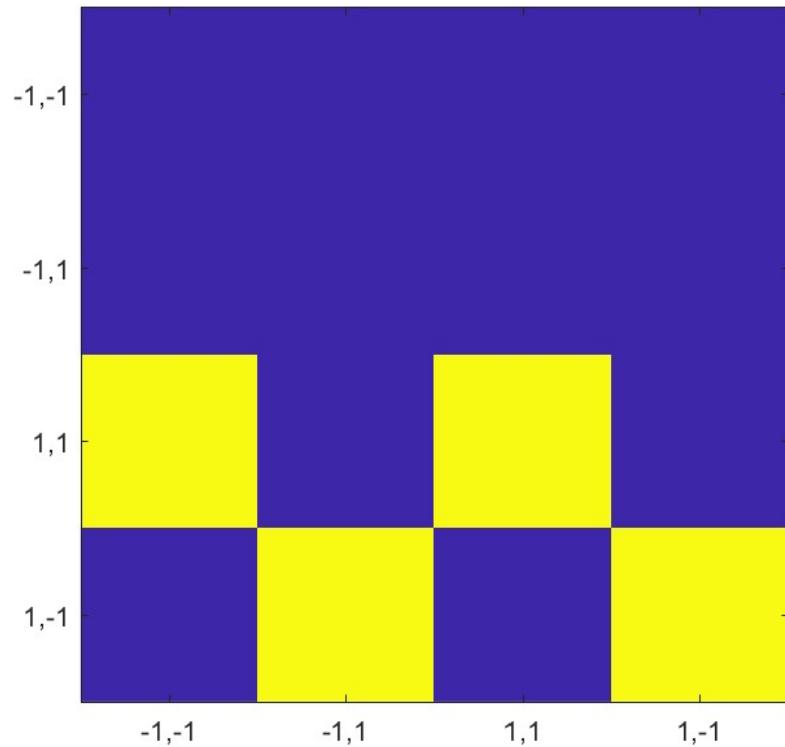
- All patterns are local minima (stationary)
  - But recall from perturbed patterns is random

# Three *non-orthogonal* 4-bit patterns



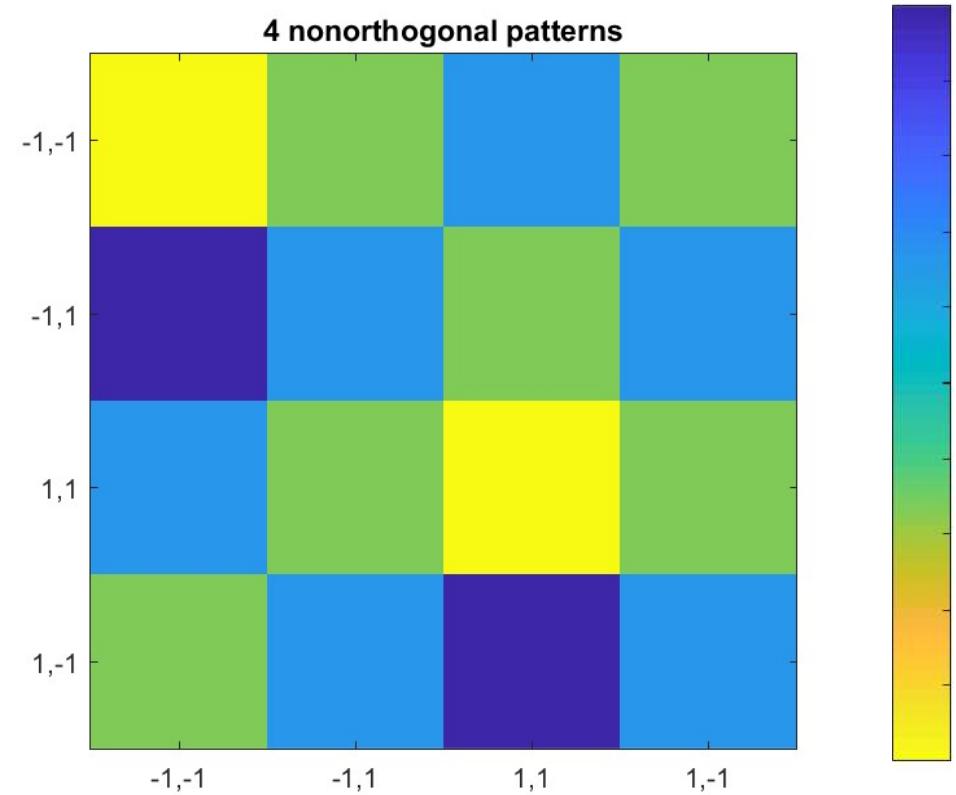
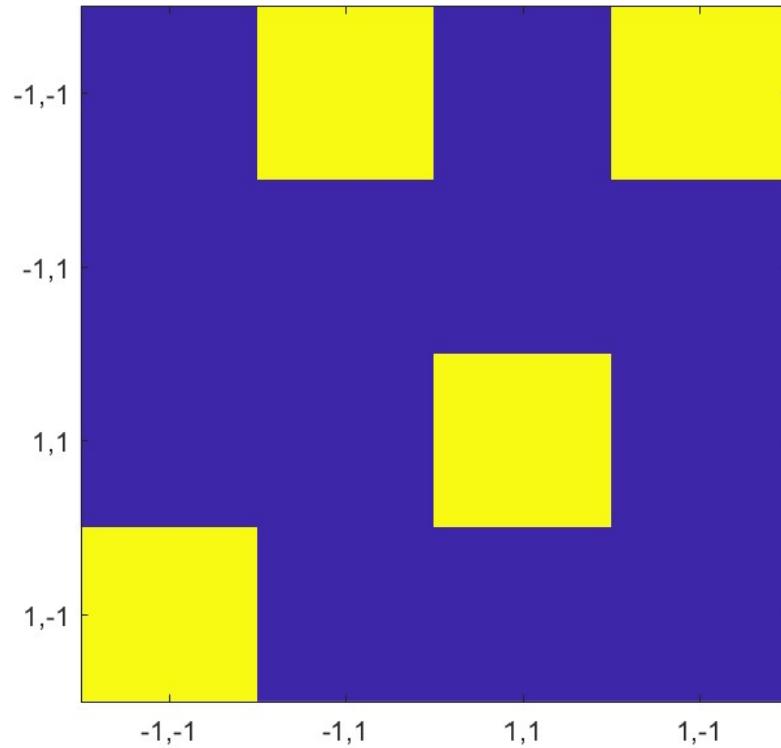
- Patterns in the corner are not recalled
  - They end up being attracted to the -1,-1 pattern
  - Note some “ghosts” ended up in the “well” of other patterns
    - So one of the patterns has stronger recall than the other two

# *Four orthogonal 4-bit patterns*



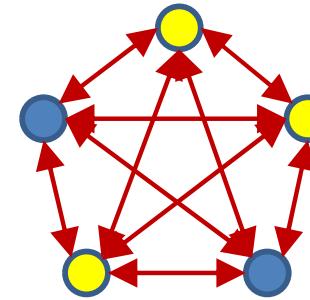
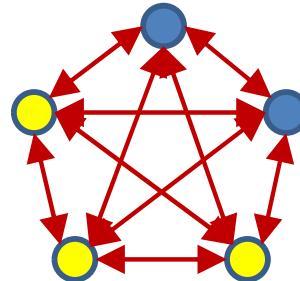
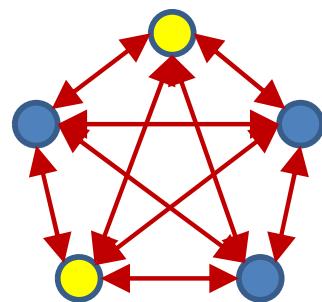
- All patterns are stationary, but none are stable
  - Total wipe out

# *Four nonorthogonal 4-bit patterns*



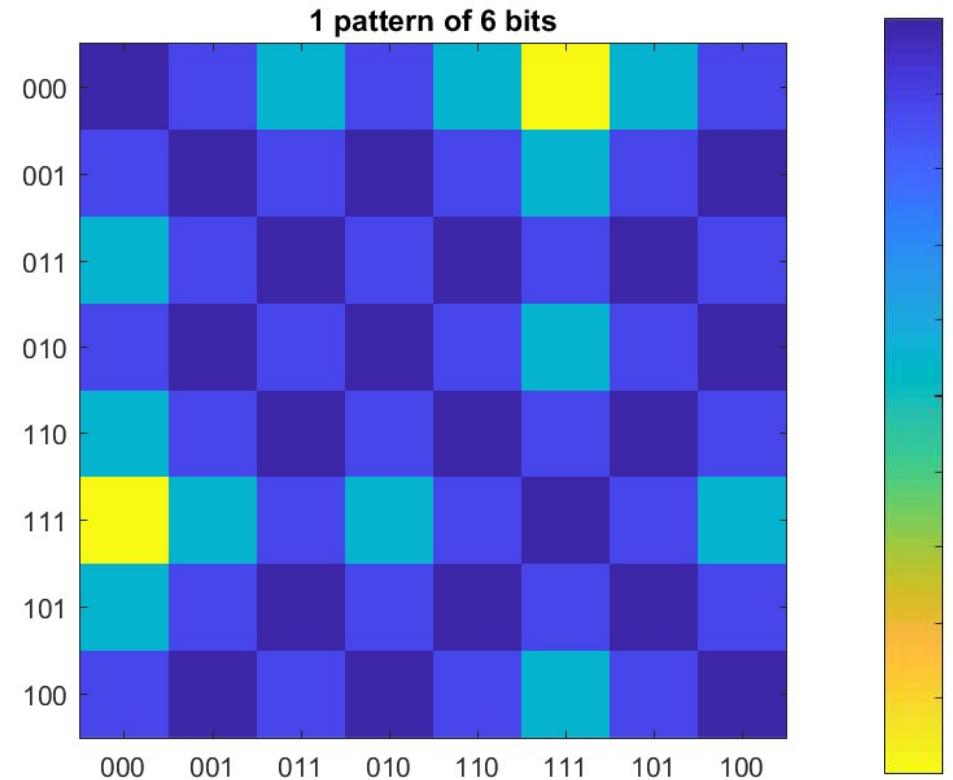
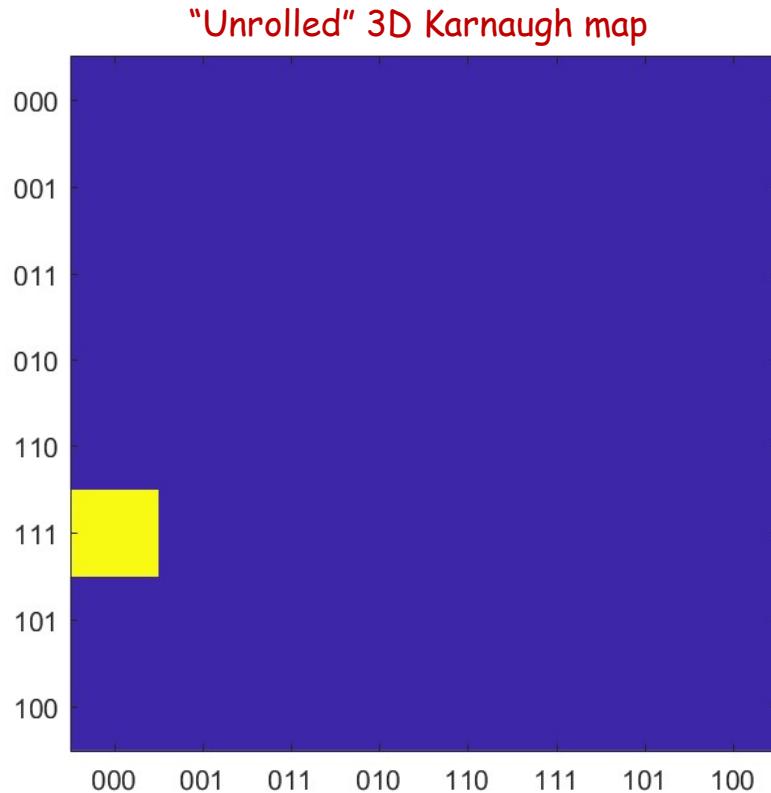
- One stable pattern
  - “Collisions” when the ghost of one pattern occurs next to another

# How many patterns can we store?



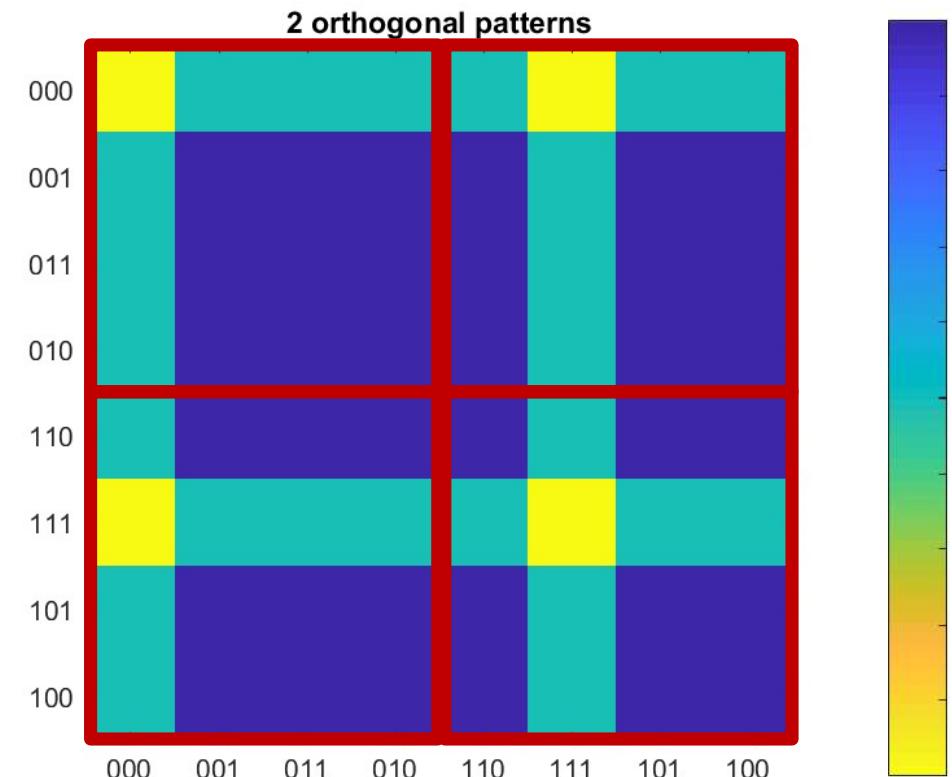
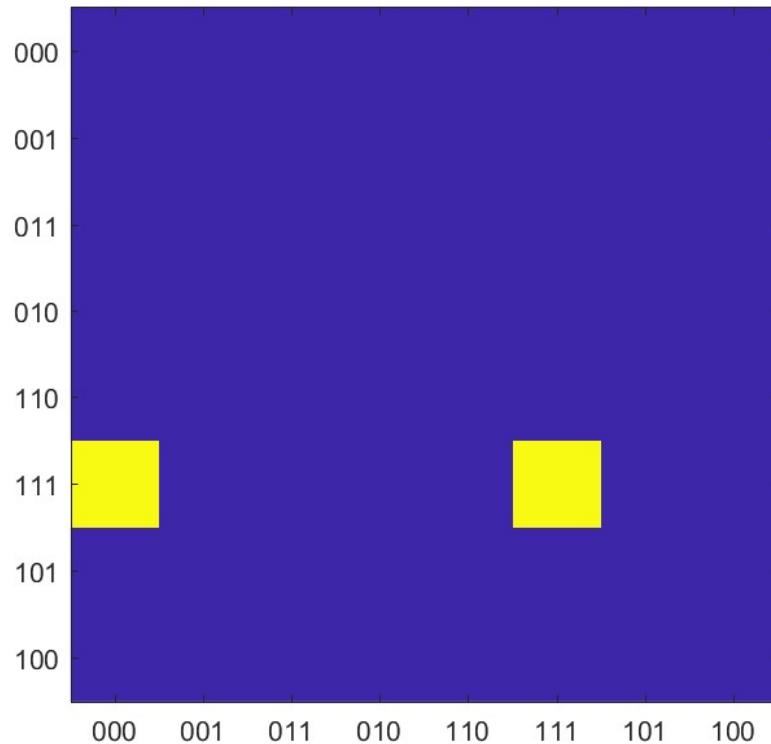
- Hopfield: For a network of  $N$  neurons can store up to  $0.14N$  random patterns
- Apparently a fuzzy statement
  - What does it really mean to say “stores”  $0.14N$  random patterns?
    - Stationary? Stable? No other local minima?
  - What if the patterns to store are not random?
- $N=4$  may not be a good case ( $N$  too small)

# A 6-bit pattern



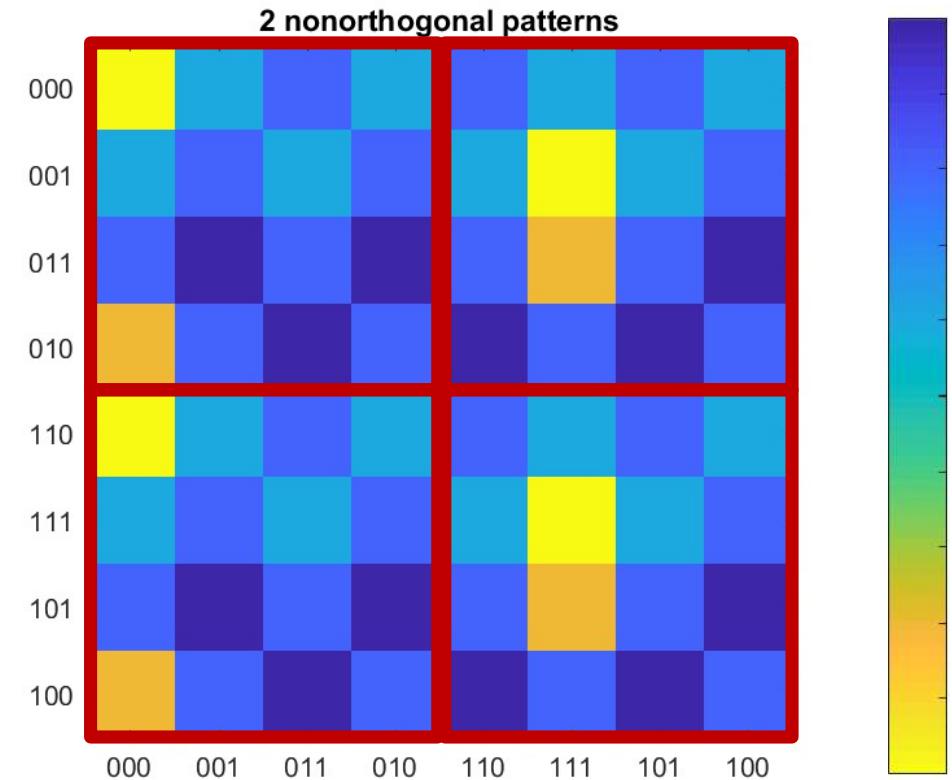
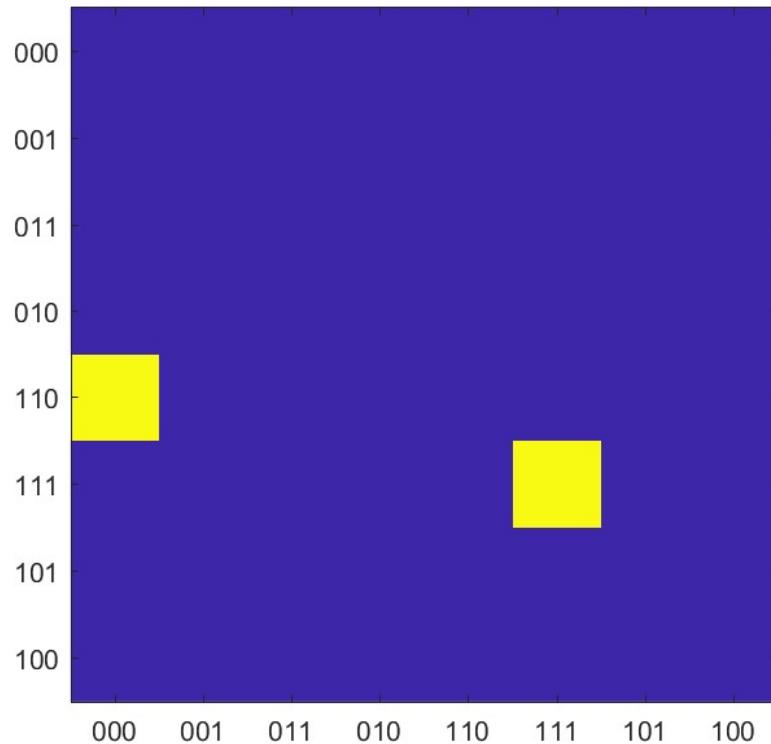
- Perfectly stationary and stable
- But many spurious local minima..
  - Which are “fake” memories

# Two orthogonal 6-bit patterns



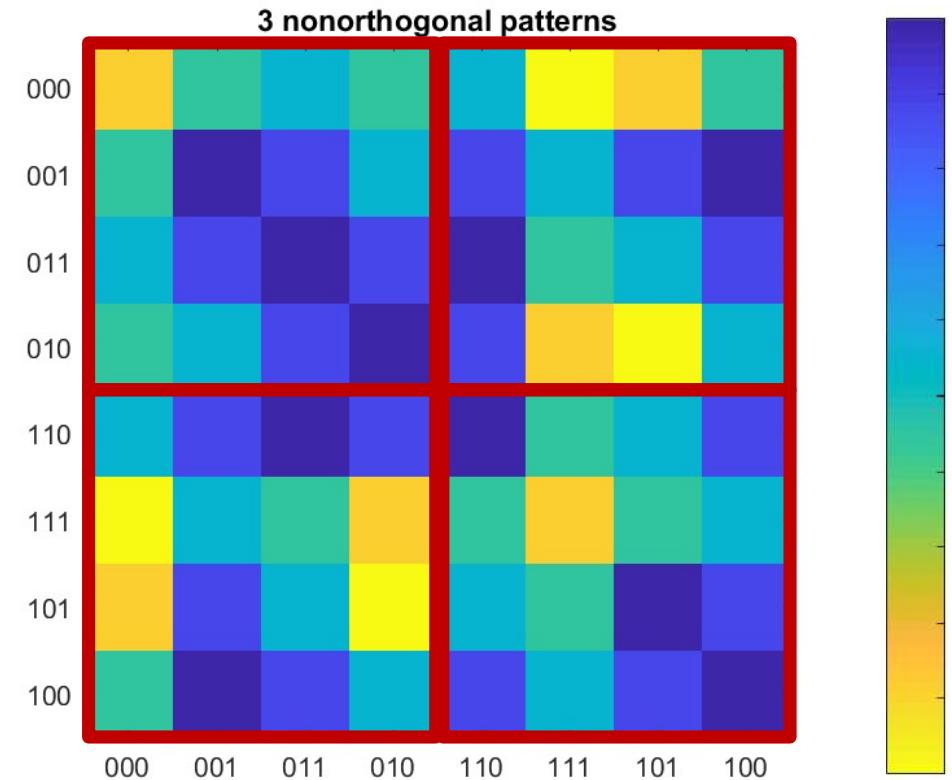
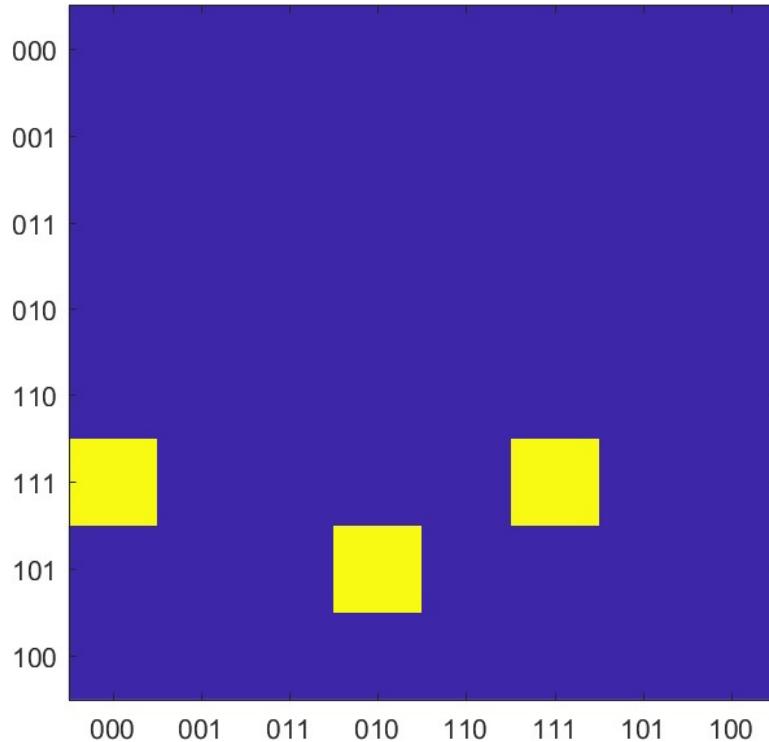
- Perfectly stationary and stable
- Several spurious “fake-memory” local minima..
  - Figure overstates the problem: actually a 3-D Kmap

# Two non-orthogonal 6-bit patterns



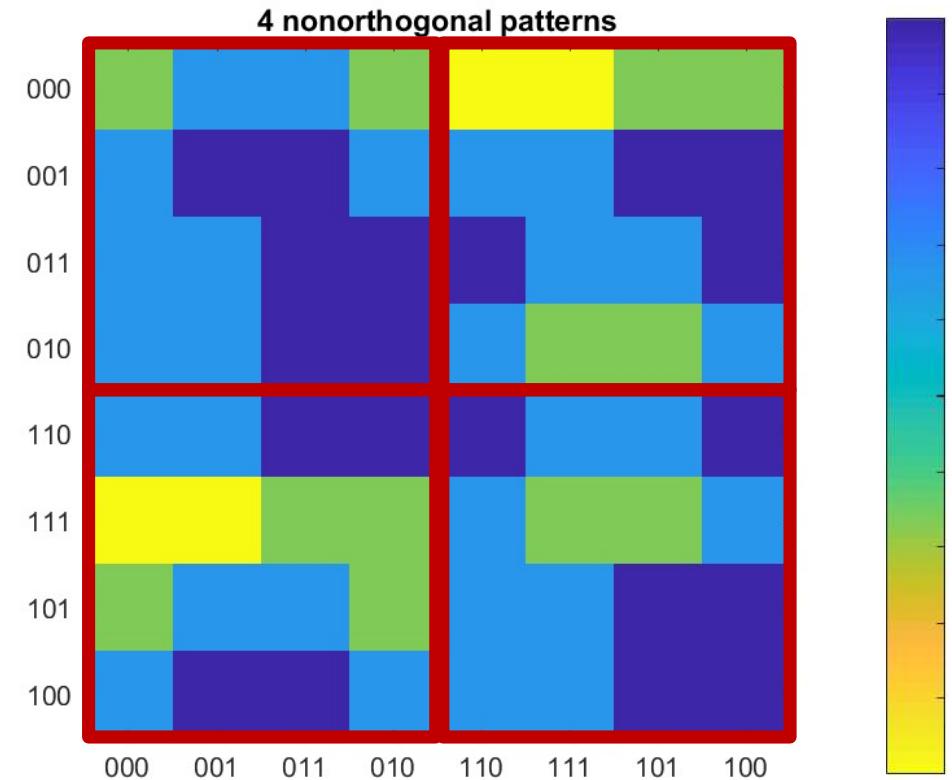
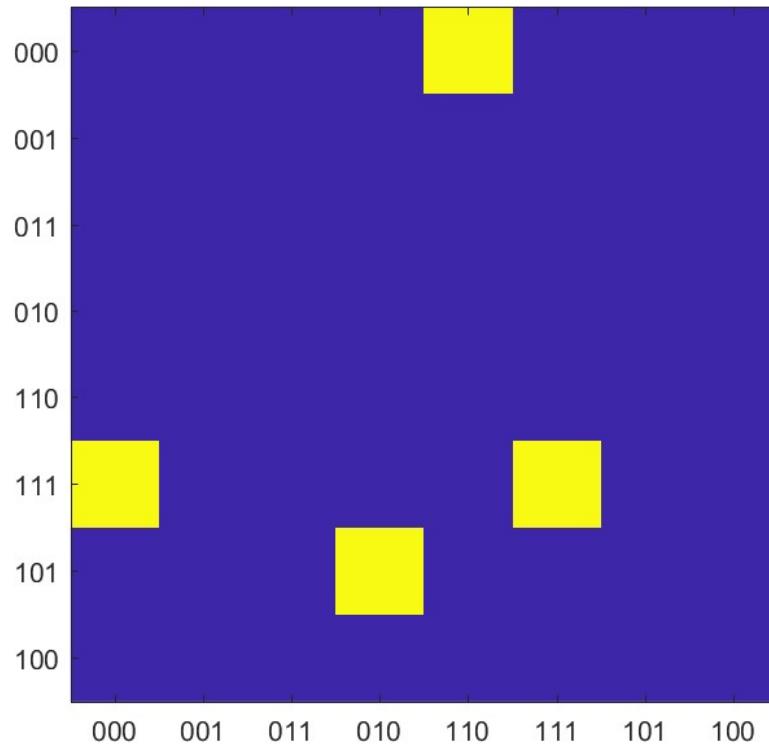
- Perfectly stationary and stable
- Some spurious “fake-memory” local minima..
  - But every stored pattern has “bowl”
  - *Fewer* spurious minima than for the orthogonal case

# Three *non-orthogonal* 6-bit patterns



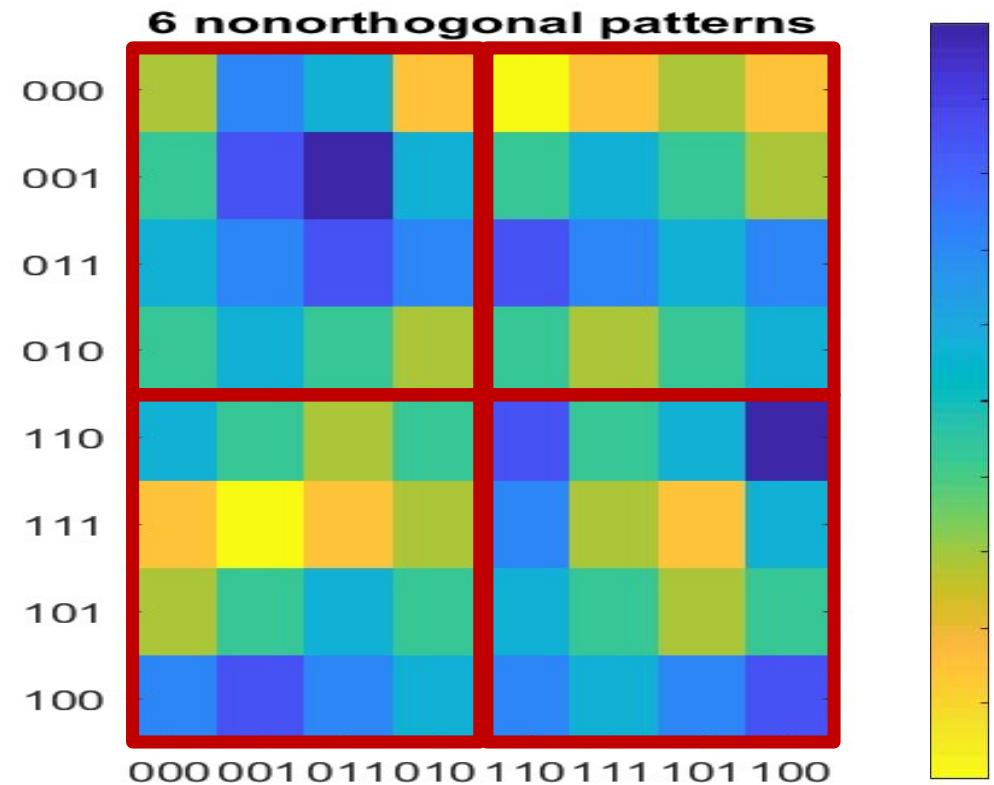
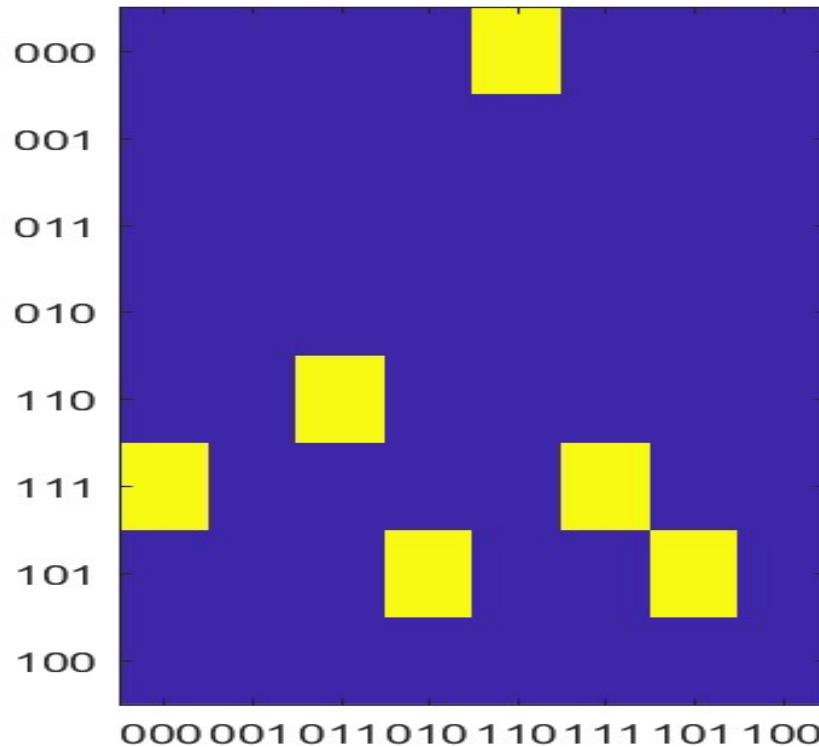
- Note: Cannot have 3 or more orthogonal 6-bit patterns..
- Patterns are perfectly stationary and stable ( $K > 0.14N$ )
- Some spurious “fake-memory” local minima..
  - But every stored pattern has “bowl”
  - *Fewer* spurious minima than for the orthogonal 2-pattern case

# Four *non-orthogonal* 6-bit patterns



- Patterns are perfectly stationary for  $K > 0.14N$
- *Fewer* spurious minima than for the orthogonal 2-pattern case
  - Most fake-looking memories are in fact ghosts..

# Six non-orthogonal 6-bit patterns

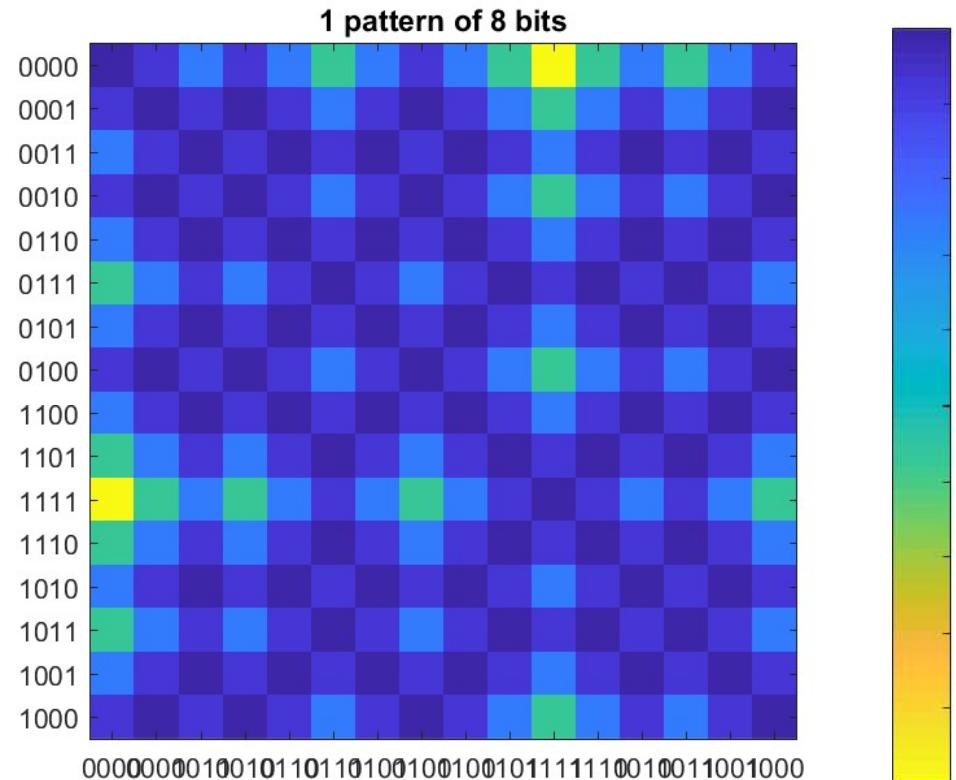
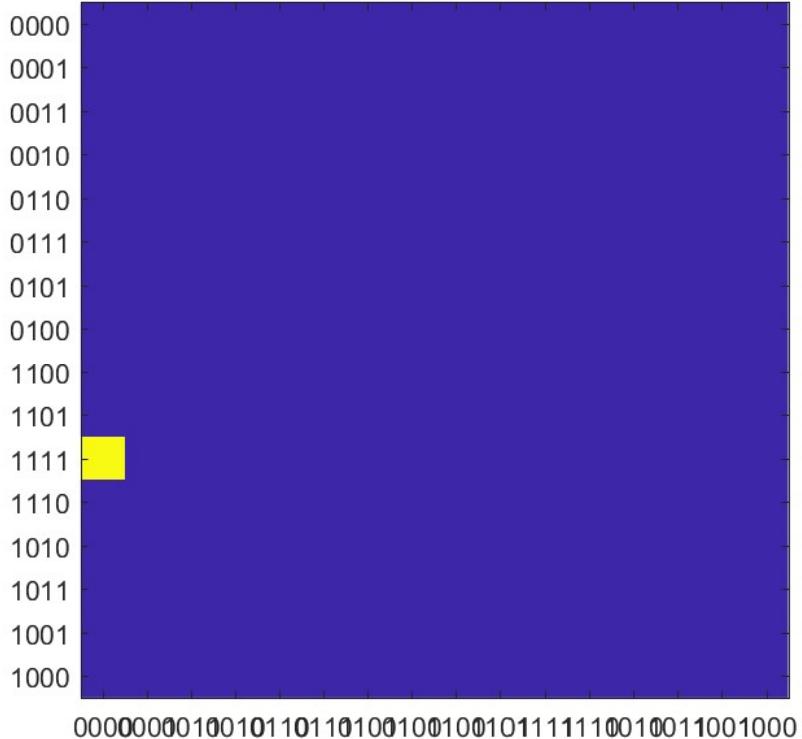


- Breakdown largely due to interference from “ghosts”
- But multiple patterns are stationary, and often stable
  - For  $K \gg 0.14N$

# More visualization..

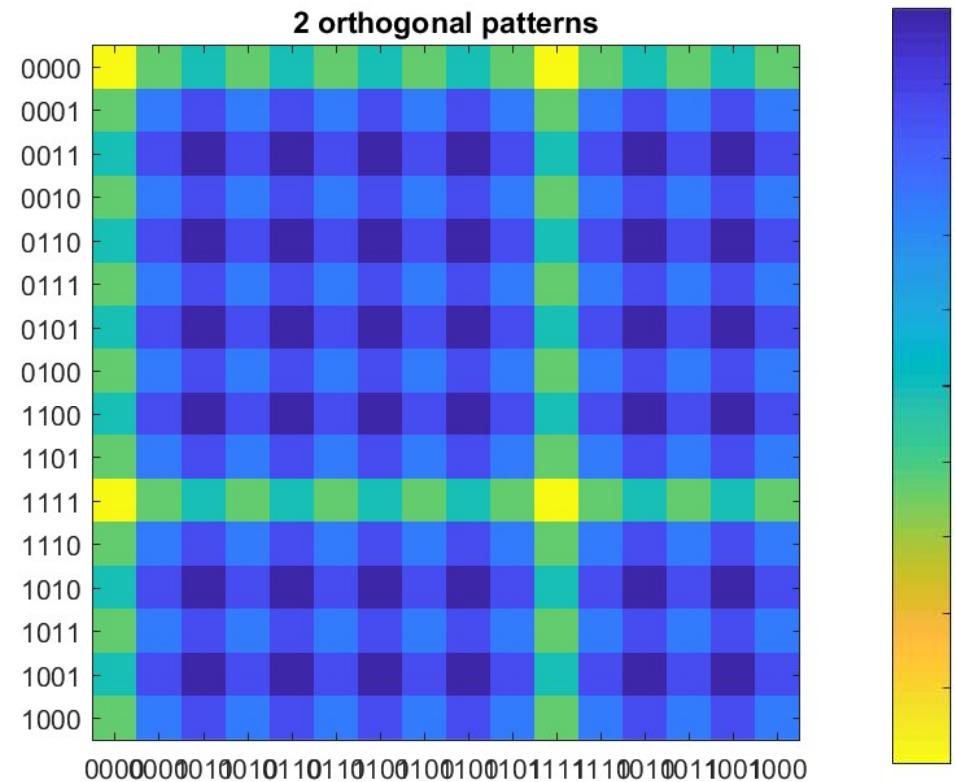
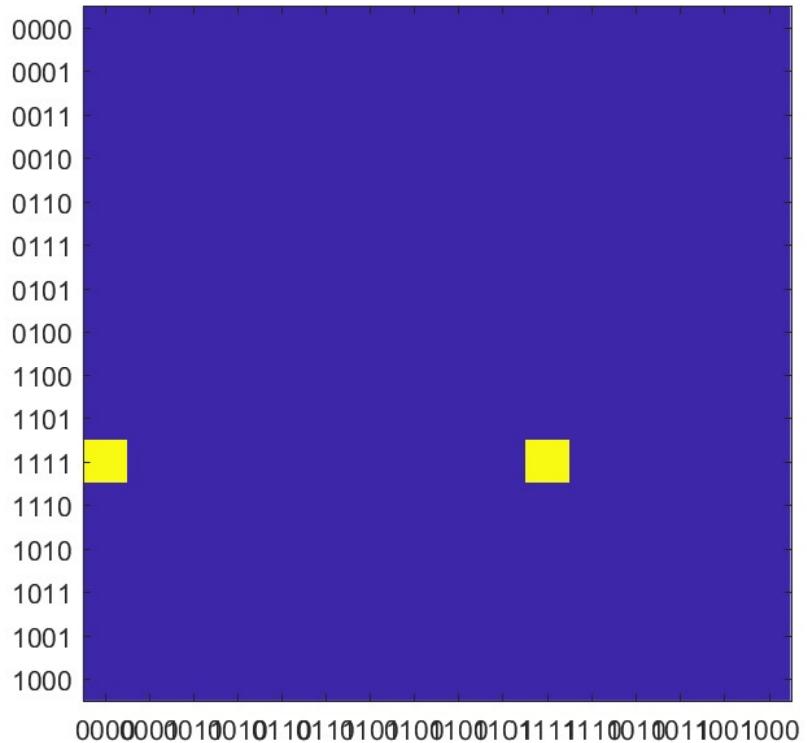
- Lets inspect a few 8-bit patterns
  - Keeping in mind that the Karnaugh map is now a 4-dimensional tesseract

# One 8-bit pattern



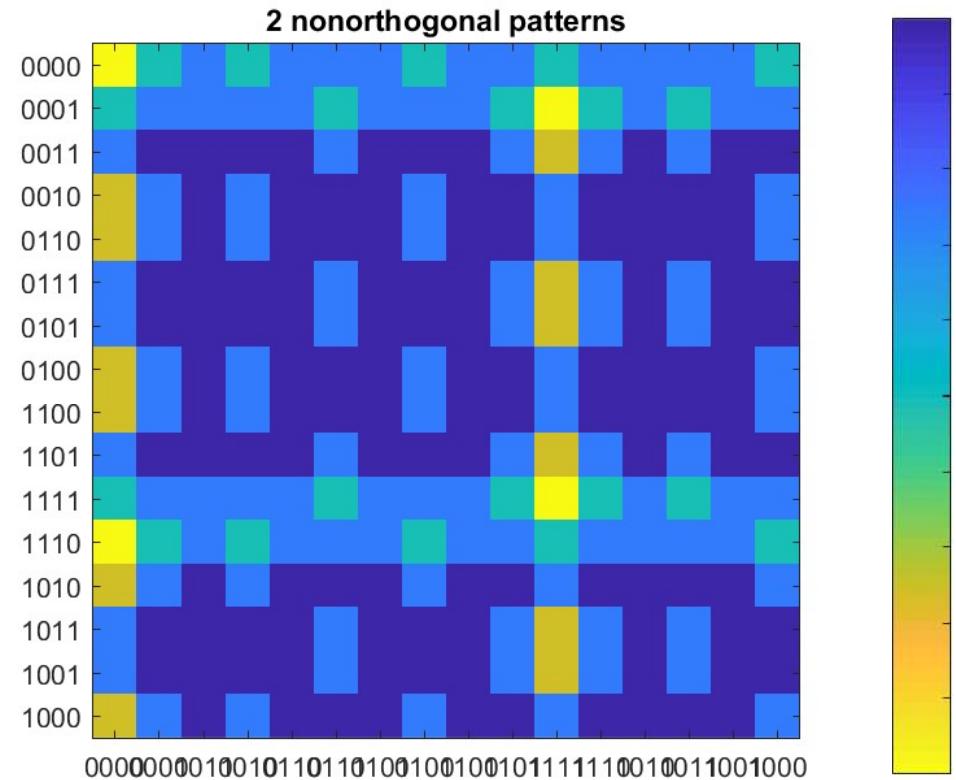
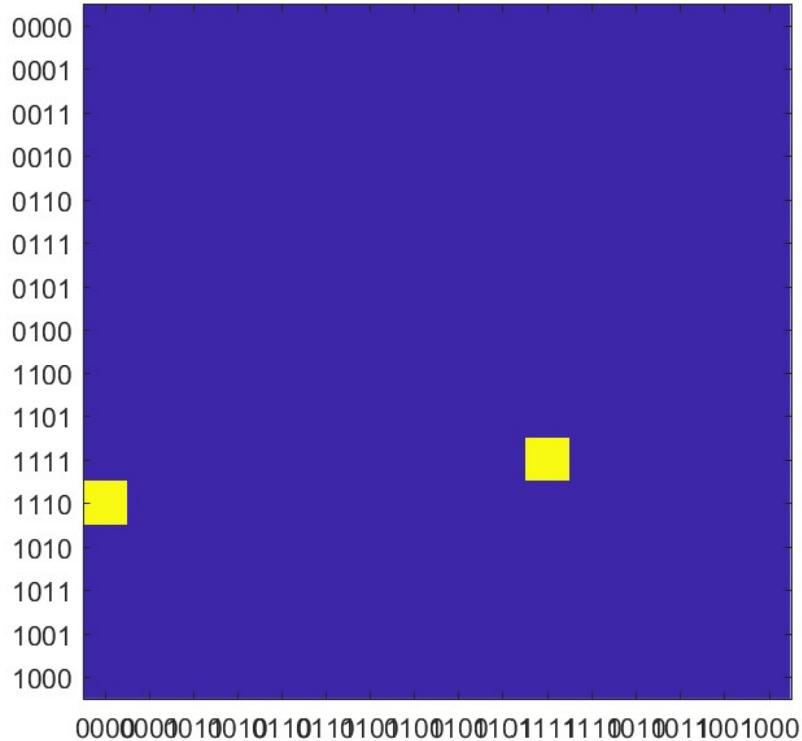
- Its actually cleanly stored, but there are a few spurious minima

# Two orthogonal 8-bit patterns



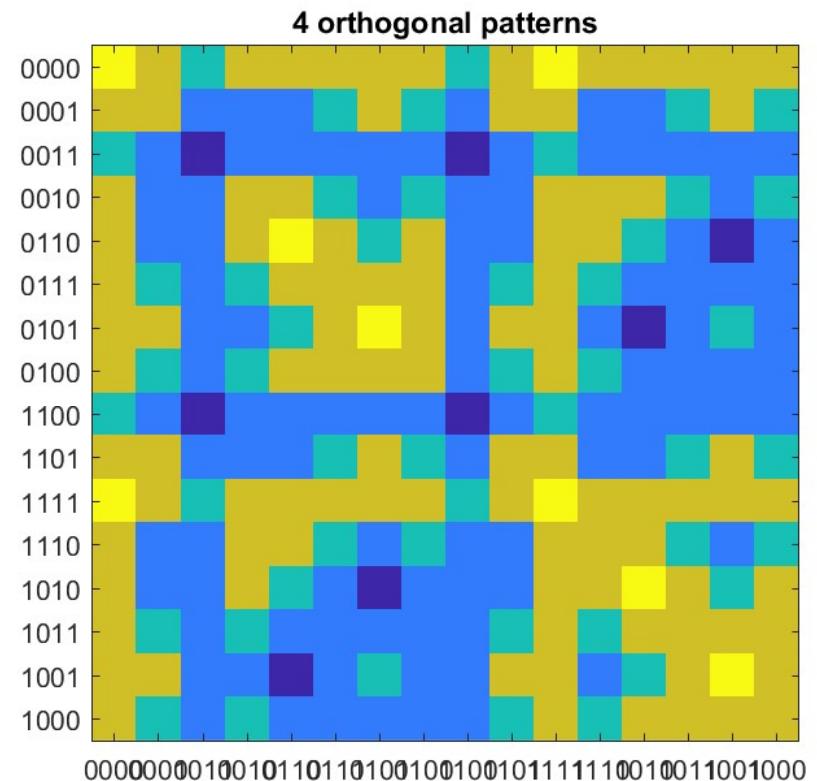
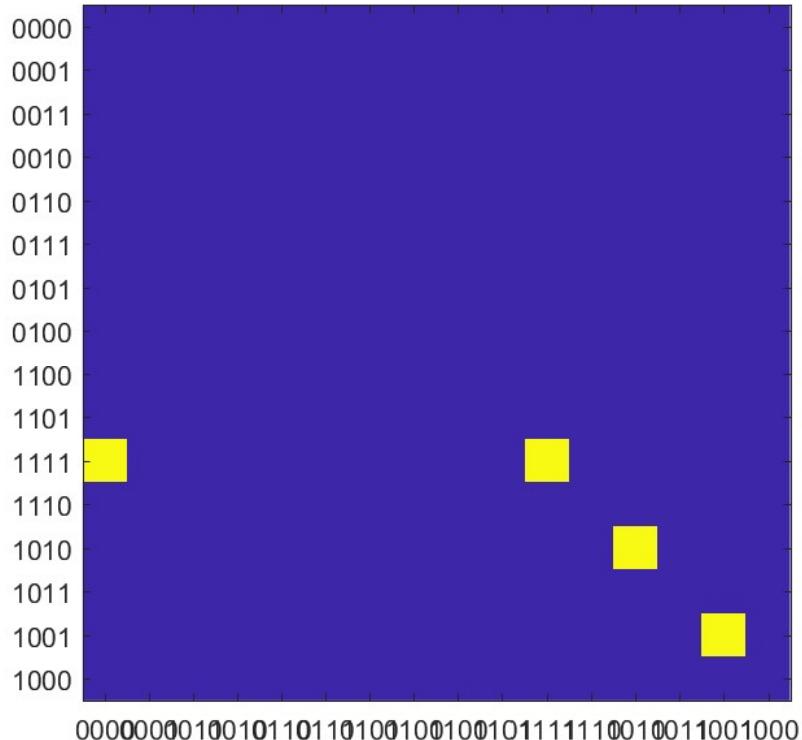
- Both have regions of attraction
- Some spurious minima

# Two non-orthogonal 8-bit patterns



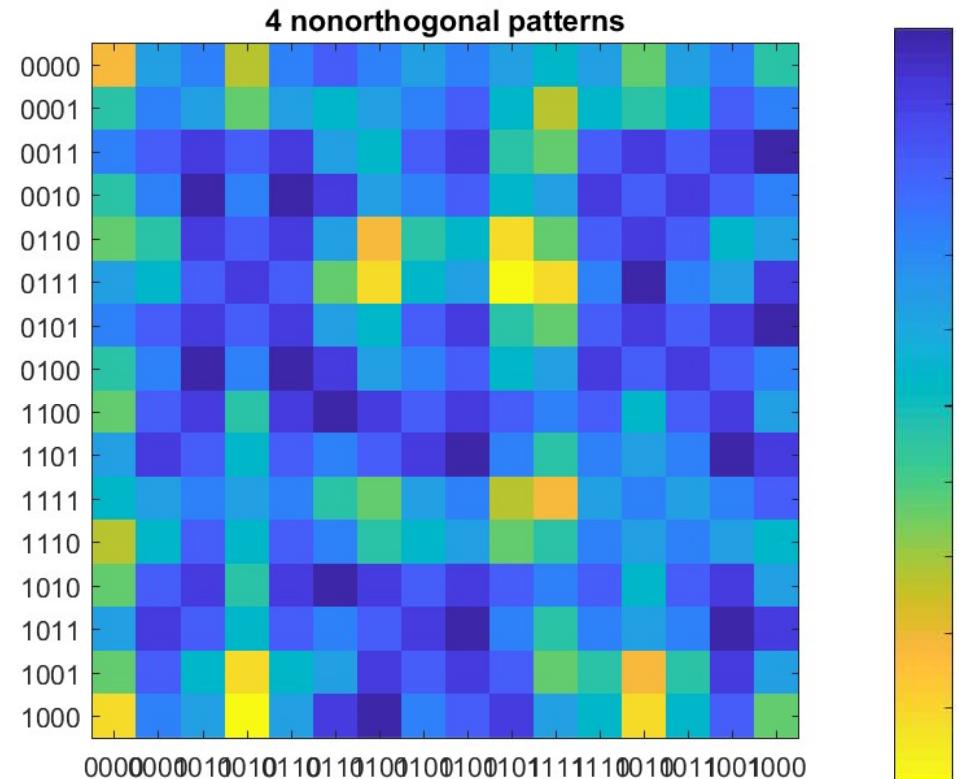
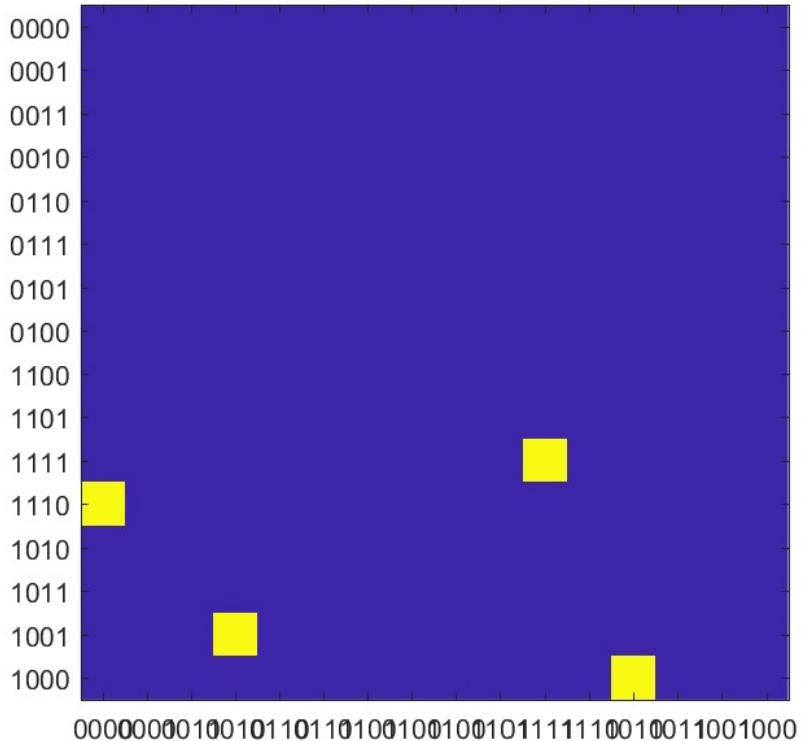
- Actually have fewer spurious minima
  - Not obvious from visualization..

# Four orthogonal 8-bit patterns



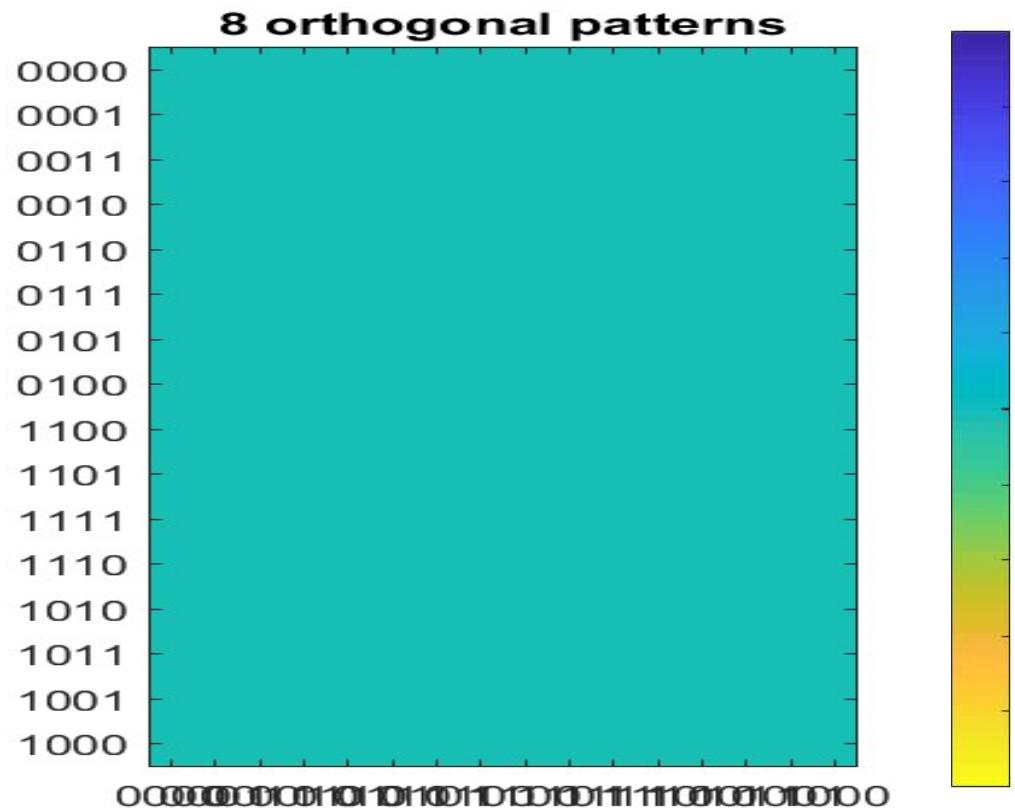
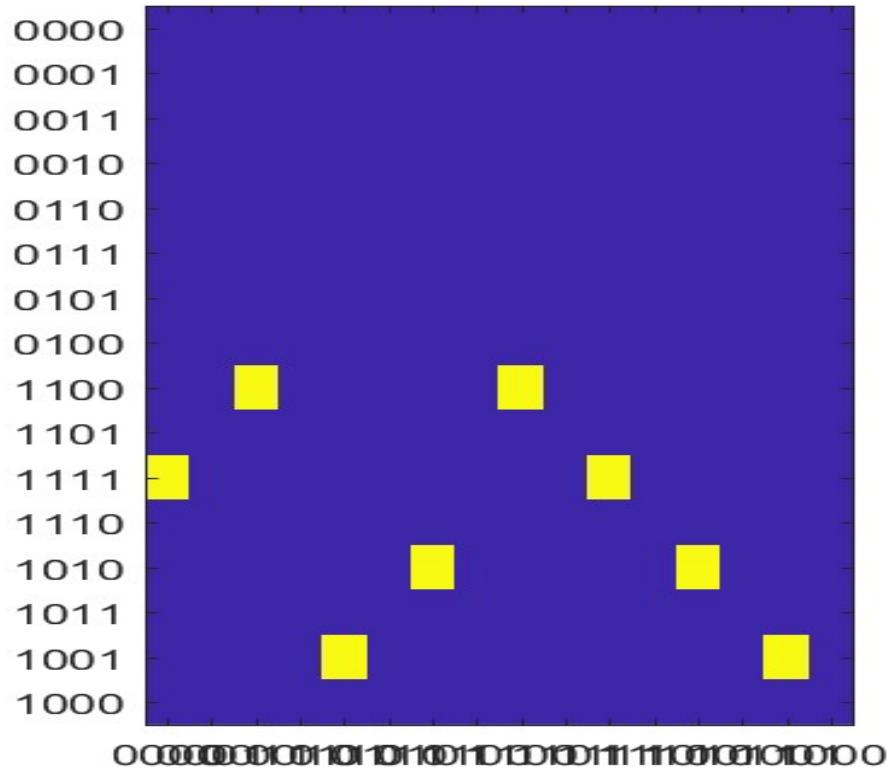
- Successfully stored

# Four non-orthogonal 8-bit patterns



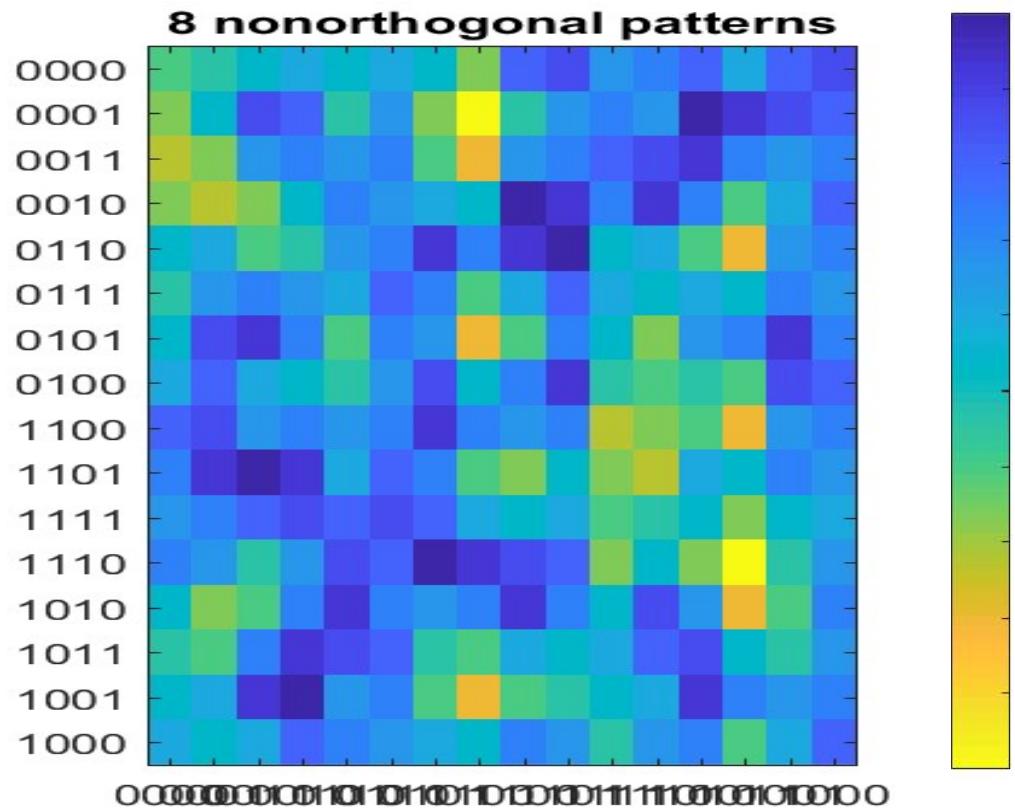
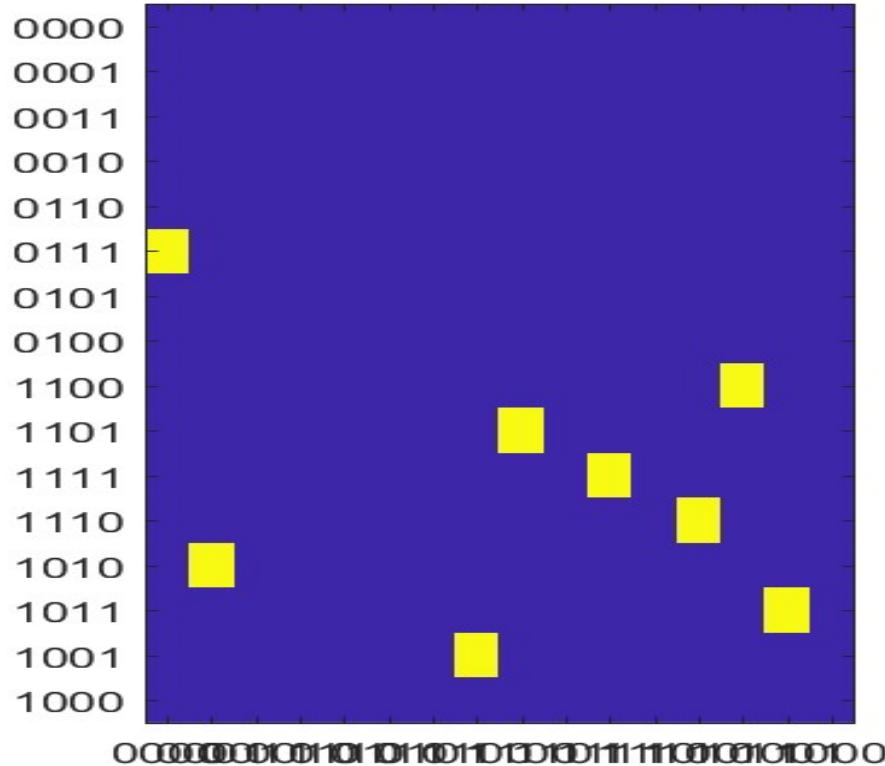
- Stored with interference from ghosts..

# Eight orthogonal 8-bit patterns



- Wipeout

# Eight non-orthogonal 8-bit patterns

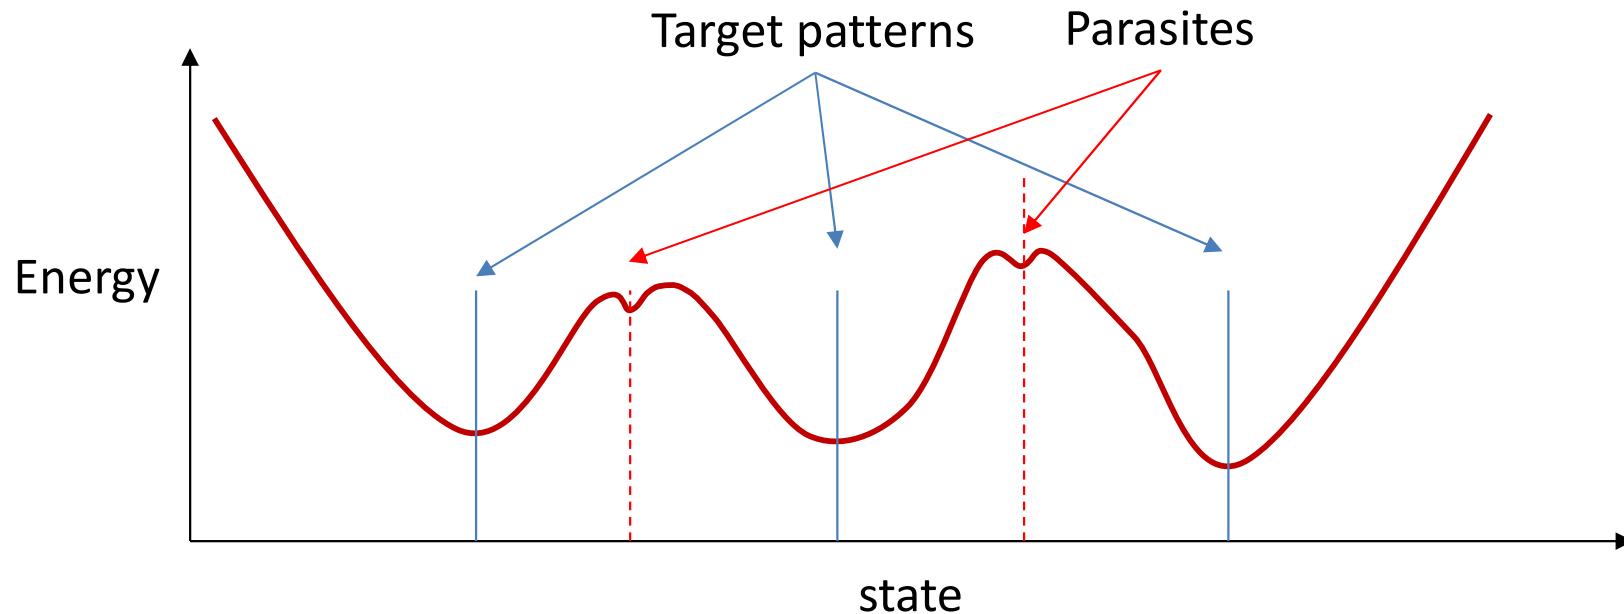


- Nothing stored
  - Neither stationary nor stable

# Observations

- Many “parasitic” patterns
  - Undesired patterns that also become stable or attractors
- Apparently, a capacity to store *more* than  $0.14N$  patterns

# Parasitic Patterns



- Parasitic patterns can occur because sums of odd numbers of stored patterns are also stable for Hebbian learning:
  - $\mathbf{y}_{parasite} = \text{sign}(\mathbf{y}_a + \mathbf{y}_b + \mathbf{y}_c)$
- They are also from other random local energy minima from the weights matrices themselves

# Capacity

- Seems possible to store  $K > 0.14N$  patterns
  - i.e. obtain a weight matrix  $W$  such that  $K > 0.14N$  patterns are stationary
  - Possible to make more than  $0.14N$  patterns at-least 1-bit stable
- Patterns that are *non-orthogonal* easier to remember
  - I.e. patterns that are *closer* are easier to remember than patterns that are farther!!
- Can we attempt to get greater control on the process than Hebbian learning gives us?
  - Can we do *better* than Hebbian learning?
    - Better capacity and fewer spurious memories?

# Story so far

- A Hopfield network is a loopy binary net with symmetric connections
  - Neurons try to align themselves to the local field caused by other neurons
- Given an initial configuration, the patterns of neurons in the net will evolve until the “energy” of the network achieves a local minimum
  - The network acts as a *content-addressable* memory
    - Given a damaged memory, it can evolve to recall the memory fully
- The network must be designed to store the desired memories
  - Memory patterns must be *stationary* and *stable* on the energy contour
- Network memory can be trained by Hebbian learning
  - Guarantees that a network of  $N$  bits trained via Hebbian learning can store  $0.14N$  random patterns with less than 0.4% probability that they will be unstable
- However, empirically it appears that we may sometimes be able to store *more than*  $0.14N$  patterns

# Poll 3: @1798

Mark all that are true

- We can try to “assign” memories to a Hopfield network through Hebbian learning of the weights matrix
- All patterns learned through Hebbian learning will be “remembered”
- The N-bit Hopfield network has the capacity to remember up to  $0.14N$  patterns

# Poll 3

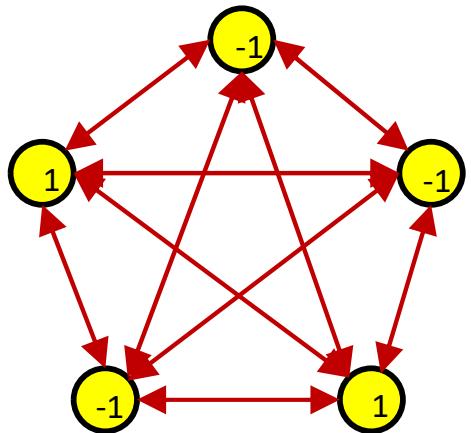
Mark all that are true

- We can try to “assign” memories to a Hopfield network through Hebbian learning of the weights matrix
- All patterns learned through Hebbian learning will be “remembered”
- The N-bit Hopfield network has the capacity to remember up to  $0.14N$  patterns

# Bold Claim

- I can *always* store (upto) N orthogonal patterns such that they are stationary!
  - Why?
- I can avoid spurious memories by adding some noise during recall!

# Recap: Hebbian Learning to Store a Specific Pattern



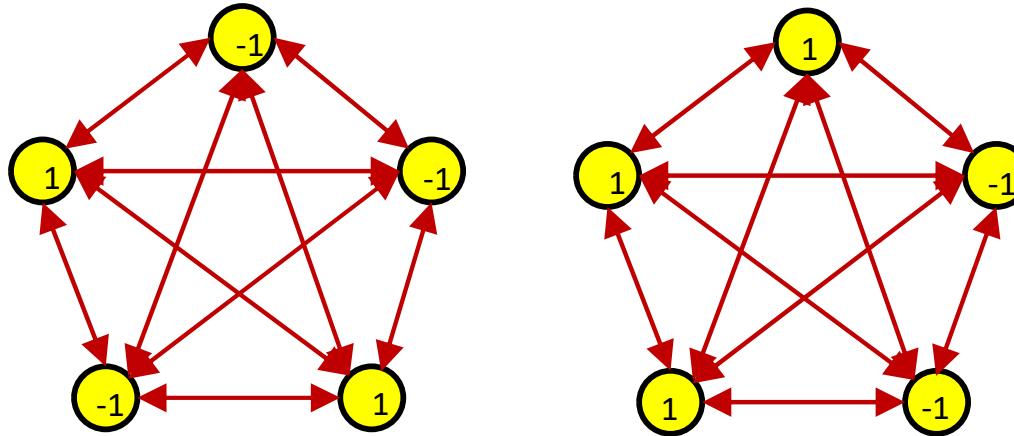
HEBBIAN LEARNING:

$$w_{ji} = y_j y_i$$

$$\mathbf{W} = \mathbf{y}_p \mathbf{y}_p^T - \mathbf{I}$$

- For a single stored pattern, Hebbian learning results in a network for which the target pattern is a global minimum

# Storing multiple patterns



- Let  $\mathbf{y}_p$  be the vector representing  $p$ -th pattern
- Let  $\mathbf{Y} = [\mathbf{y}_1 \ \mathbf{y}_2 \ \dots]$  be a matrix with all the stored patterns
- Then..

$$\mathbf{W} = \frac{1}{N_p} \sum_p (\mathbf{y}_p \mathbf{y}_p^T - \mathbf{I}) = \frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I}$$

Number of patterns

$\mathbf{W}$  is a positive semi-definite matrix

# A minor adjustment

- Note behavior of  $\mathbf{E}(\mathbf{y}) = \mathbf{y}^T \mathbf{W} \mathbf{y}$  with

$$\mathbf{W} = \frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I}$$

- Is identical to behavior with

$$\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$$

- Since

$$\mathbf{y}^T \left( \frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I} \right) \mathbf{y} = \frac{1}{N_p} \mathbf{y}^T \mathbf{Y} \mathbf{Y}^T \mathbf{y} - N$$

- But  $\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$  is easier to analyze. Hence in the following slides we will use  $\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$

Energy landscape  
only differs by  
an additive constant and  
a scaling  
  
Location  
of minima remain same

# A minor adjustment

- Note behavior of  $\mathbf{E}(\mathbf{y}) = \mathbf{y}^T \mathbf{W} \mathbf{y}$  with

$$\mathbf{W} = \frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I}$$

Both have the same Eigen vectors

behavior with

$$\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$$

- Since

Energy landscape only differs by an additive constant and a scaling

Location of minima remain same

$$\mathbf{y}^T \left( \frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I} \right) \mathbf{y} = \frac{1}{N_p} \mathbf{y}^T \mathbf{Y} \mathbf{Y}^T \mathbf{y} - N$$

- But  $\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$  is easier to analyze. Hence in the following slides we will use  $\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$

# A minor adjustment

- Note behavior of  $\mathbf{E}(\mathbf{y}) = \mathbf{y}^T \mathbf{W} \mathbf{y}$  with

$$\mathbf{W} = \frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I}$$

Both have the same Eigen vectors

behavior with

$$\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$$

Energy landscape only differs by an additive constant and a scaling

Location of minima remain same

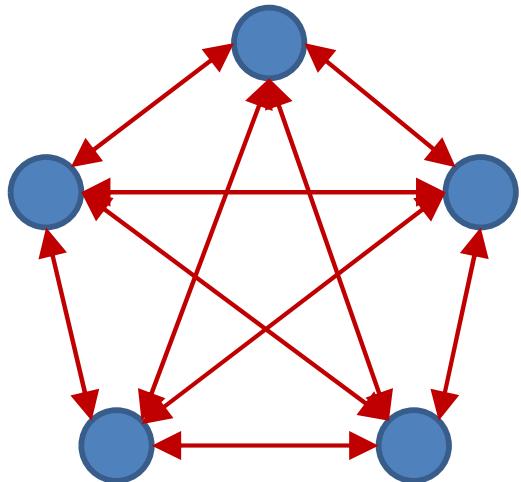
- So:

NOTE: This is a positive semidefinite matrix

$$(\frac{1}{N_p} \mathbf{Y} \mathbf{Y}^T - \mathbf{I}) \mathbf{y} = \frac{1}{N_p} \mathbf{y}^T \mathbf{Y} \mathbf{Y}^T \mathbf{y} - N$$

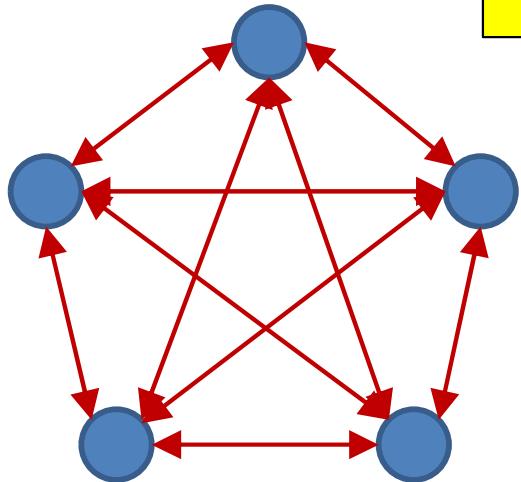
- But  $\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$  is easier to analyze. Hence in the following slides we will use  $\mathbf{W} = \mathbf{Y} \mathbf{Y}^T$

# Consider the energy function



$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$

# Consider the energy function



This is a quadratic!

For Hebbian learning  
W is positive semidefinite

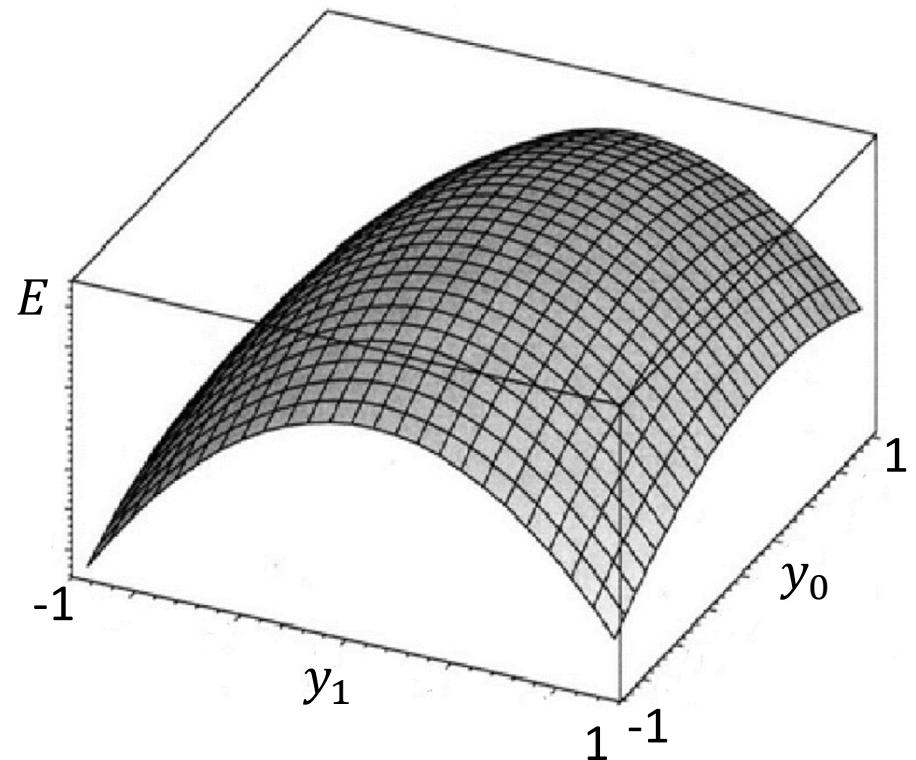
E is concave

$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$

- The Energy function is concave if **W** is positive (semi) definite

# The Energy function

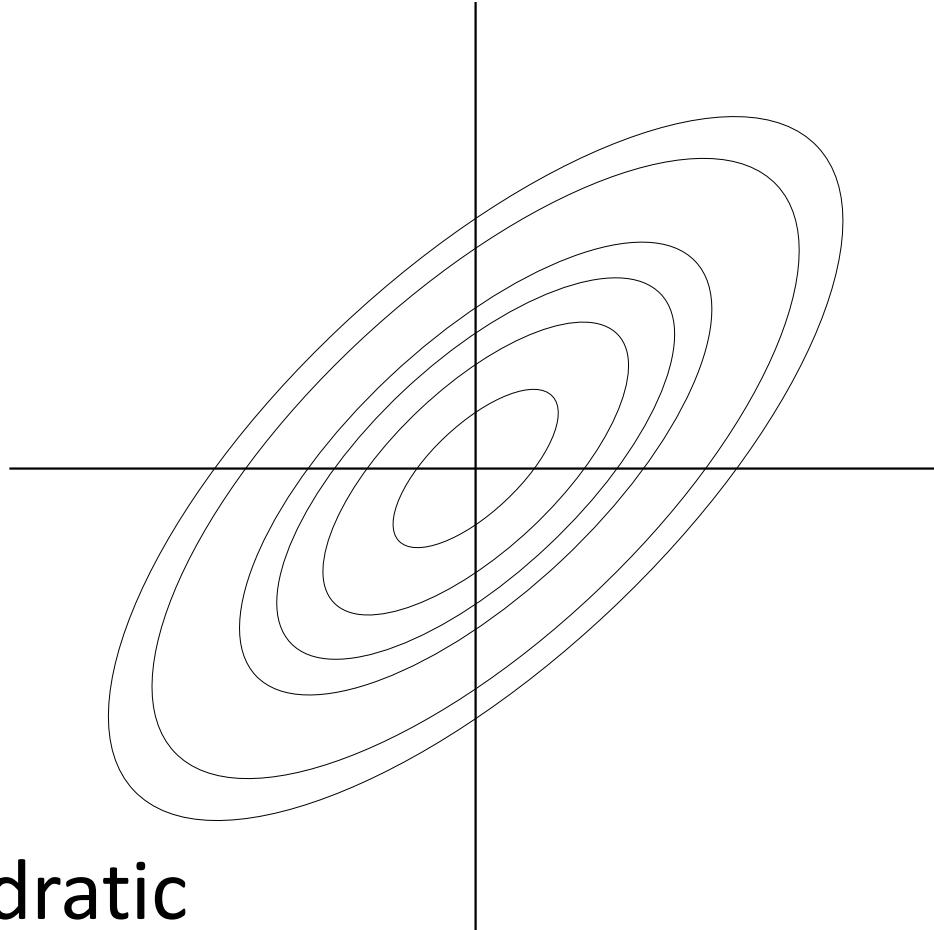
$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$



- $E$  is a concave quadratic

# The Energy function

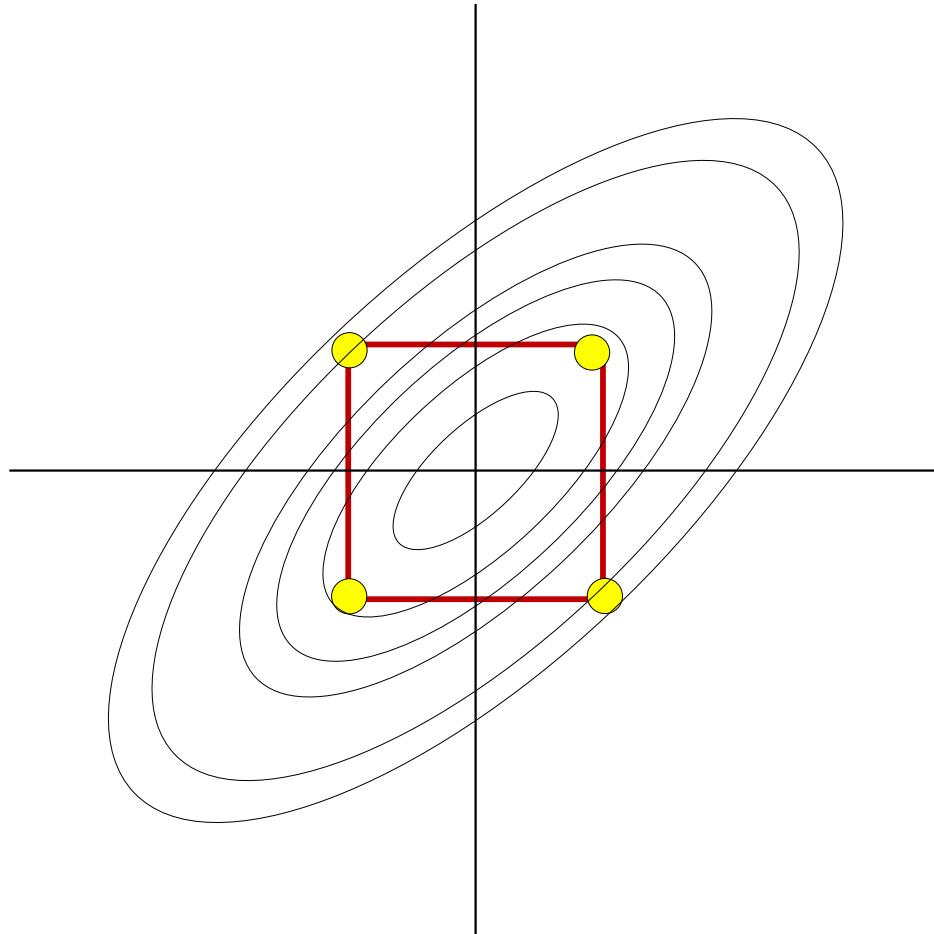
$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$



- $E$  is a concave quadratic
  - Shown from above (assuming 0 bias)

# The energy function

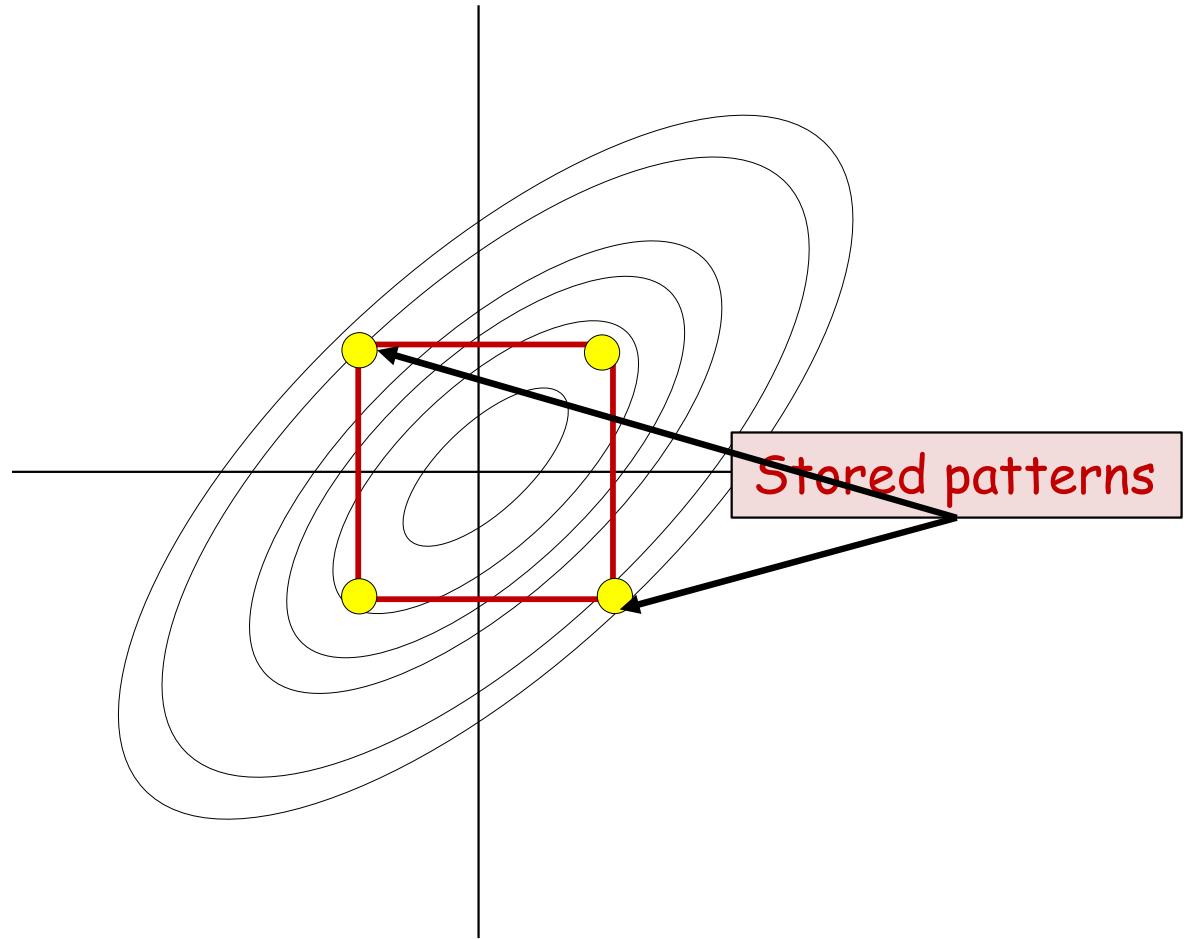
$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$



- $E$  is a concave quadratic
  - Shown from above (assuming 0 bias)
- The minima will lie on the boundaries of the hypercube
  - But components of  $\mathbf{y}$  can only take values  $\pm 1$
  - I.e.  $\mathbf{y}$  lies on the corners of the unit hypercube

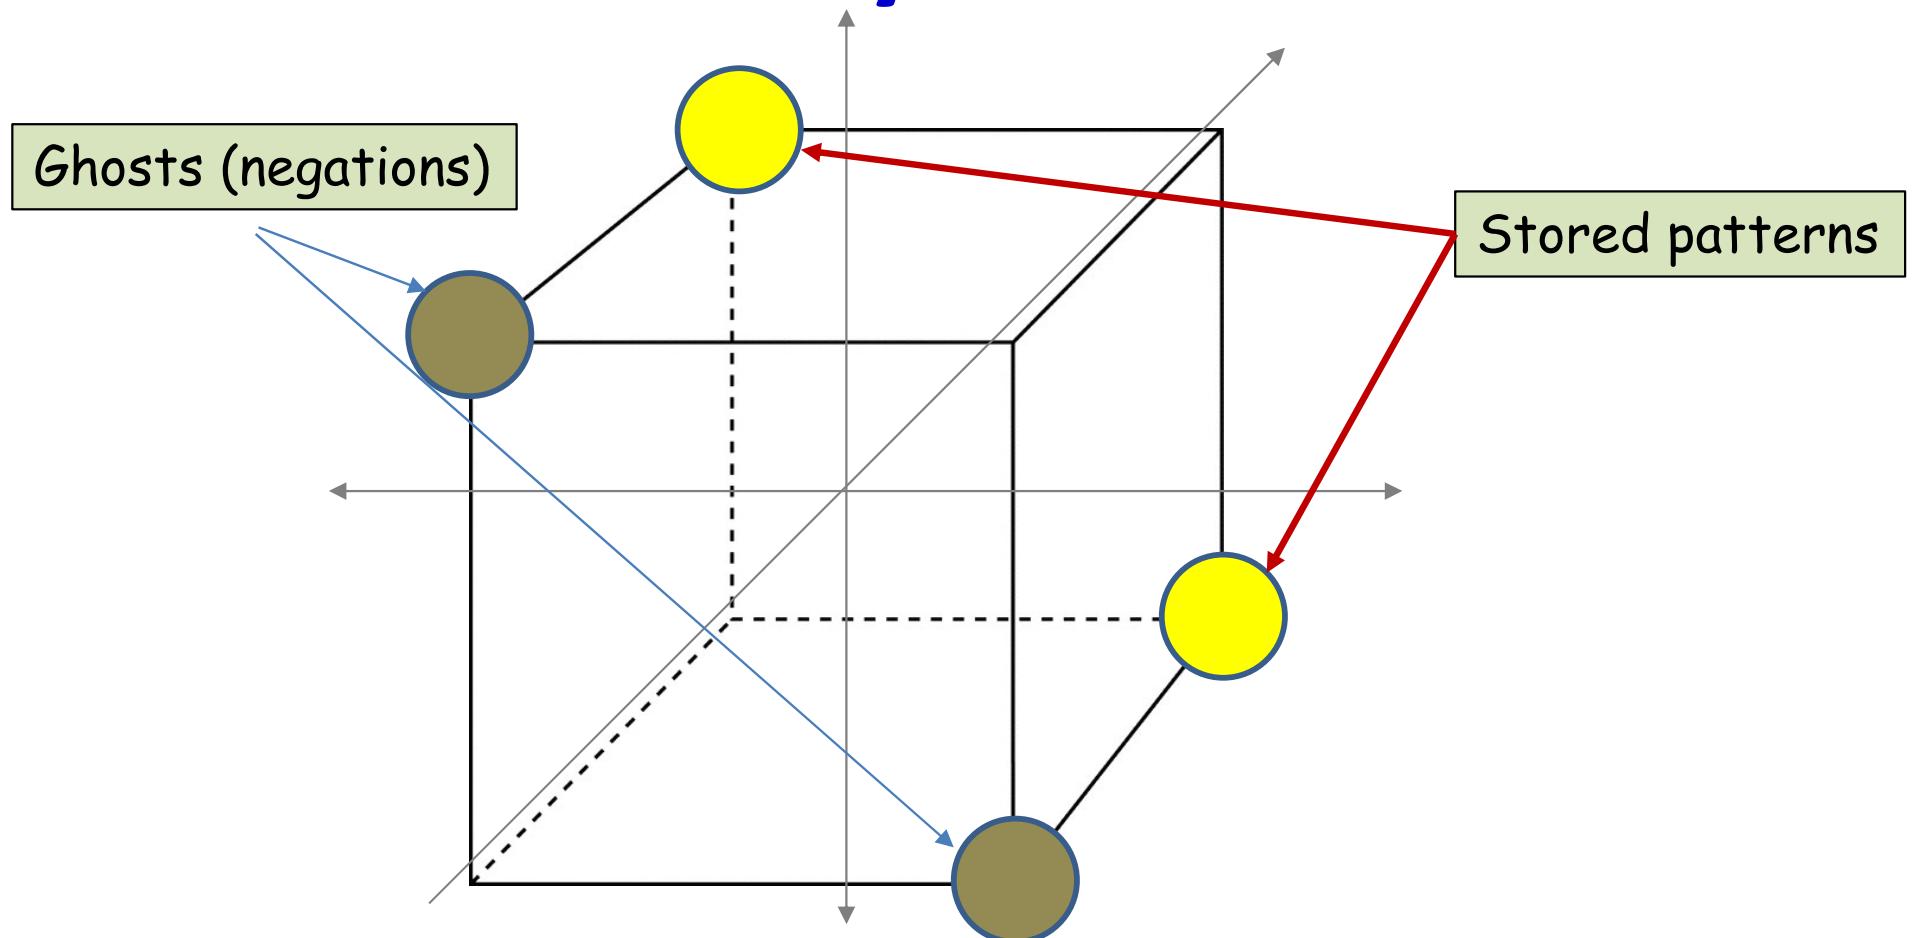
# The energy function

$$E = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$



- The stored values of  $\mathbf{y}$  are the ones where all adjacent corners are lower on the quadratic

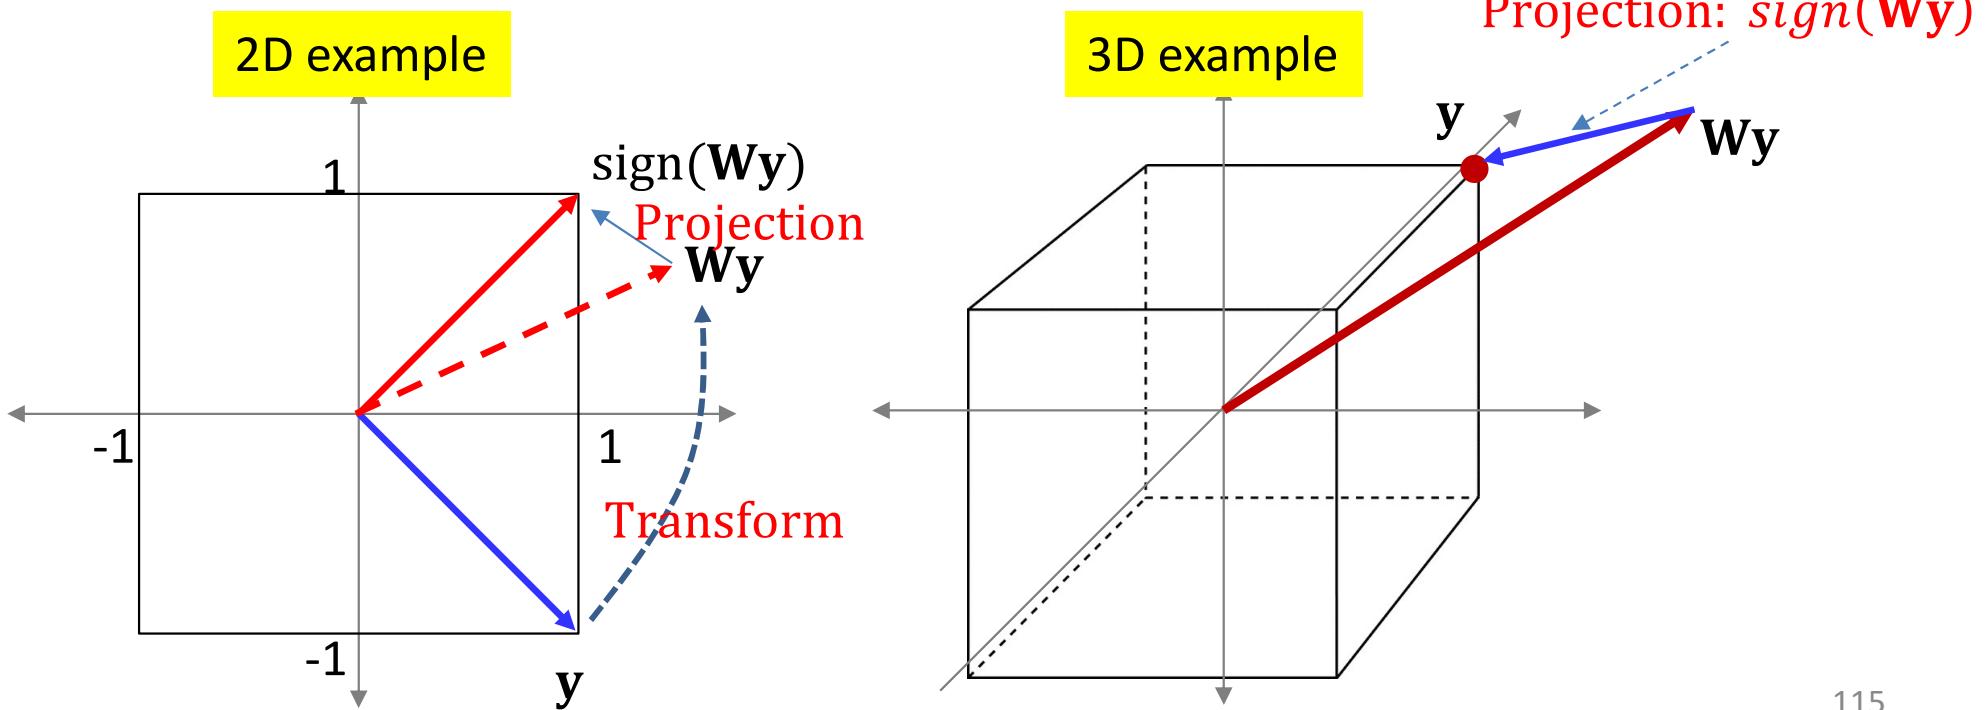
# Patterns you can store



- All patterns are on the corners of a hypercube
  - If a pattern is stored, its “ghost” is stored as well
  - Intuitively, patterns must ideally be maximally far apart
    - Though this doesn’t seem to hold for Hebbian learning

# Evolution of the network

- Note: for real vectors  $sign(\mathbf{y})$  is a projection
  - Projects  $\mathbf{y}$  onto the nearest corner of the hypercube
  - It “quantizes” the space into orthants
- Response to field:  $\mathbf{y} \leftarrow sign(\mathbf{W}\mathbf{y})$ 
  - Each step rotates the vector  $\mathbf{y}$  and then projects it onto the nearest corner



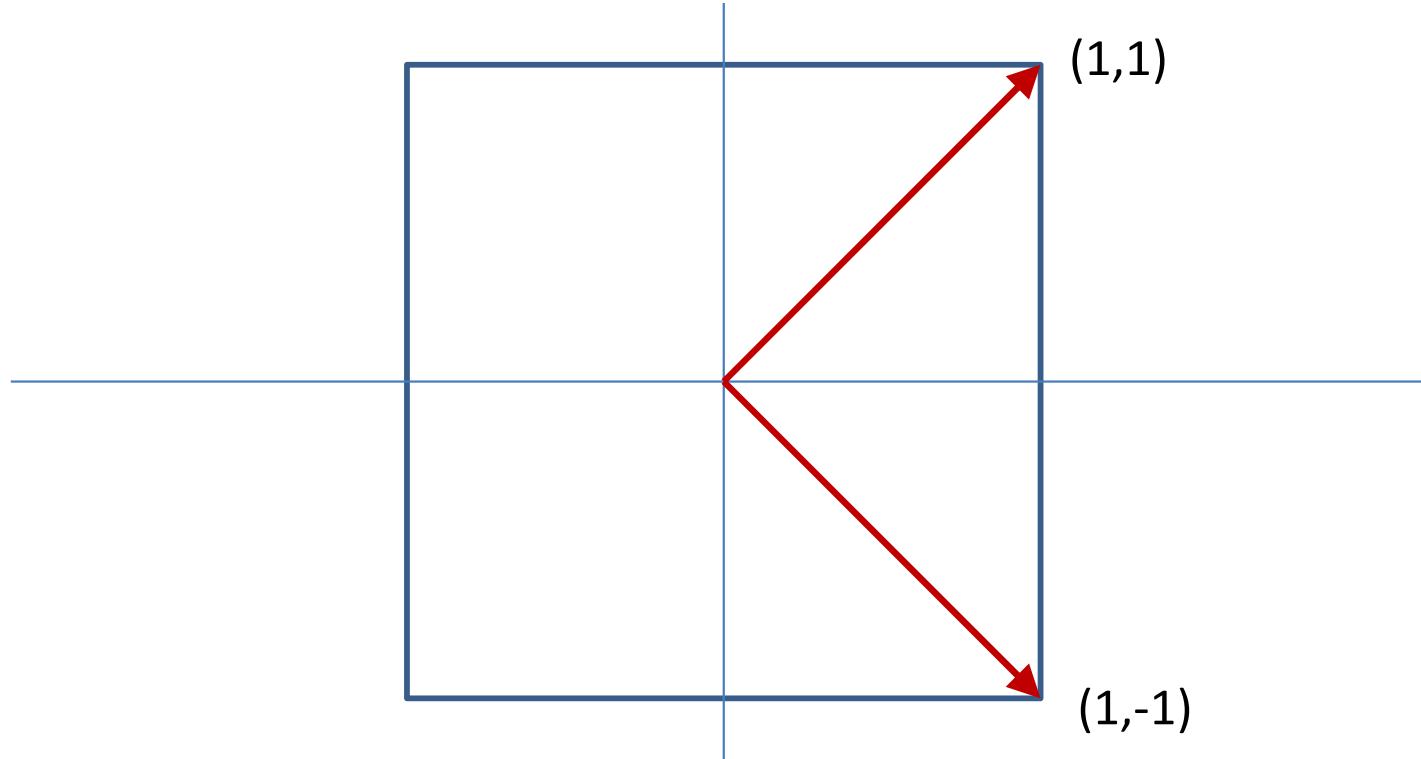
# Storing patterns

- A pattern  $\mathbf{y}_P$  is stored if:
  - $sign(\mathbf{W}\mathbf{y}_p) = \mathbf{y}_p$  for all target patterns
- Training: Design  $\mathbf{W}$  such that this holds
- Simple solution:  $\mathbf{y}_p$  is an Eigenvector of  $\mathbf{W}$ 
  - And the corresponding Eigenvalue is positive
$$\mathbf{W}\mathbf{y}_p = \lambda\mathbf{y}_p$$
  - More generally  $orthant(\mathbf{W}\mathbf{y}_p) = orthant(\mathbf{y}_p)$
- How many such  $\mathbf{y}_p$  can we have?

# Random fact that should interest you

- Number of ways of selecting two  $N$ -bit binary patterns  $y_1$  and  $y_2$  such that they differ from one another in exactly  $N/2$  bits is  $\mathcal{O}(2^{\frac{3N}{2}})$
- The size of the largest set of  $N$ -bit binary patterns  $\{y_1, y_2, \dots\}$  that *all* differ from one another in exactly  $N/2$  bits is at most  $N$ 
  - Trivial proof.. ☺

# Only N patterns?



- Symmetric weight matrices have orthogonal Eigen vectors
- You can have max  $N$  orthogonal vectors in an  $N$ -dimensional space

# random fact that should interest you

- The Eigenvectors of any symmetric matrix  $W$  are orthogonal
- The Eigenvalues may be positive or negative

# Storing more than one pattern

- Requirement: Given  $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_P$ 
  - Design  $\mathbf{W}$  such that
    - $sign(\mathbf{W}\mathbf{y}_p) = \mathbf{y}_p$  for all target patterns
    - There are no other *binary* vectors for which this holds
- What is the largest number of patterns that can be stored?

# Storing $K$ orthogonal patterns

- Simple solution: Design  $\mathbf{W}$  such that  $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_K$  are the Eigen vectors of  $\mathbf{W}$ 
  - Let  $\mathbf{Y} = [\mathbf{y}_1 \mathbf{y}_2 \dots \mathbf{y}_K]$

$$\mathbf{W} = \mathbf{Y}\Lambda\mathbf{Y}^T$$

- $\lambda_1, \dots, \lambda_K$  are positive
  - For  $\lambda_1 = \lambda_2 = \lambda_K = 1$  this is exactly the Hebbian rule
- The patterns are provably stationary

# Hebbian rule

- In reality
  - Let  $\mathbf{Y} = [\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K \ \mathbf{r}_{K+1} \ \mathbf{r}_{K+2} \ \dots \ \mathbf{r}_N]$

$$\mathbf{W} = \mathbf{Y} \Lambda \mathbf{Y}^T$$

- $\mathbf{r}_{K+1} \ \mathbf{r}_{K+2} \ \dots \ \mathbf{r}_N$  are orthogonal to  $\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K$
- $\lambda_1 = \lambda_2 = \lambda_K = 1$
- $\lambda_{K+1}, \dots, \lambda_N = 0$

# Storing $N$ orthogonal patterns

- When we have  $N$  orthogonal (or near orthogonal) patterns  $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N$

- $Y = [\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_N]$

$$\mathbf{W} = \mathbf{Y}\Lambda\mathbf{Y}^T$$

- $\lambda_1 = \lambda_2 = \lambda_N = 1$

- The Eigen vectors of  $\mathbf{W}$  span the space
- Also, for any  $\mathbf{y}_k$

$$\mathbf{W}\mathbf{y}_k = \mathbf{y}_k$$

# Storing $N$ orthogonal patterns

- The  $N$  orthogonal patterns  $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N$  *span the space*
- Any pattern  $\mathbf{y}$  can be written as

$$\mathbf{y} = a_1 \mathbf{y}_1 + a_2 \mathbf{y}_2 + \cdots + a_N \mathbf{y}_N$$

$$\mathbf{W}\mathbf{y} = a_1 \mathbf{W}\mathbf{y}_1 + a_2 \mathbf{W}\mathbf{y}_2 + \cdots + a_N \mathbf{W}\mathbf{y}_N$$

$$= a_1 \mathbf{y}_1 + a_2 \mathbf{y}_2 + \cdots + a_N \mathbf{y}_N = \mathbf{y}$$

- All patterns are stationary
  - Everything is a stationary memory
  - Completely useless network

# Storing $K$ orthogonal patterns

- Even if we store fewer than  $N$  patterns

- Let  $Y = [\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K \ \mathbf{r}_{K+1} \ \mathbf{r}_{K+2} \ \dots \ \mathbf{r}_N]$

$$W = Y \Lambda Y^T$$

- $\mathbf{r}_{K+1} \ \mathbf{r}_{K+2} \ \dots \ \mathbf{r}_N$  are orthogonal to  $\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K$
  - $\lambda_1 = \lambda_2 = \lambda_K = 1$
  - $\lambda_{K+1}, \dots, \lambda_N = 0$
- Any pattern that is *entirely* in the subspace spanned by  $\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K$  is also stable (same logic as earlier)
- Only patterns that are *partially* in the subspace spanned by  $\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K$  are unstable
  - Get projected onto subspace spanned by  $\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K$

# Problem with Hebbian Rule

- Even if we store fewer than  $N$  patterns
  - Let  $Y = [\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K \ \mathbf{r}_{K+1} \ \mathbf{r}_{K+2} \ \dots \ \mathbf{r}_N]$

$$W = Y\Lambda Y^T$$

- $\mathbf{r}_{K+1} \ \mathbf{r}_{K+2} \ \dots \ \mathbf{r}_N$  are orthogonal to  $\mathbf{y}_1 \ \mathbf{y}_2 \ \dots \ \mathbf{y}_K$
- $\lambda_1 = \lambda_2 = \lambda_K = 1$
- Problems arise because Eigen values are all 1.0
  - Ensures stationarity of vectors in the subspace
  - All stored patterns are equally important
  - What if we get rid of this requirement?

# Hebbian rule and general (non-orthogonal) vectors

$$w_{ji} = \sum_{p \in \{p\}} y_i^p y_j^p$$

- What happens when the patterns are *not* orthogonal
- What happens when the patterns are presented *more* than once
  - Different patterns presented different numbers of times
  - Equivalent to having unequal Eigen values..
- Can we predict the evolution of any vector  $\mathbf{y}$ 
  - Hint: For real valued vectors, use Lanczos iterations
    - Can write  $\mathbf{Y}_P = \mathbf{U}_P \Lambda \mathbf{V}_P^T$ ,  $\rightarrow \mathbf{W} = \mathbf{U}_P \Lambda^2 \mathbf{U}_P^T$
    - Tougher for binary vectors (NP)

# The bottom line

- With a network of  $N$  units (i.e.  $N$ -bit patterns)
- The maximum number of stationary patterns is actually *exponential* in  $N$ 
  - McEliece and Posner, 84'
  - E.g. when we had the Hebbian net with  $N$  orthogonal base patterns, *all* patterns are stationary
- For a *specific* set of  $K$  patterns, we can *always* build a network for which all  $K$  patterns are stable provided  $K \leq N$ 
  - Mostafa and St. Jacques 85'
    - For large  $N$ , the upper bound on  $K$  is actually  $N/4\log N$ 
      - McEliece et. Al. 87'
    - **But this may come with many “parasitic” memories**

# The bottom line

- With an network of  $N$  units (i.e.  $N$ -bit patterns)
- The maximum number of stable patterns is actually *exponential* in  $N$ 
  - McEliece and Posner, 84'
  - E.g. when we had the  $K = N$  base patterns, *all* patterns are stable
- For a *specific* set of  $K$  patterns, we can *always* build a network for which all  $K$  patterns are stable provided  $K \leq N$ 
  - Mostafa and St. Jacques 85'
    - For large  $N$ , the upper bound on  $K$  is actually  $N/4\log N$
    - McEliece et. Al. 87'
  - **But this may come with many “parasitic” memories**

How do we find this  
network?

# The bottom line

- With an network of  $N$  units (i.e.  $N$ -bit patterns)
- The maximum number of stable patterns is actually *exponential* in  $N$ 
  - McEliece and Posner, 84'
  - E.g. when we had the  $K = N$  base patterns, *all* patterns are stable
- For a *specific* set of  $K$  patterns, we can *always* build a network for which all  $K$  patterns are stable provided  $K \leq N$ 
  - Mostafa and St. Jacques 85'
    - For large  $N$ , the upper bound on  $K$  is actually  $\approx N$
    - McEliece et. Al. 87'
  - But this may come with many “parasitic” memories

How do we find this  
network?

Can we do something  
about this?

# Story so far

- Hopfield nets with  $N$  neurons can store up to  $0.14N$  random patterns through Hebbian learning with 0.996 probability of recall
  - The recalled patterns are the Eigen vectors of the weights matrix with the highest Eigen values
- Hebbian learning assumes all patterns to be stored are equally important
  - For orthogonal patterns, the patterns are the Eigen vectors of the constructed weights matrix
  - All Eigen values are identical
- In theory the number of stationary states in a Hopfield network can be exponential in  $N$
- The number of *intentionally* stored patterns (stationary *and* stable) can be as large as  $N$ 
  - But comes with many parasitic memories

